

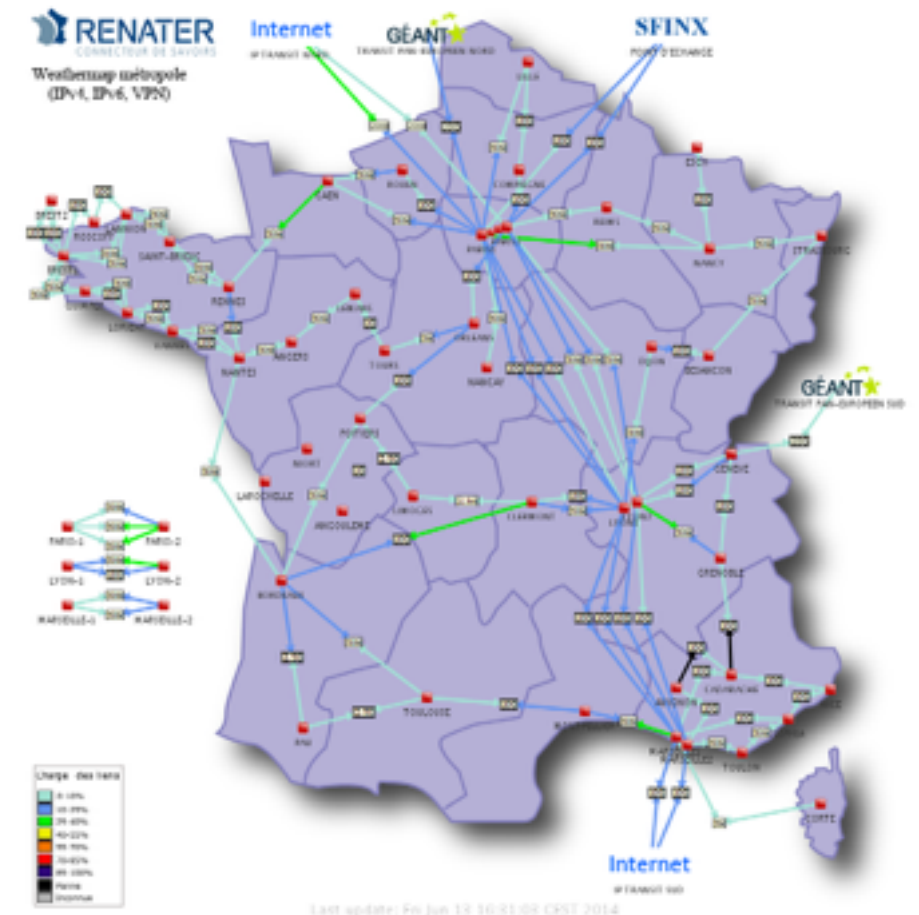
# Distributing OpenStack

To build geographically distributed clouds

Jonathan Pastor  
*[jonathan.pastor@inria.fr](mailto:jonathan.pastor@inria.fr)*

# Context

- Cloud computing has become very popular.
- Ever-increasing demand => ever-increasing infrastructure size.
- PB: scalability, reliability, energy but also security, jurisdiction and network overhead.
- Decentralise the production of computing ressources (Discovery project, <http://beyondtheclouds.github.io/>).

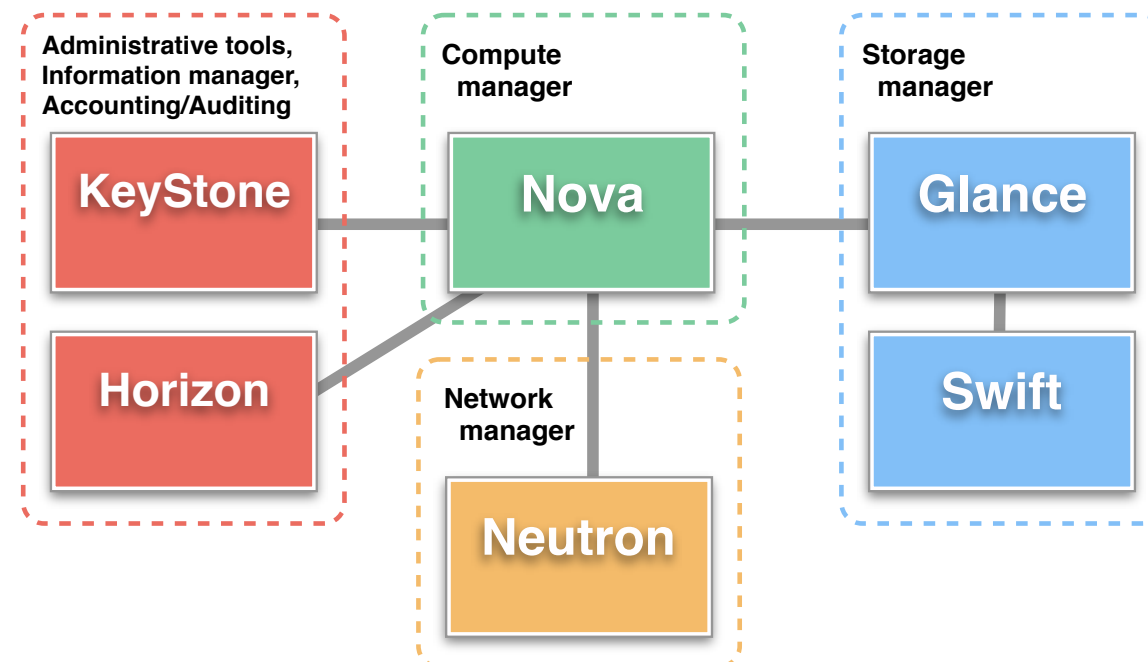


# LUC-OS

- Locality Based Utility Computing OS (LUC-OS):
- A **fully distributed** Cloud-OS that enables to use and operate a massively distributed infrastructure at WAN scale, leveraging **locality properties** in order to organise efficient **cooperations**.
- To address fault tolerance and energy concerns.
- To address the network overhead, micro/nano DCs will be located on ISP point of presence.

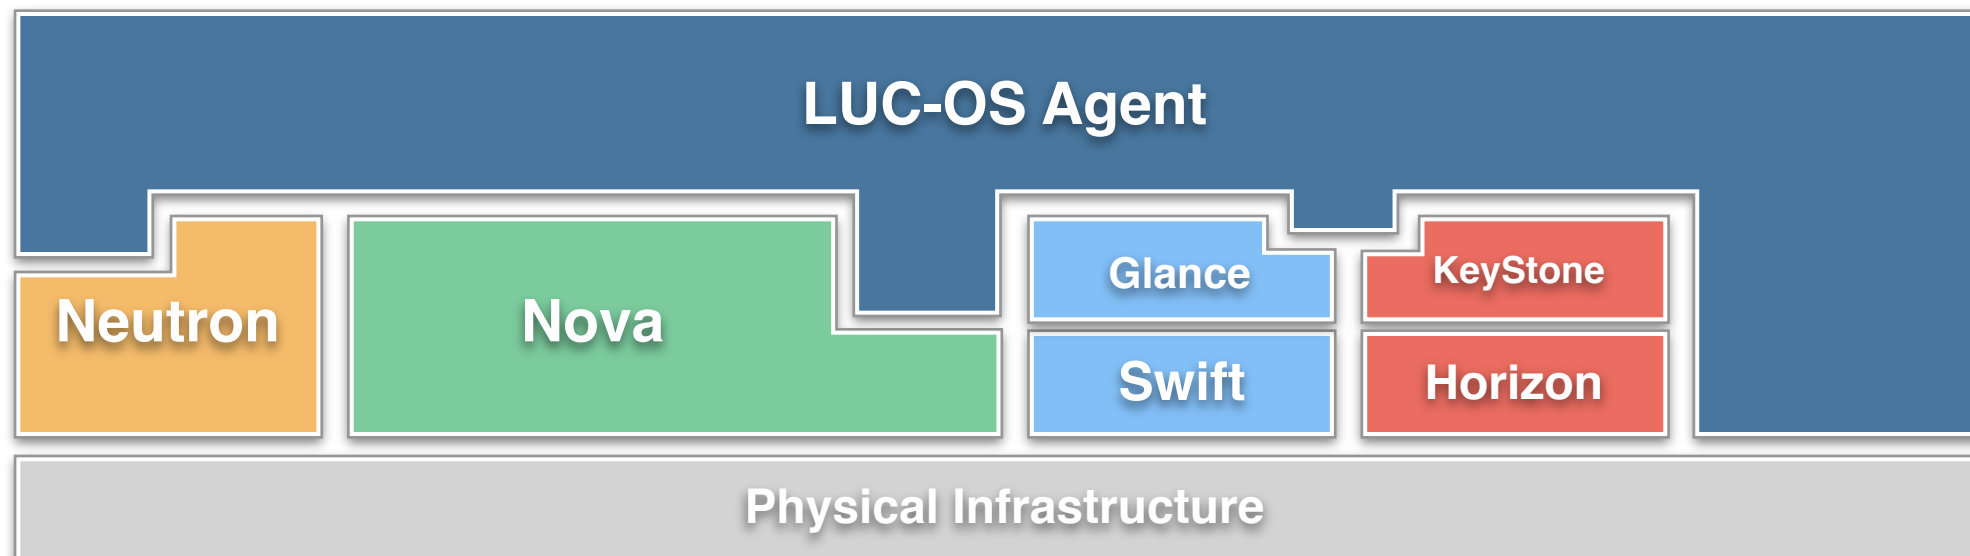
# OpenStack

- OpenStack is an IaaS manager of choice.
- It leverages an architecture composed of several services.
- Nova is the controlling service of OpenStack.



# Designing the LUC-OS on top of OpenStack

- The LUC-OS will rely on a multi-agent architecture.
- Some services of the LUC-OS may entirely reuse implementation from OpenStack (**Swift**).  
[https://www.swiftstack.com/docs/admin/cluster\\_management/regions.html](https://www.swiftstack.com/docs/admin/cluster_management/regions.html)
- Some services will “adapt” OpenStack to the LUC-OS (**Nova**).




# How to distribute OpenStack?

- We focus on distributing Nova
- Several ways to reach this objective:
  - 1) Leverage Classic HA deployment to build distributed Clouds (Standard).
  - 2) Using Cells (Cern).
  - 3) Replacing MySQL.

# Classic HA deployment

- Nova uses relational databases as local backend (sqlalchemy).
- Database is replicated and synchronised on each nova controller nodes.
- Proposal: Use HA and deploy at least one node controller per site.



<https://github.com/madkiss/openstackinaction4>

# Classic HA deployment

+

- Already working (with few controller nodes)

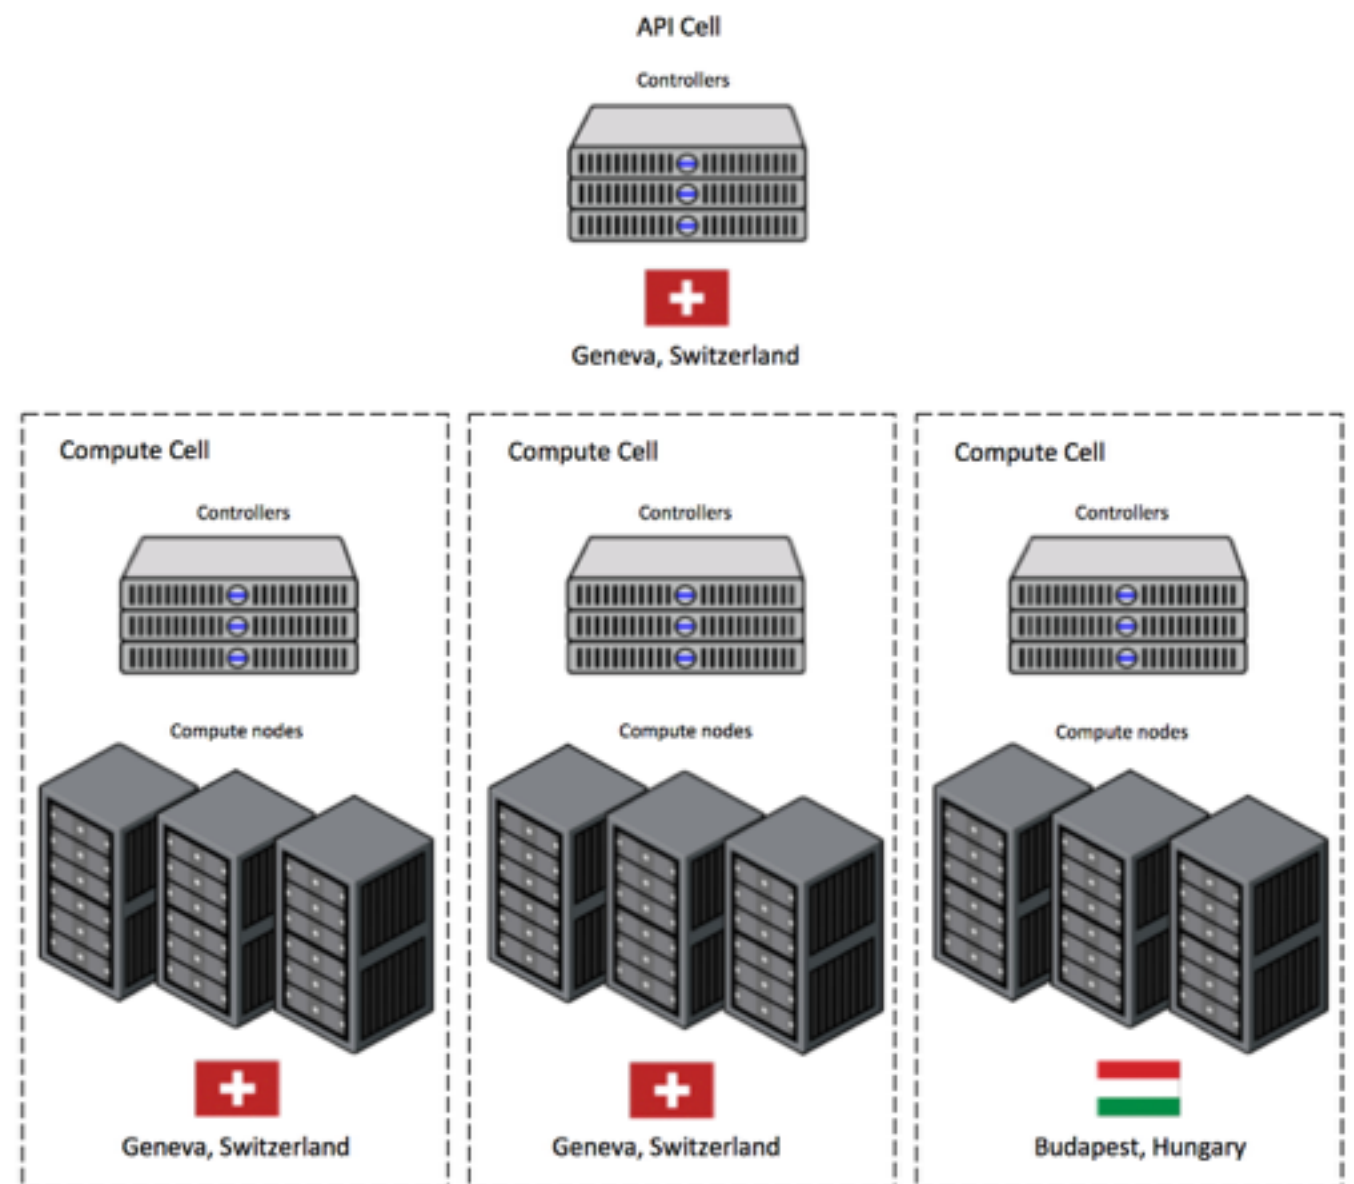
-

- **Scalability:**  
“**n controller nodes** means that for each write, there will be **n-1 writes** of other nodes”



# Using Cells (Cern)

- Hierarchical structure: a top cell contains nova-api and children cells contains remaining services including MySQL.
- Already used in Cern infrastructure:
  - *50 000 cores*
  - *3 different sites*
- <http://openstack-in-production.blogspot.fr/2014/03/cern-cloud-architecture-update-for.html>



# Using Cells (Cern)

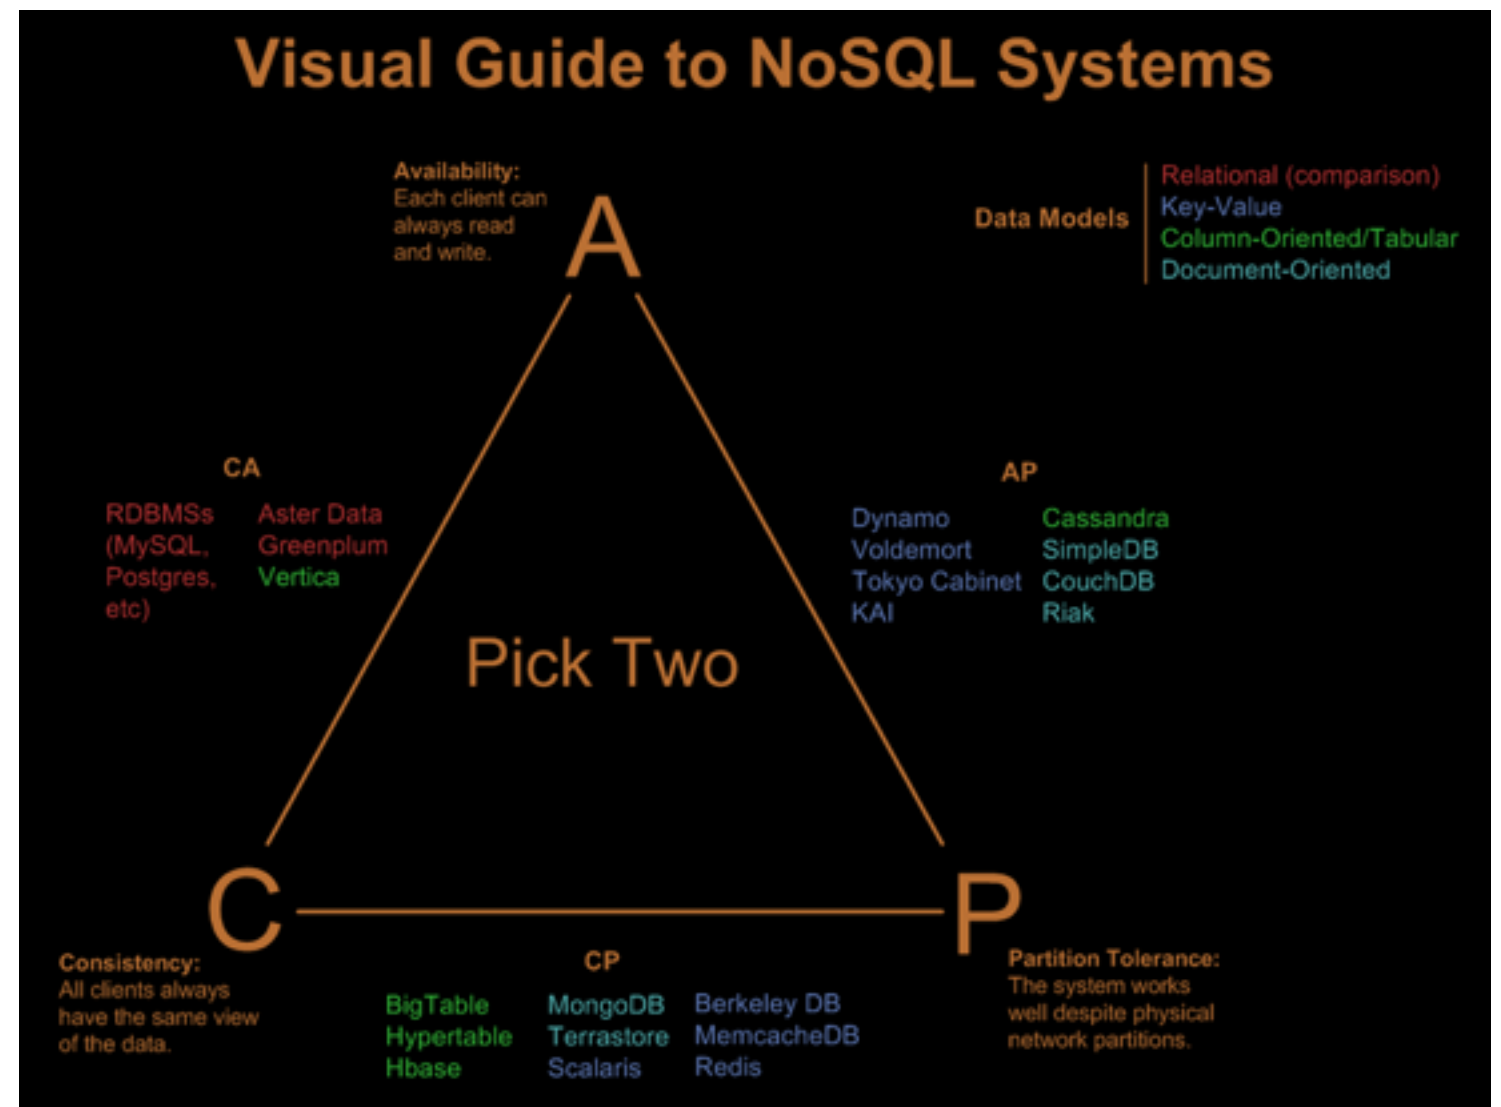
- It is possible to build the LUC-OS (flat infrastructure) on top of cells:
- **n servers** means **1 top cell**, and **n-1 cells that contains nova-controller and nova-compute**.
- It doesn't fit with the LUC-OS needs: top cell is a centralised (scalability, SPOF, ...).
- Proposal: distributing the top cell and cells become the pivot between the LUC-OS and cloud infrastructure

# Using Cells (Cern)

- ***Drawbacks:***
  - Require to modify the code of top cell (fork of OpenStack cells and follow each modification).
- ***Long term:***
  - Require to extend the protocol of cells each time we add capabilities to Nova compute (LUC-OS).

# Replacing MySQL

- MySQL + Galera:  
Availability and consistency.
- Do we need consistency or eventual consistency?
- Replacement of MySQL by a NoSQL solution.



# Replacing MySQL


- Nova services doesn't access directly the database: they calls functions located in:  
**nova/nova/db/api.py**
- These function manipulates databases and return python objects from:

```
/Users/jonathan/vm3/stack/nova/nova/objects
__init__.py      external_event.py  instance_fault.py  quotas.py
aggregate.py     fields.py          instance_group.py  security_group.py
base.py          fixed_ip.py        instance_info_cache.py security_group_rule.py
block_device.py  flavor.py          keypair.py         service.py
compute_node.py  floating_ip.py     migration.py       virtual_interface.py
dns_domain.py    instance.py        network.py
ec2.py           instance_action.py pci_device.py
```

# Replacing MySQL

- Nova already include the possibility of using a different db backend (only sqlalchemy yet).
- **nova/nova/db/api.py** => provides functions that manipulates the database backend

```
def instance_get(context, instance_id, columns_to_join=None):  
    """Get an instance or raise if it does not exist."""  
    print "Hello from ->instance_get"  
    return IMPL.instance_get(context, instance_id,  
                             columns_to_join=columns_to_join)
```



- It means that we can add a new backend:

```
_BACKEND_MAPPING = {  
    'sqlalchemy': 'nova.db.sqlalchemy.api',  
    'discovery': 'nova.db.discovery.api'  
}
```

# Replacing MySQL

- We propose to implement all functions contained in **nova/nova/db/api.py** to use a NoSQL database such as Riak, Scalaris or Cassandra.
- As current implementation (sqlalchemy) stores object with relation, we have to keep these relations with the new implementation.
- This would enable to use existing HA deployment (cf slide 4) without the active replication.

# Conclusion

- We propose to replace MySQL by a NoSQL solution.
- If this works with Nova, we can extend this to KeyStone.
- No dependency with the Cell concept.



# Bibliography

- [https://www.swiftstack.com/docs/admin/cluster\\_management/regions.html](https://www.swiftstack.com/docs/admin/cluster_management/regions.html)
- <https://github.com/madkiss/openstackinaction4>
- <http://openstack-in-production.blogspot.fr/2014/03/cern-cloud-architecture-update-for.html>