

# OpenStack on the Edge BoF Session - Boston 2017

Adrien Lebre

Fog Computing / Edge Computing/ Massively Distributed Clouds Working Group

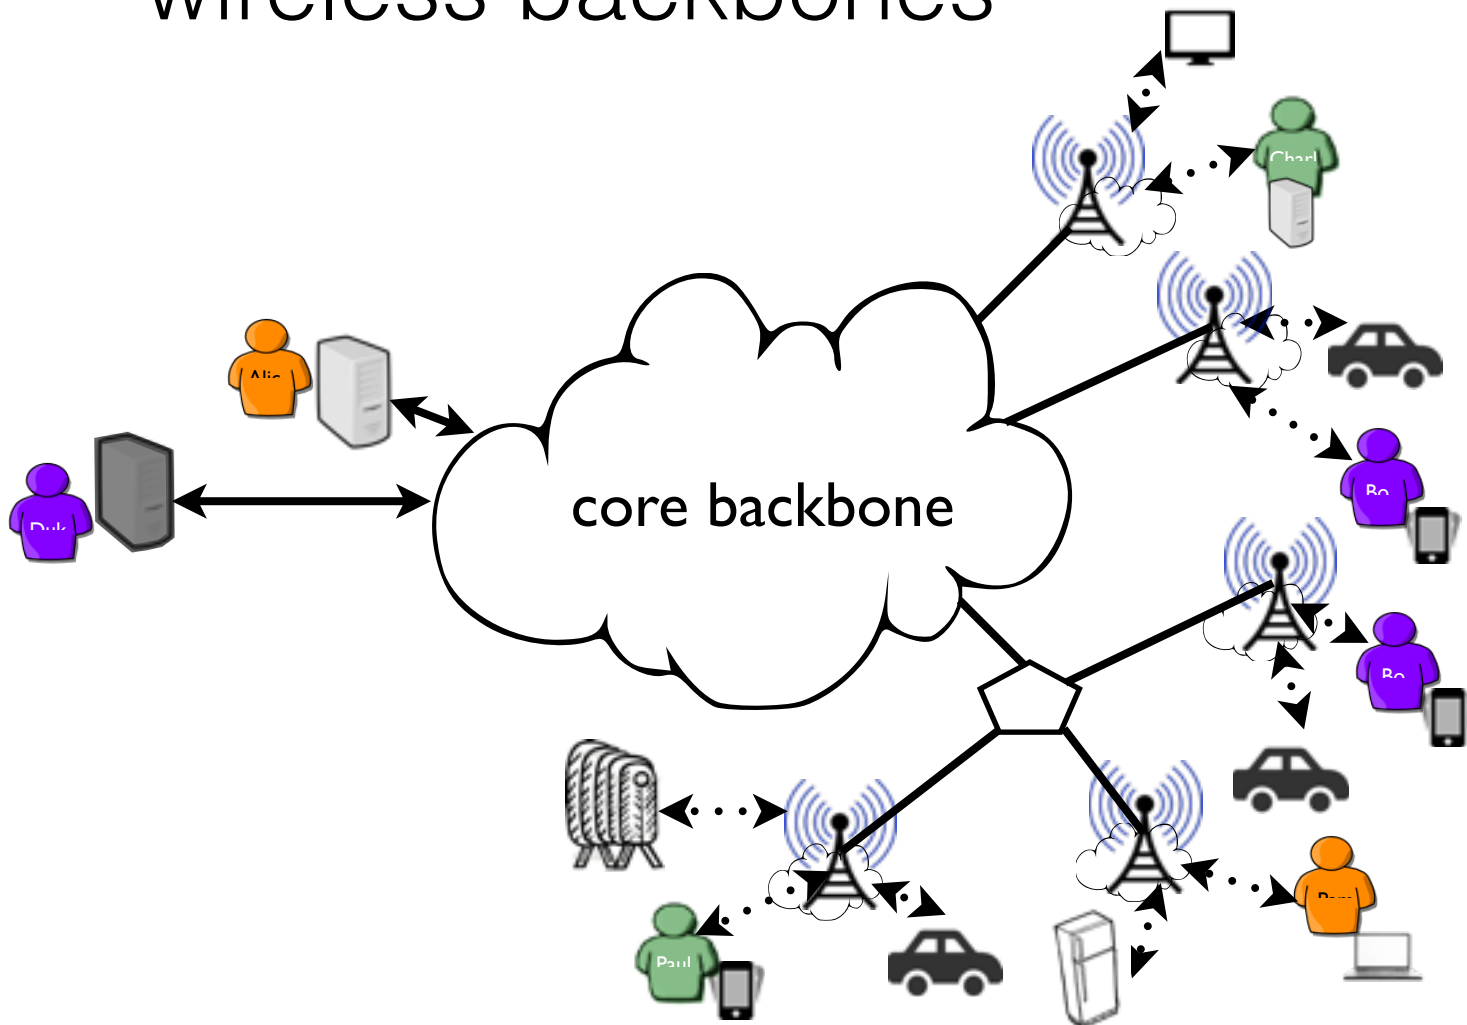
<https://etherpad.openstack.org/p/BOS-Fog-Edge-MassivelyDistributed-BoF>  
(or just go on google and looks for "Fog Edge OpenStack")

# Fog/Edge/MDC Infrastructures

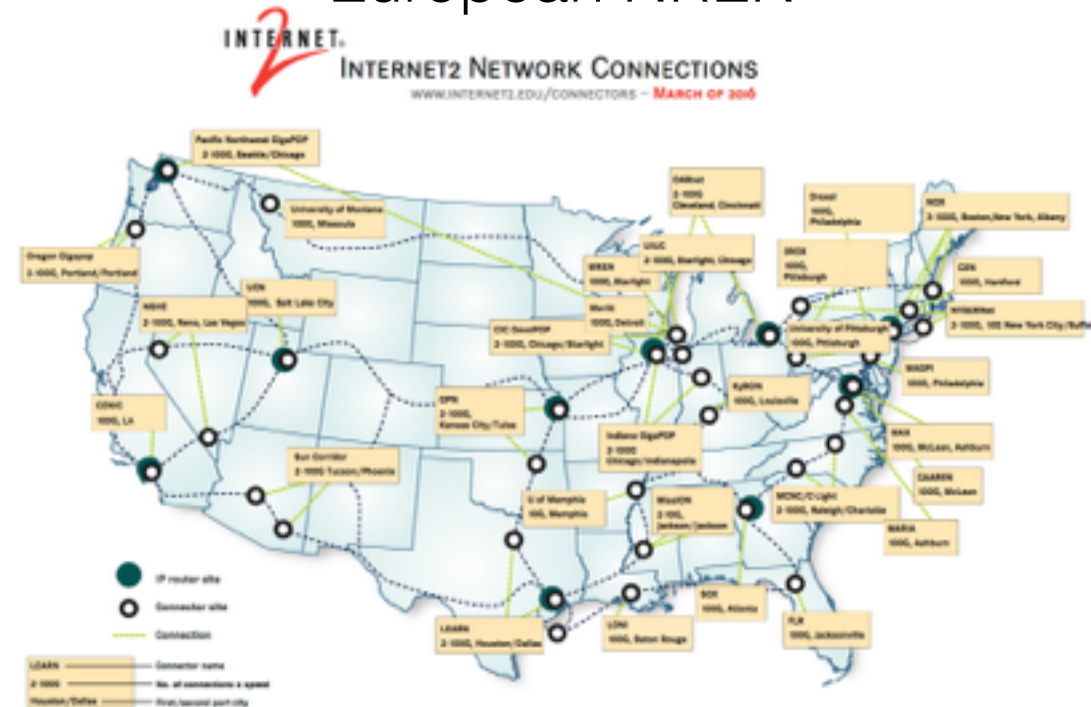
- Leverage network backbones

Extend any point of presence of network backbones (aka PoP) with servers (from network hubs up to major DSLAMs that are operated by telecom companies, network institutions...).

- Extend to the edge by including wireless backbones



# European NREN



USA NREN



# Micro/Nano DCs



Deployment of a PoP of the Orange French backbone




Sagrada Familia microDC  
(Barcelona, Spain)



MDC Industry - Brazil



# Micro/Nano DCs



Localized or **micro data centers** are a fact of life, but by **applying a self-contained, scalable and remotely managed solution and process**, CIOs **can reduce costs, improve agility, and introduce new levels of compliance and service continuity.**

Creating micro data centers is something companies have done for years, but often in an ad hoc manner.

**Gartner 2015**




Sagrada Familia microDC  
(Barcelona, Spain)



MDC Industry - Brazil

# What's about Brokering Approaches?


- Sporadic (hybrid computing/cloud bursting) almost ready for production
- While standards are coming (OCCI...), current brokers are rather limited to simple usages and not advanced administration operations






# What's about Brokering Approaches?

- Sporadic (hybrid computing/cloud bursting) almost ready for production
- While standards are coming (OCCI...), current brokers are rather limited to simple usages and not advanced administration operations



# What's about Brokering Approaches?


- Sporadic (hybrid computing/cloud bursting) almost ready for production
- While standards are coming (OCCI...), current brokers are rather limited to simple usages and not advanced administration operations



Advanced brokers must reimplement standard IaaS mechanisms while facing the API limitation

# Would OpenStack be the solution?


- Do not reinvent the wheel... it is too late





# Would OpenStack be the solution?

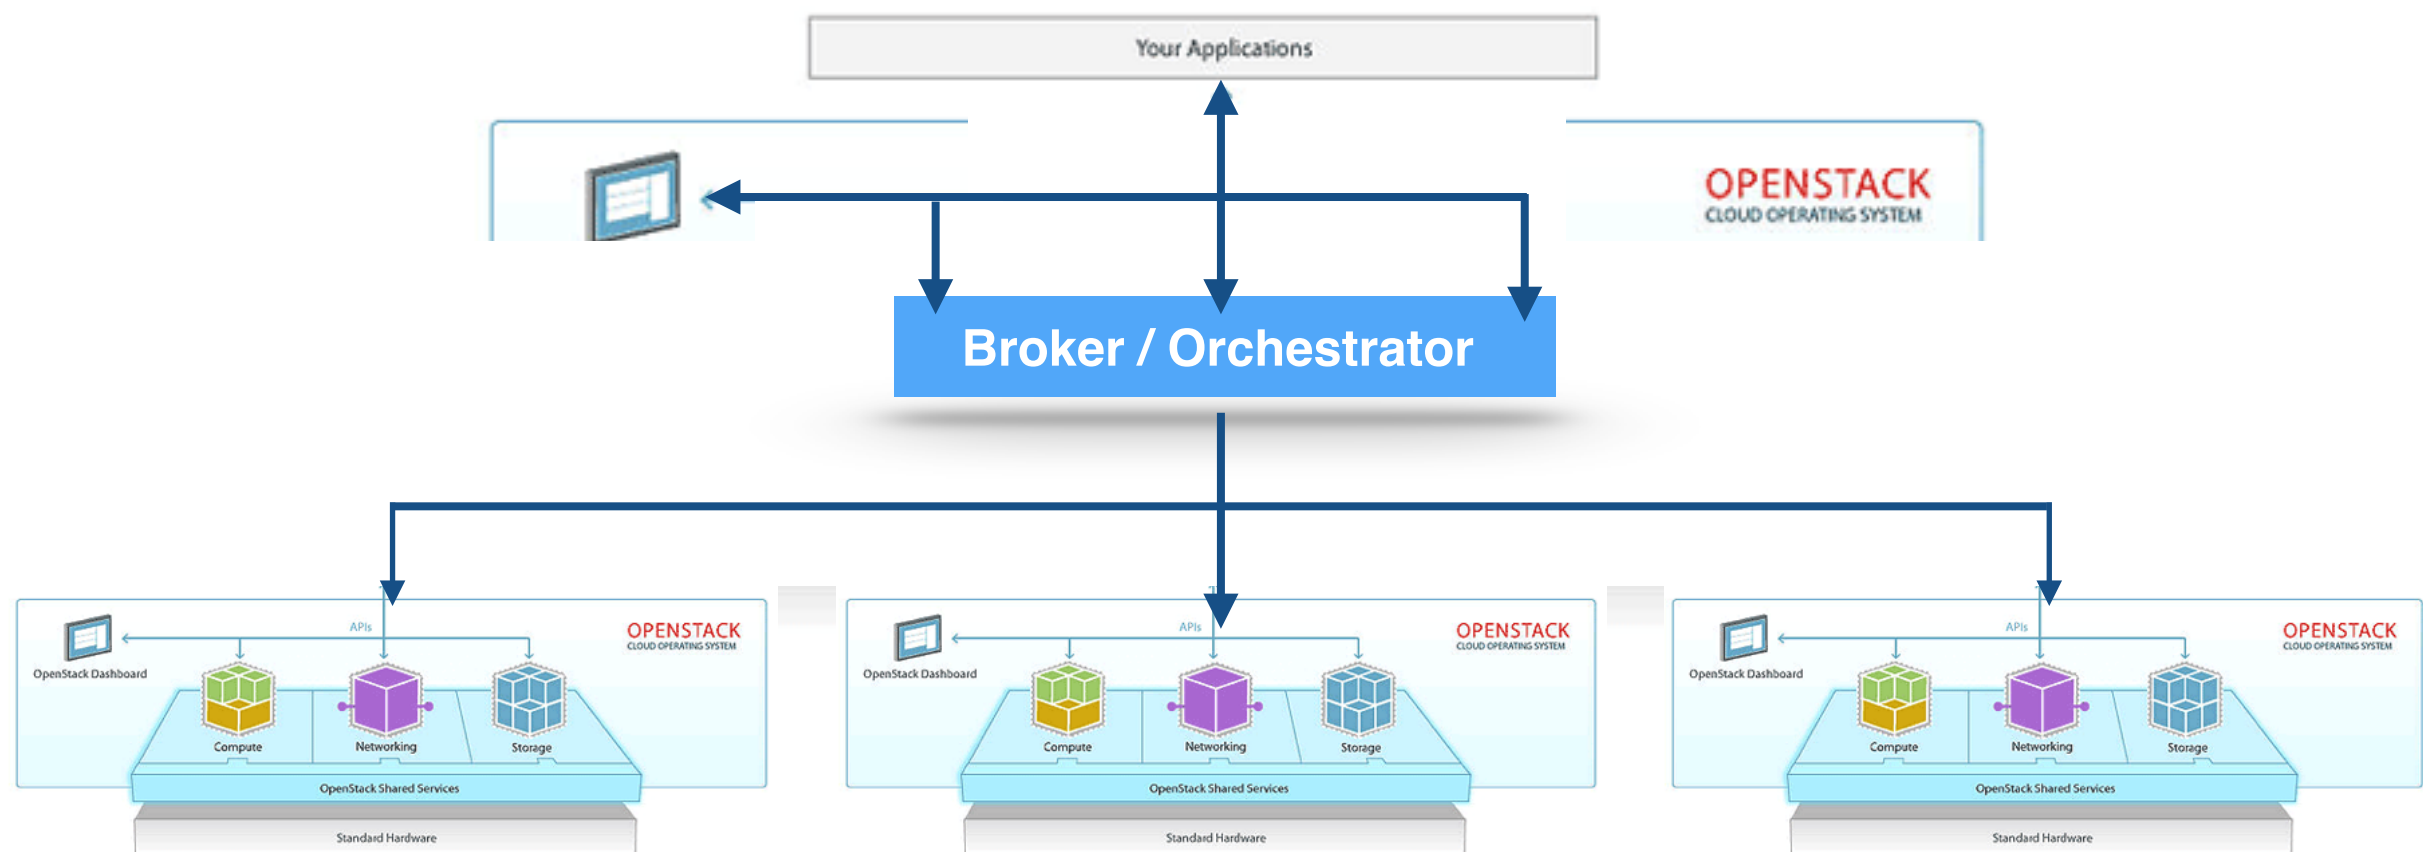
- Do not reinvent the wheel... it is too late



# Would OpenStack be the solution?

- Do not reinvent the wheel... it is too late
- Few proposals to federate/operate distinct OpenStack DCs

Top/Down: add a substrate to pilot independent OpenStack instances.

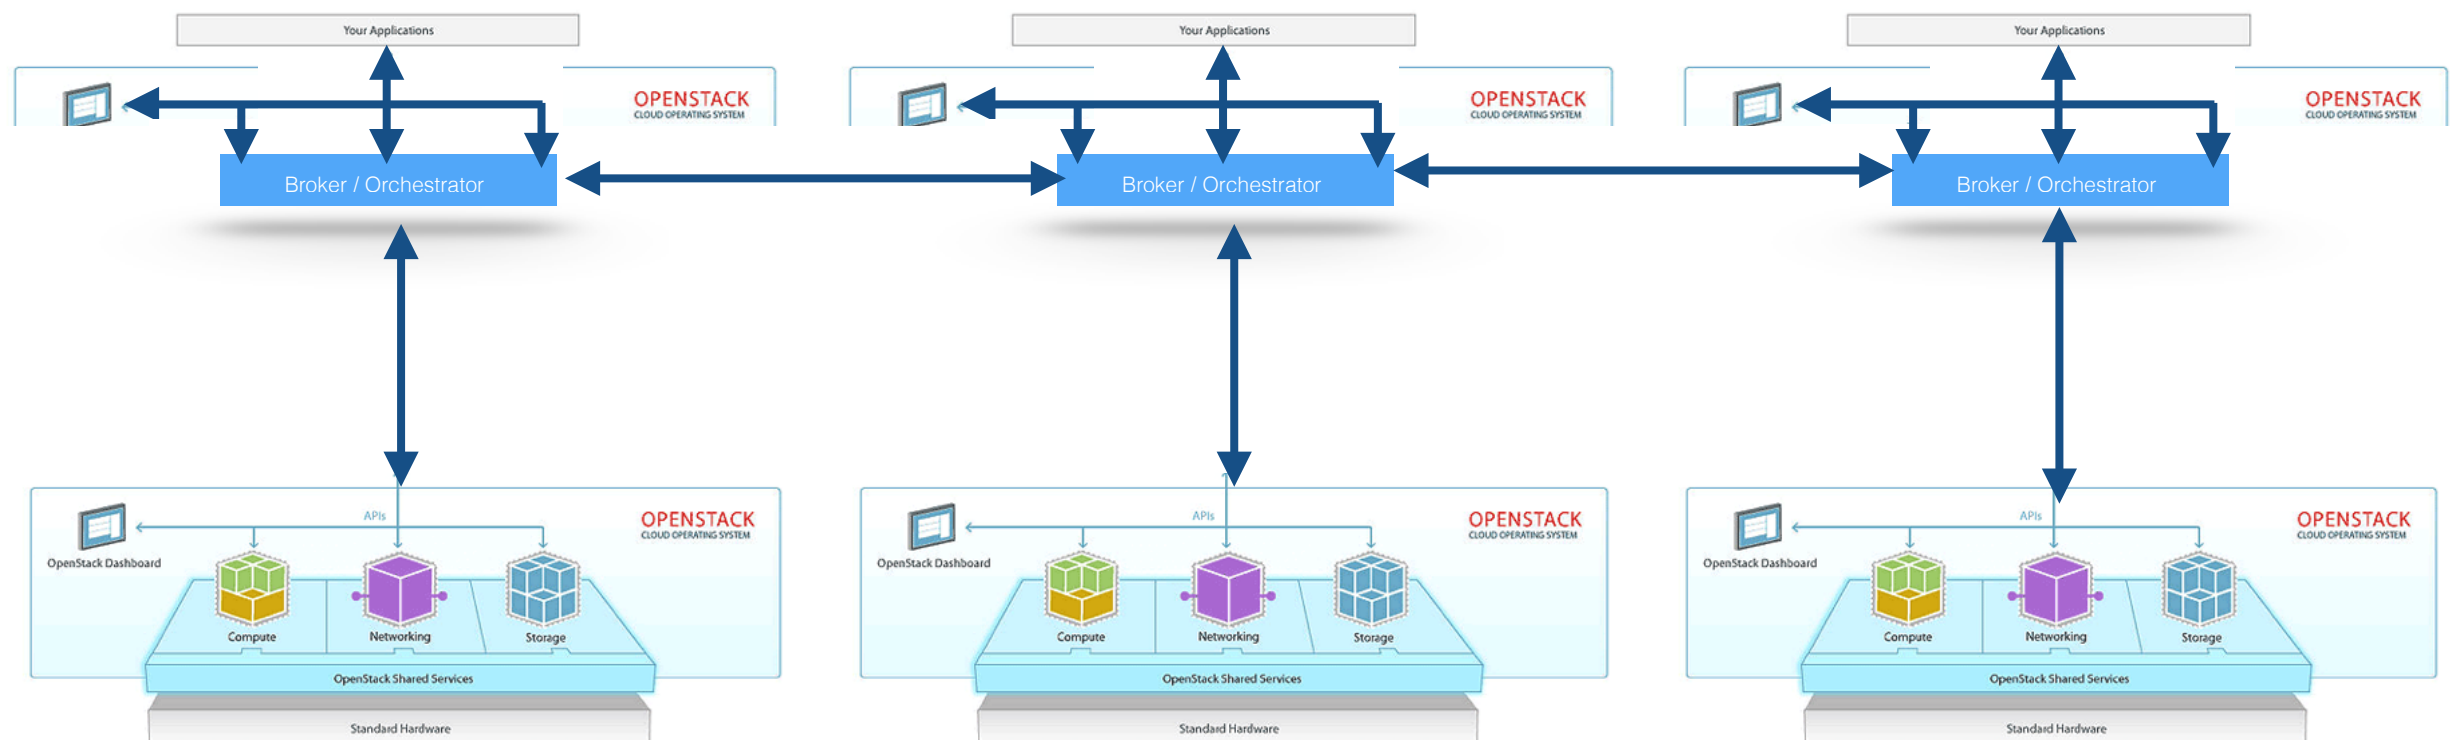




# Would OpenStack be the solution?

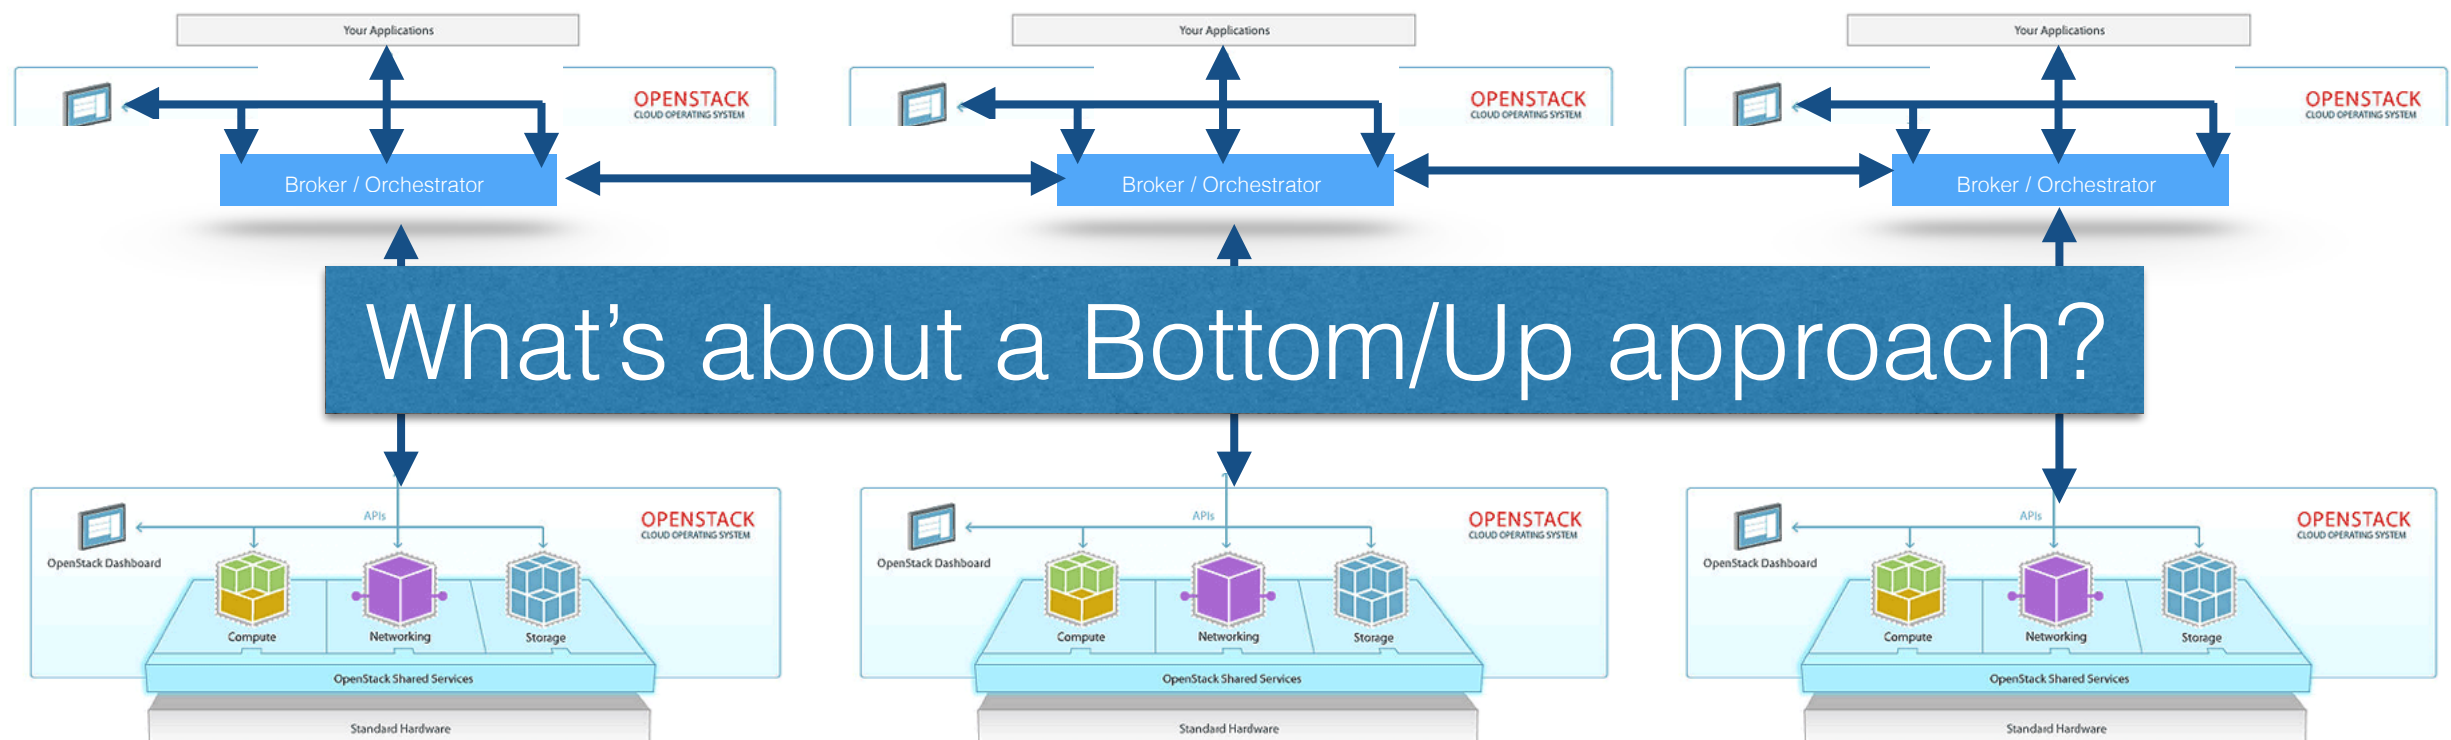
- Do not reinvent the wheel... it is too late
- Few proposals to federate/operate distinct OpenStack DCs

Top/Down: add a substrate to pilot independent OpenStack instances.



# Would OpenStack be the solution?

- Do not reinvent the wheel... it is too late
  - Few proposals to federate/operate distinct OpenStack DCs
- Top/Down: add a substrate to pilot independent OpenStack instances.






# Would OpenStack be the solution?

- Do not reinvent the wheel... it is too late
- Few proposals to federate/operate distinct OpenStack DCs

Bottom/Up - investigate whether/how OpenStack core services can be cooperative by default using Self\* and P2P mechanisms




Natively distributed/cooperative

# Would OpenStack be the solution?

- Do not reinvent the wheel... it is too late
- Few proposals to federate/operate distinct OpenStack DCs

Bottom/Up - investigate whether/how OpenStack core services can be cooperative by default using Self\* and P2P mechanisms



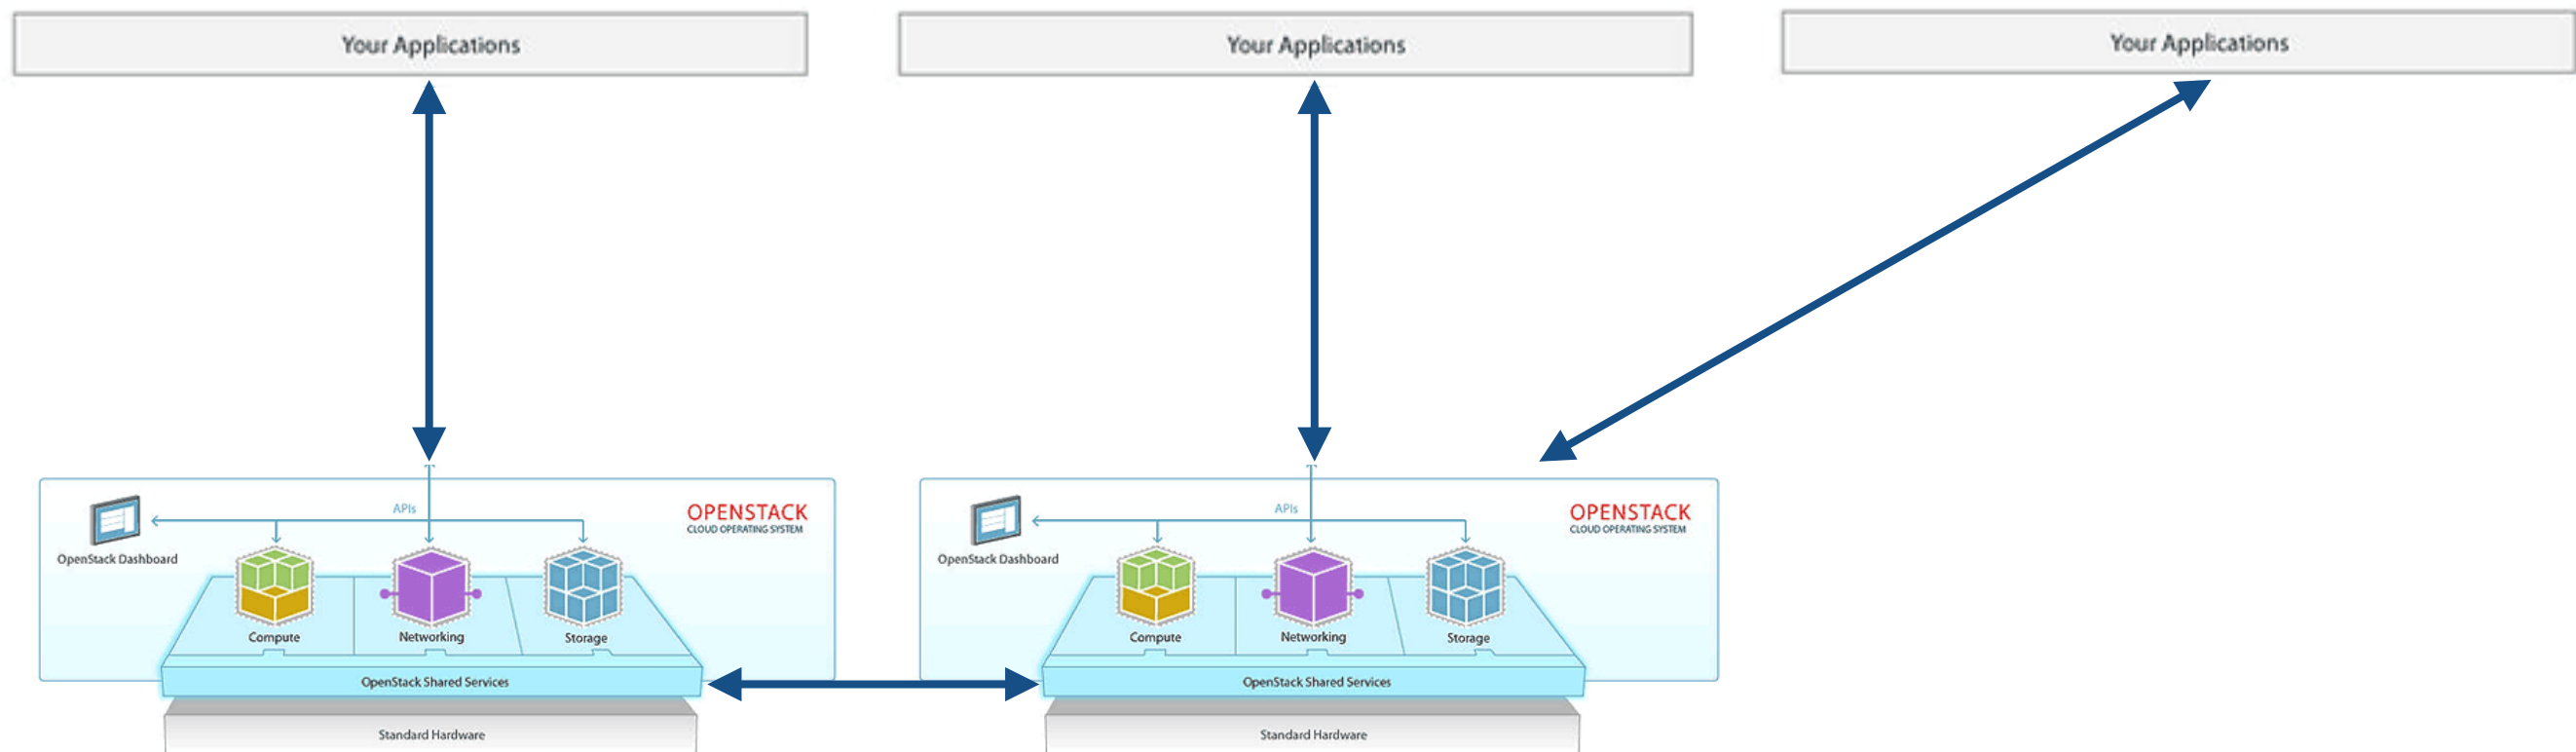
Natively distributed/cooperative



# Would OpenStack be the solution?

- Do not reinvent the wheel... it is too late
- Few proposals to federate/operate distinct OpenStack DCs

Bottom/Up - investigate whether/how OpenStack core services can be cooperative by default using Self\* and P2P mechanisms

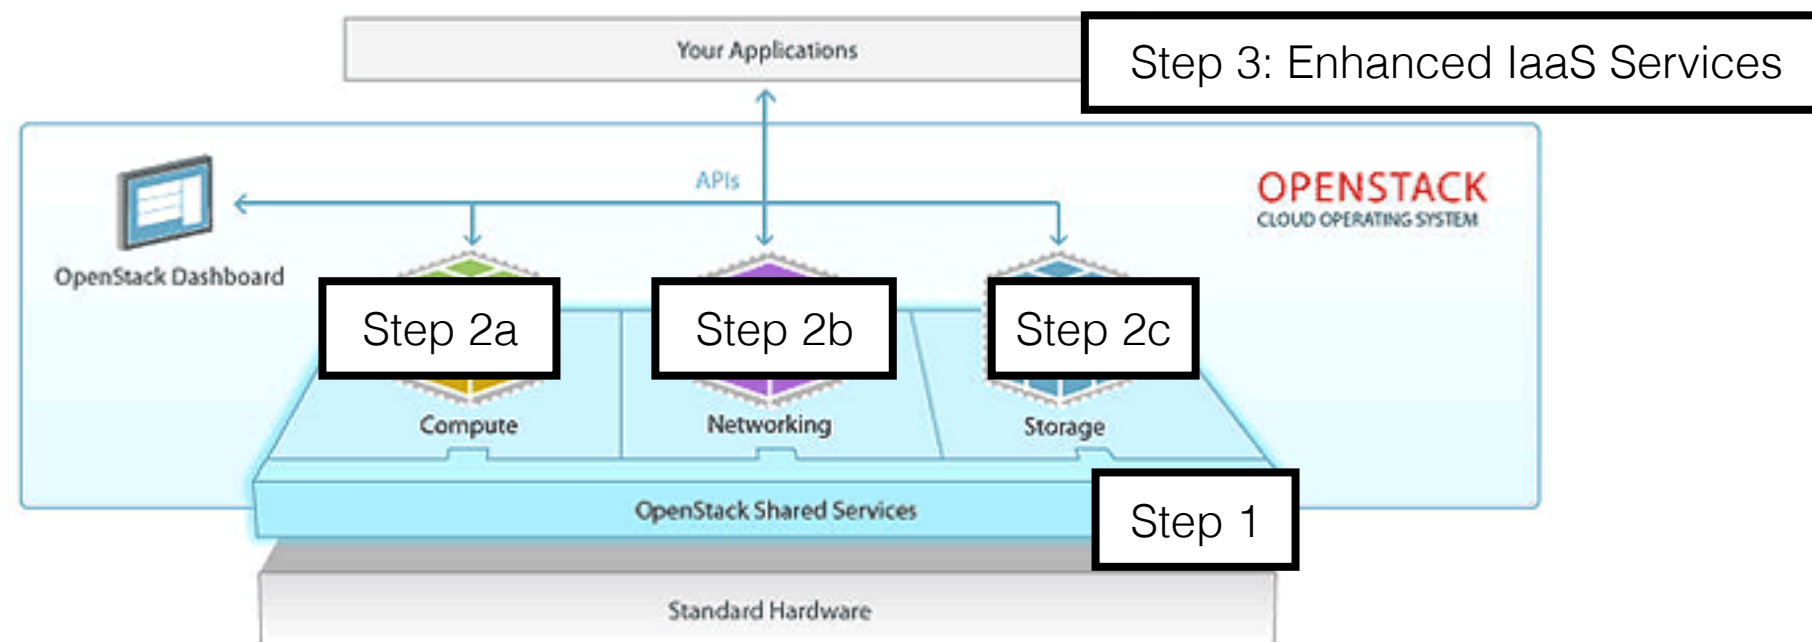


Natively distributed/cooperative

# Would OpenStack be the solution?

- Do not reinvent the wheel... it is too late
- Few proposals to federate/operate distinct OpenStack DCs

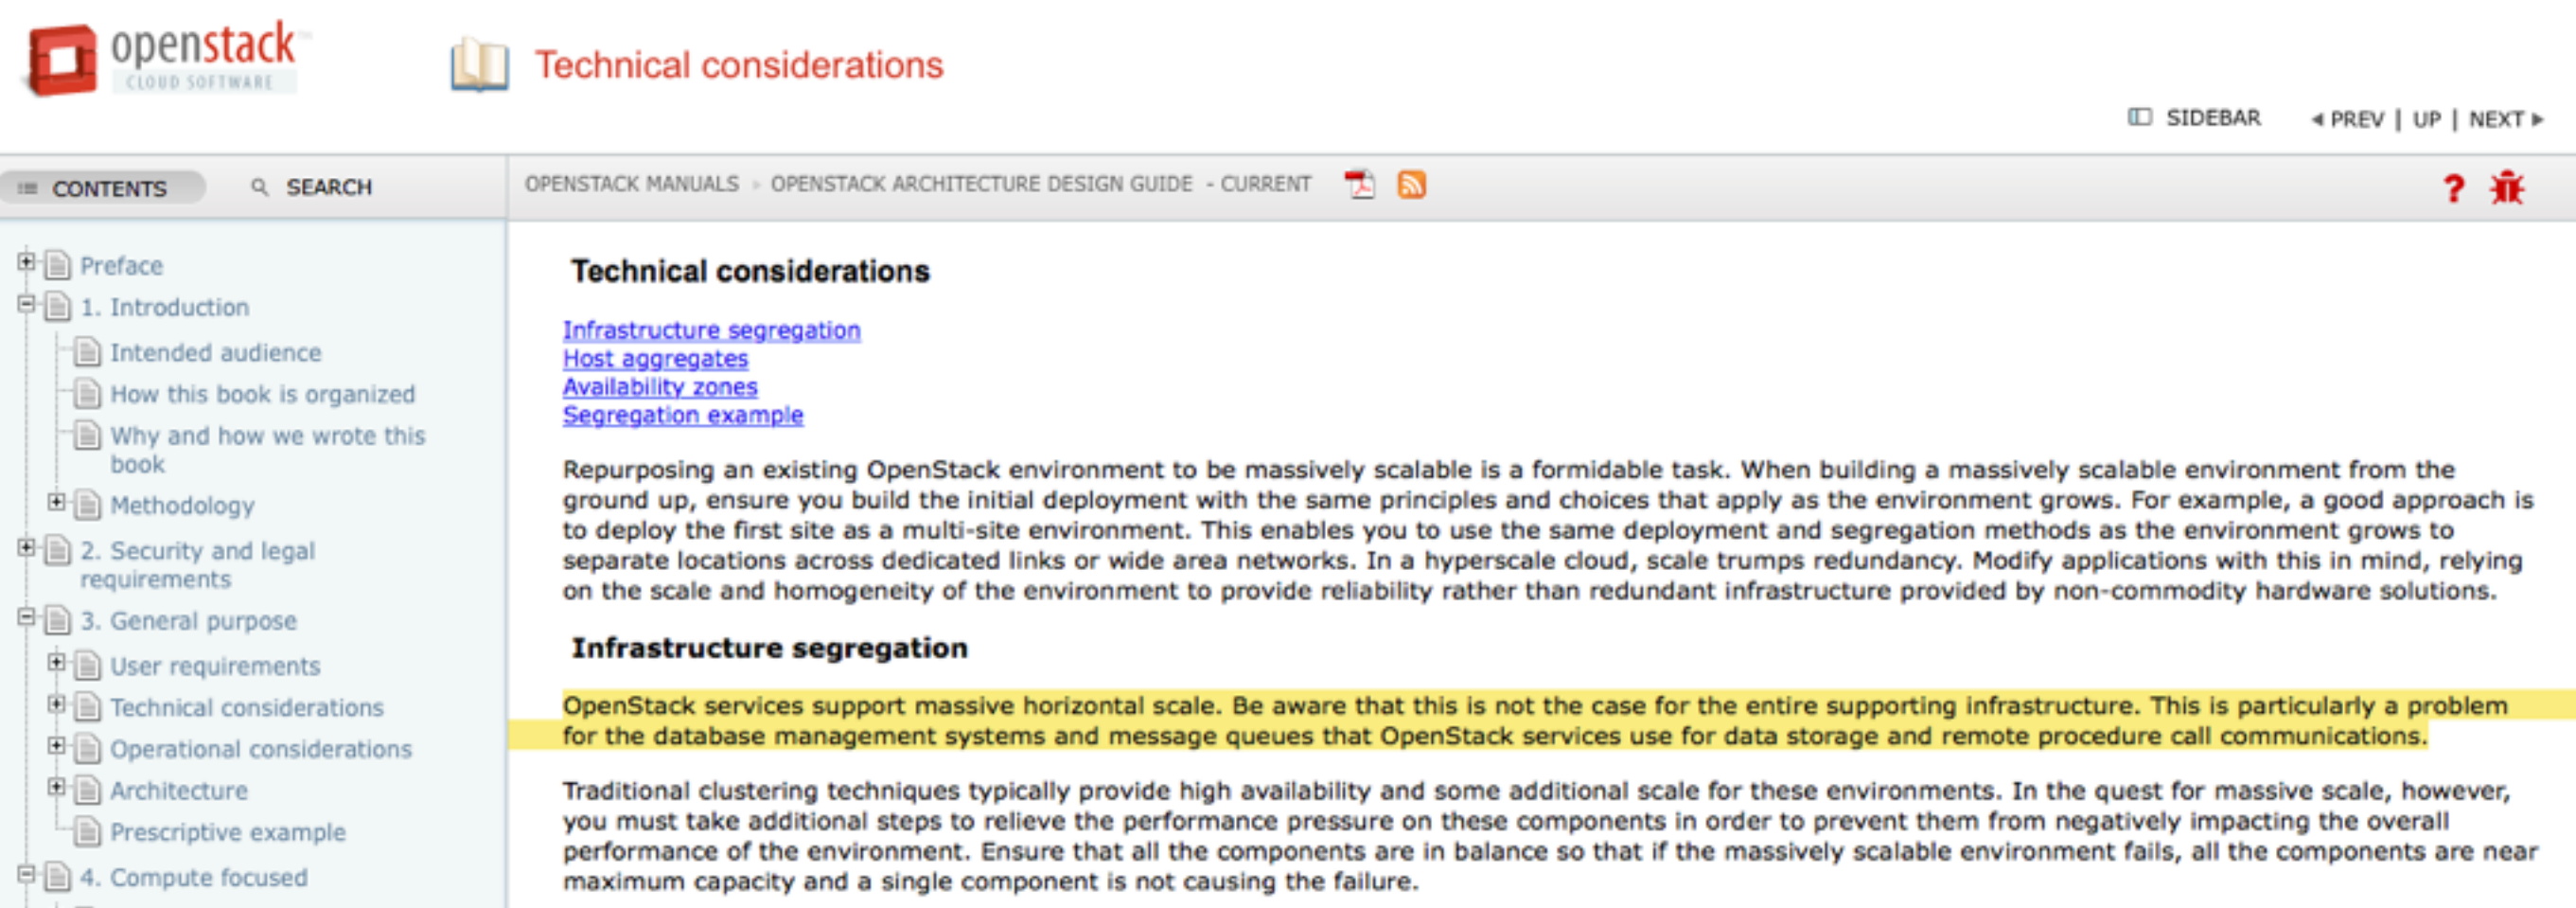
Bottom/Up - investigate whether/how OpenStack core services can be cooperative by default using Self\* and P2P mechanisms





# Looking back to the Future

- Austin Summit - May 2016 - Nova PoC



The screenshot displays the OpenStack Cloud Software website. The top navigation bar includes the OpenStack logo, a book icon, and the text 'Technical considerations'. On the right side of the top bar are links for 'SIDEBAR', 'PREV', 'UP', 'NEXT', a help icon, and a search icon. Below the top bar is a secondary navigation bar with 'CONTENTS', 'SEARCH', and a breadcrumb trail: 'OPENSTACK MANUALS > OPENSTACK ARCHITECTURE DESIGN GUIDE - CURRENT'. The left sidebar contains a table of contents with expandable sections: 'Preface', '1. Introduction' (with sub-items 'Intended audience', 'How this book is organized', 'Why and how we wrote this book'), 'Methodology', '2. Security and legal requirements', '3. General purpose' (with sub-items 'User requirements', 'Technical considerations', 'Operational considerations', 'Architecture', 'Prescriptive example'), and '4. Compute focused'. The main content area is titled 'Technical considerations' and contains several links: 'Infrastructure segregation', 'Host aggregates', 'Availability zones', and 'Segregation example'. The main text discusses repurposing an existing OpenStack environment for massive scalability, emphasizing the need for a multi-site environment and the importance of scale over redundancy. A highlighted section titled 'Infrastructure segregation' states that OpenStack services support massive horizontal scale, but the supporting infrastructure (database management systems and message queues) does not, which is a problem for data storage and remote procedure call communications. The text concludes by noting that traditional clustering techniques provide high availability and scale, but additional steps are needed to relieve performance pressure on these components to prevent them from negatively impacting the overall environment performance.

**Technical considerations**

[Infrastructure segregation](#)  
[Host aggregates](#)  
[Availability zones](#)  
[Segregation example](#)

Repurposing an existing OpenStack environment to be massively scalable is a formidable task. When building a massively scalable environment from the ground up, ensure you build the initial deployment with the same principles and choices that apply as the environment grows. For example, a good approach is to deploy the first site as a multi-site environment. This enables you to use the same deployment and segregation methods as the environment grows to separate locations across dedicated links or wide area networks. In a hyperscale cloud, scale trumps redundancy. Modify applications with this in mind, relying on the scale and homogeneity of the environment to provide reliability rather than redundant infrastructure provided by non-commodity hardware solutions.

**Infrastructure segregation**

OpenStack services support massive horizontal scale. Be aware that this is not the case for the entire supporting infrastructure. This is particularly a problem for the database management systems and message queues that OpenStack services use for data storage and remote procedure call communications.

Traditional clustering techniques typically provide high availability and some additional scale for these environments. In the quest for massive scale, however, you must take additional steps to relieve the performance pressure on these components in order to prevent them from negatively impacting the overall performance of the environment. Ensure that all the components are in balance so that if the massively scalable environment fails, all the components are near maximum capacity and a single component is not causing the failure.

# Looking back to the Future



## Technical considerations

SIDEBAR ◀ PREV | UP | NEXT ▶

OPENSTACK MANUALS > OPENSTACK ARCHITECTURE DESIGN GUIDE - CURRENT



### Technical considerations

[Infrastructure segregation](#)

[Host aggregates](#)

[Availability zones](#)

[Segregation example](#)

Repurposing an existing OpenStack environment to be massively scalable is a formidable task. When building a massively scalable environment from the ground up, ensure you build the initial deployment with the same principles and choices that apply as the environment grows. For example, a good approach is to deploy the first site as a multi-site environment. This enables you to use the same deployment and segregation methods as the environment grows to separate locations across dedicated links or wide area networks. In a hyperscale cloud, scale trumps redundancy. Modify applications with this in mind, relying on the scale and homogeneity of the environment to provide reliability rather than redundant infrastructure provided by non-commodity hardware solutions.

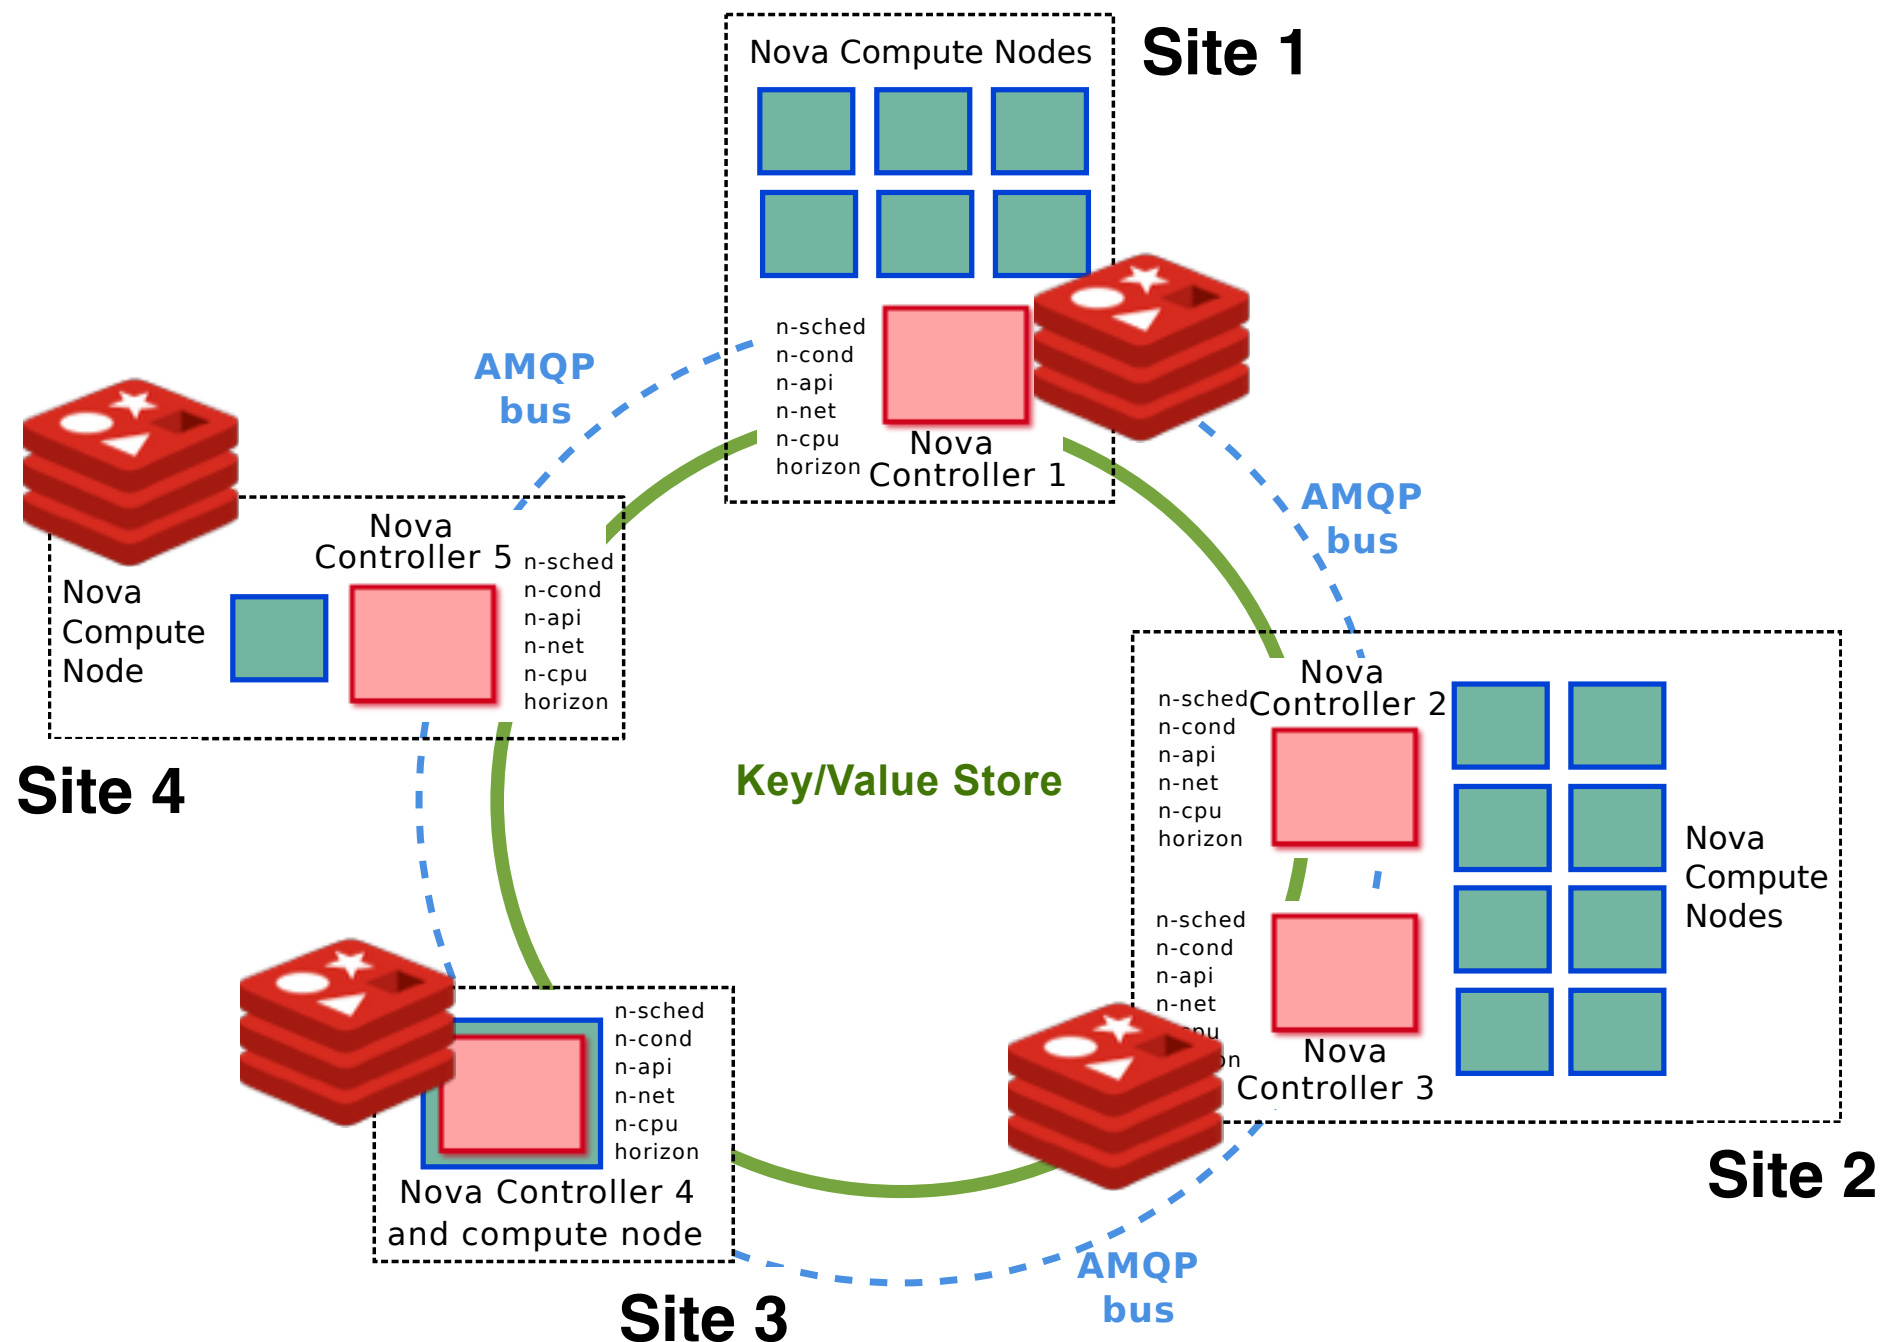
### Infrastructure segregation

OpenStack services support massive horizontal scale. Be aware that this is not the case for the entire supporting infrastructure. This is particularly a problem for the database management systems and message queues that OpenStack services use for data storage and remote procedure call communications.

Traditional clustering techniques typically provide high availability and some additional scale for these environments. In the quest for massive scale, however, you must take additional steps to relieve the performance pressure on these components in order to prevent them from negatively impacting the overall performance of the environment. Ensure that all the components are in balance so that if the massively scalable environment fails, all the components are near maximum capacity and a single component is not causing the failure.

# Looking back to the Future

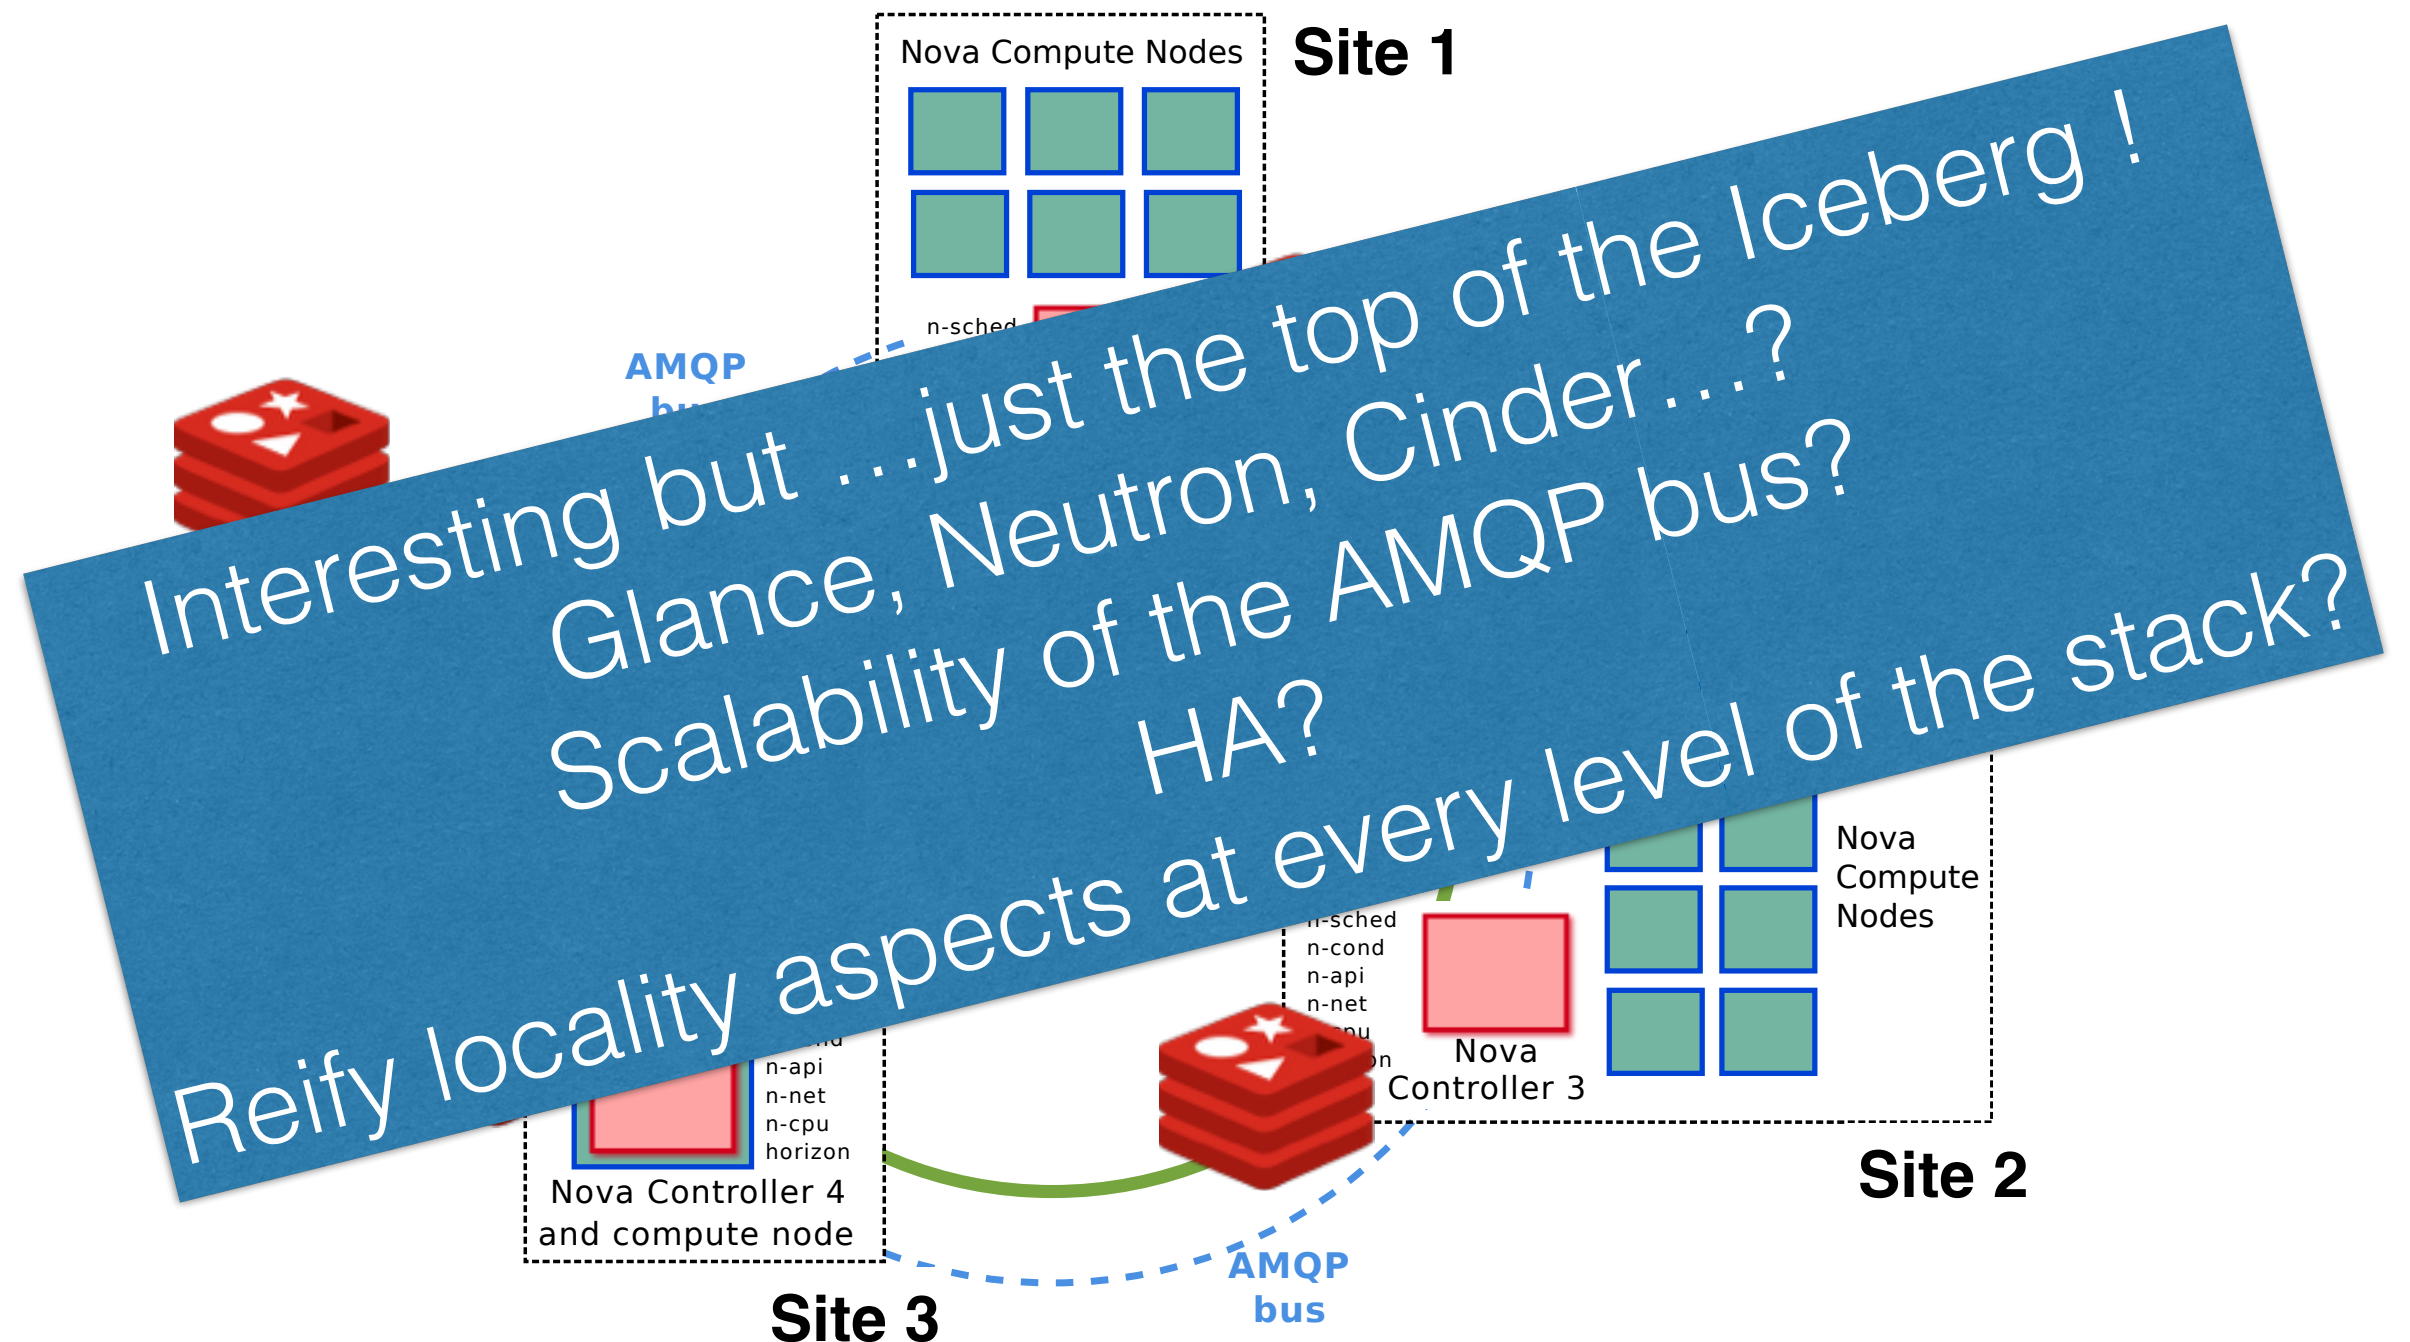
- Austin Summit - May 2016 - Nova PoC (based on Juno)  
Replaced MySQL DB by Redis (NoSQL backend)





# Looking back to the Future

- Austin Summit - May 2016 - Nova PoC (based on Juno)  
Replaced MySQL DB by Redis (NoSQL backend)



# Takeaway message

- Goal of the WG: do not reinvent the wheel (upstream first).

Study to what extent current mechanisms can handle Fog/Edge infrastructures

Propose revisions/extensions of internal mechanisms when appropriate.

Investigate how should current cloud APIs be extended to take the advantage of the geo-distribution (latency-aware applications...)

- Ongoing action: Analyze OpenStack Performance under the Fog/Edge perspective (scalability, traffic characterisation...) using EnOS (a dedicated framework for conducting performance evaluations of OpenStack)



**Toward Fog, Edge, and NFV Deployments:  
Evaluating OpenStack WANwide**

**Wednesday 2:40pm-3:20pm, Hynes CC, Level 3, Ballroom A**

- Face-to-face meeting (for current and new WG members)  
Wednesday 5:20pm - 6:00pm, Hynes CC, Level 2, MR 201



# Fog Computing / Edge Computing Massively Distributed Clouds Working Group

Adrien Lebre


BoF Session - OpenStack Summit - Boston

<https://etherpad.openstack.org/p/BOS-Fog-Edge-MassivelyDistributed-BoF>  
(or just go on google and looks for "Fog Edge OpenStack")



# Ongoing Action

- Collaboration with the Performance Team to understand OpenStack Performance (scalability, traffic characterisation...)
- EnOS: Experimental Environment for Conducting performance evaluations of OpenStack



# Ongoing Action

- Ongoing collaboration with the Performance Team to understand OpenStack Performance (scalability, traffic characterisation...)
- EnOS: Experimental Environment for Conducting performance evaluations of OpenStack
- Current focus: placement constraints/opportunities  
how many instances of each service? one global bus? one central Glance? several? Where should we locate them?...

