

# DEPLOYMENT AND RECONFIGURATION CHALLENGES

Hélène Coullon, Christian Perez, Dimitri Pertin

Ascola & Avalon

IPL Discovery




# OUTLINE

1. Motivation
2. Survey
3. Initial Deployment
4. Reconfiguration
5. Perspectives
6. IPL and impact

# MOTIVATION

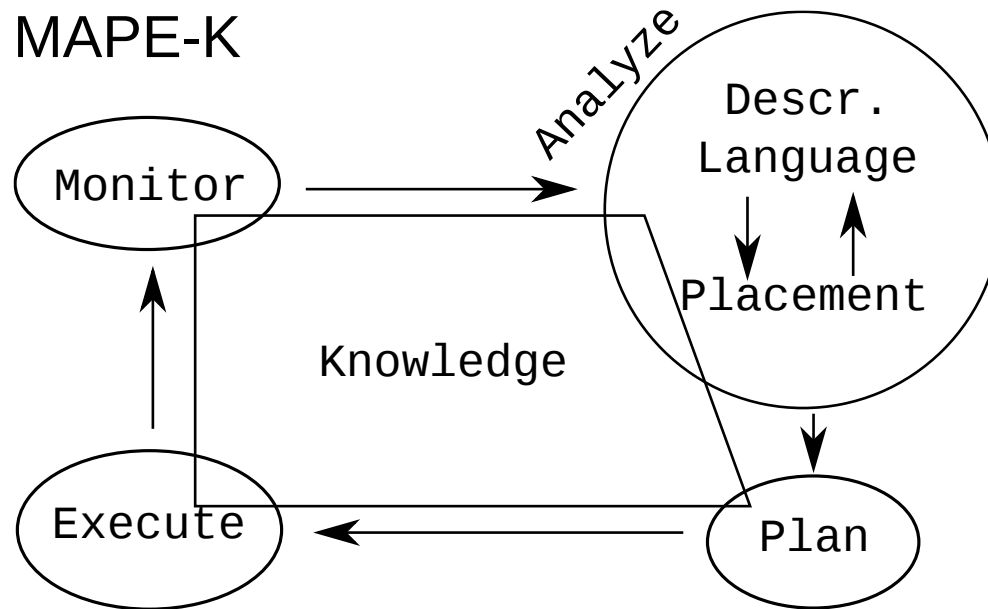
# CONTEXT

## DEPLOYMENT AUTOMATION



# DEPLOYMENT / RECONFIGURATION

## MAPE-K

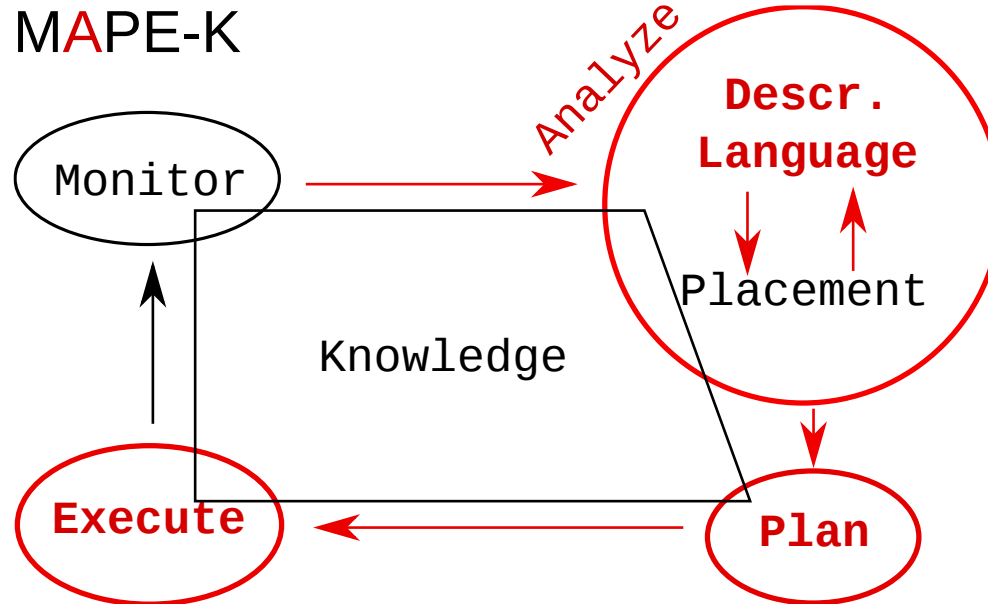


# ANALYZE

- Descr. Language
  - *What* to deploy?
  - *How* to deploy?
- Placement
  - *Where* to deploy?
  - Infrastructure/resource description

# DEPLOYMENT / RECONFIGURATION

## FOCUS



# BIG PICTURE

How to deploy/re-deploy systems and applications on infrastructures?

Expected properties:

- low-level flexible generic model
- appropriate level of expressivity
- dynamicity (reconfiguration)
- correctness and attainability
- performances and scalability



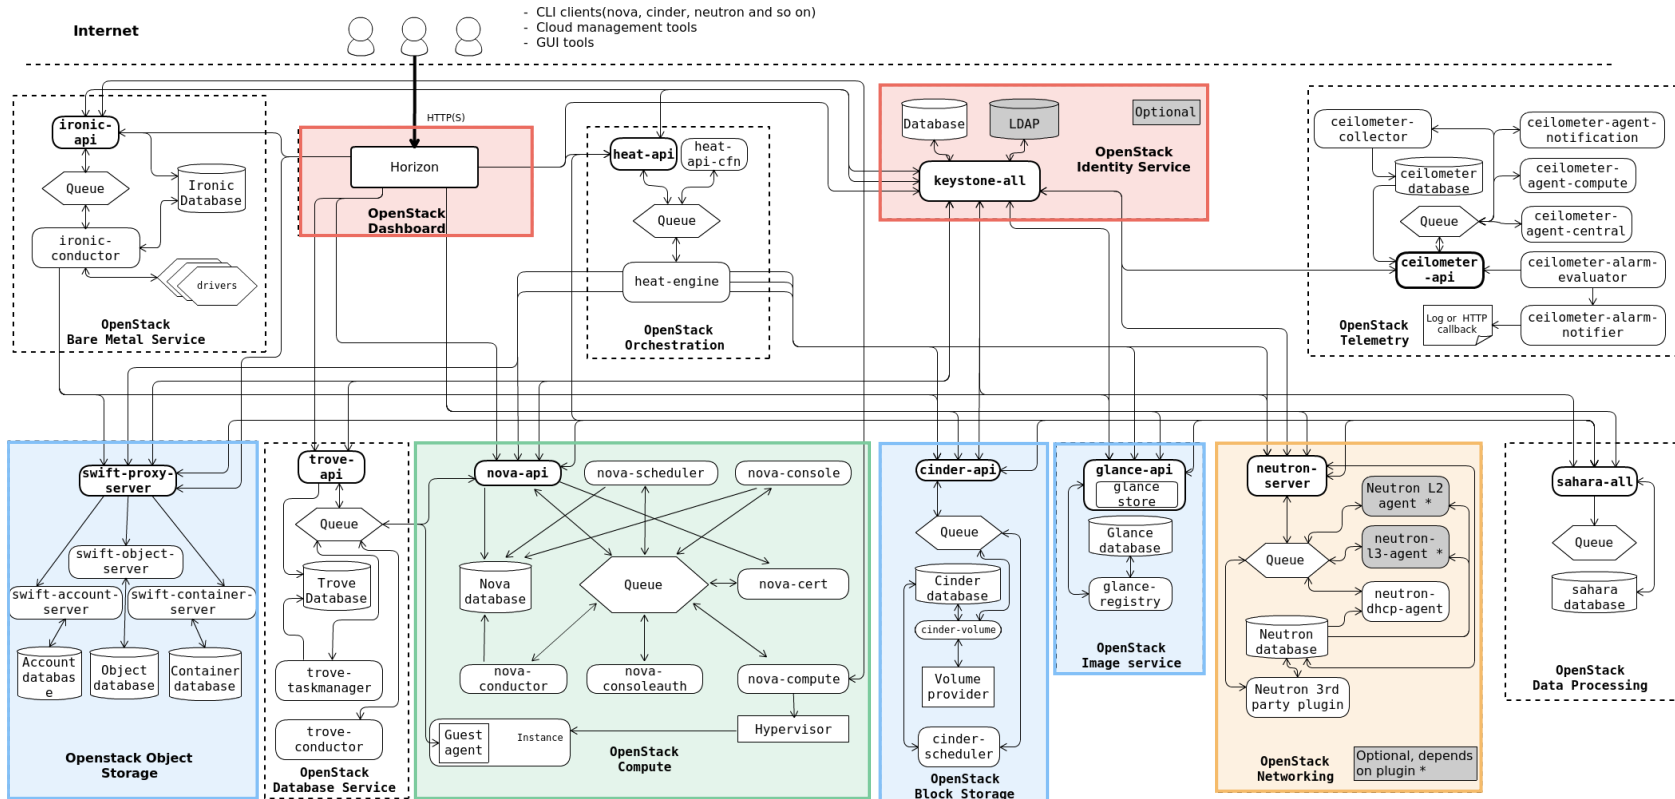
# TWO PHASES

1. Research on the initial deployment problem
  - [Dimitri Pertin](#), postdoc researcher
2. Research on the reconfiguration problem
  - [Maverick Chardet](#), PhD student (October 2017)

# USE CASES

*Any application or system*

- OpenStack and its decentralized version



# USE CASES

- Smart-\* applications composed of hybrid services
  - BigData
  - HPC
  - Stream processing
  - Virtual reality
  - etc.

**SURVEY**

# SURVEY

Hélène Coullon, Christian Perez, Dimitri Pertin.

*Production Deployment Tools for IaaS: an Overall Model  
and Survey - FiCloud 2017*

# PRODUCTION TOOLS

1. **Kolla**: *deploy production-ready OpenStack instances by leveraging Ansible and Docker*
2. **Juju**: *Canonical project to write your application life cycle and deploy it on major cloud providers*
3. **Kubernetes**: *a project designed by Google to deploy and maintain containerized applications*
4. **TripleO** (OpenStack On OpenStack): *an OpenStack project aiming at deploying OpenStack instances using OpenStack's own services*

# PRODUCTION TOOLS

## GENERICITY

	Kolla	Juju	K8s	TripleO
app. generic	No	Yes	Yes	No
env. generic	No	Yes	No	No

# PRODUCTION TOOLS

## PLACEMENT ET MONITOR

	Kolla	Juju	K8s	TripleO
Plac.	Ext.	Ext.	Int.	Int.
	Manual	Auto	Auto	Auto
Mntr.	No	Manual	1/2 Auto	1/2 Auto



# PRODUCTION TOOLS

## DESCRIPTION LANGUAGE

	Kolla	Juju	K8s	TripleO
relations	No	Yes	Yes	Yes
constraints	No	Yes	Yes	Yes
actions	No	Yes	Yes	No
reconfiguration	No	3.	1.	2.

# PRODUCTION TOOLS

## CONCLUSION

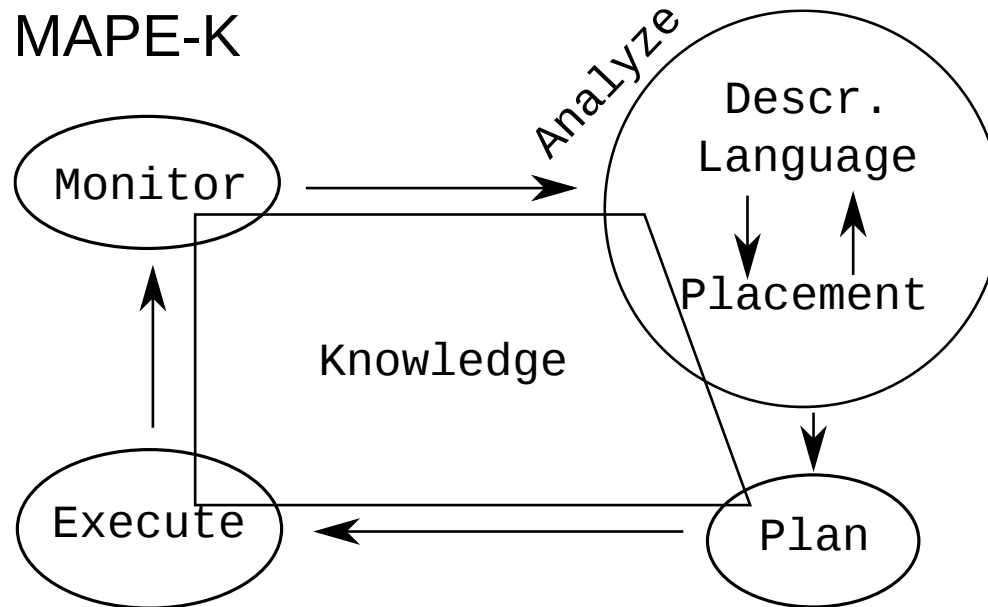
Expected properties:

- low-level flexible generic model
- appropriate level of expressivity
- dynamicity (reconfiguration)
- correctness and attainability
- performances and scalability

# INITIAL DEPLOYMENT

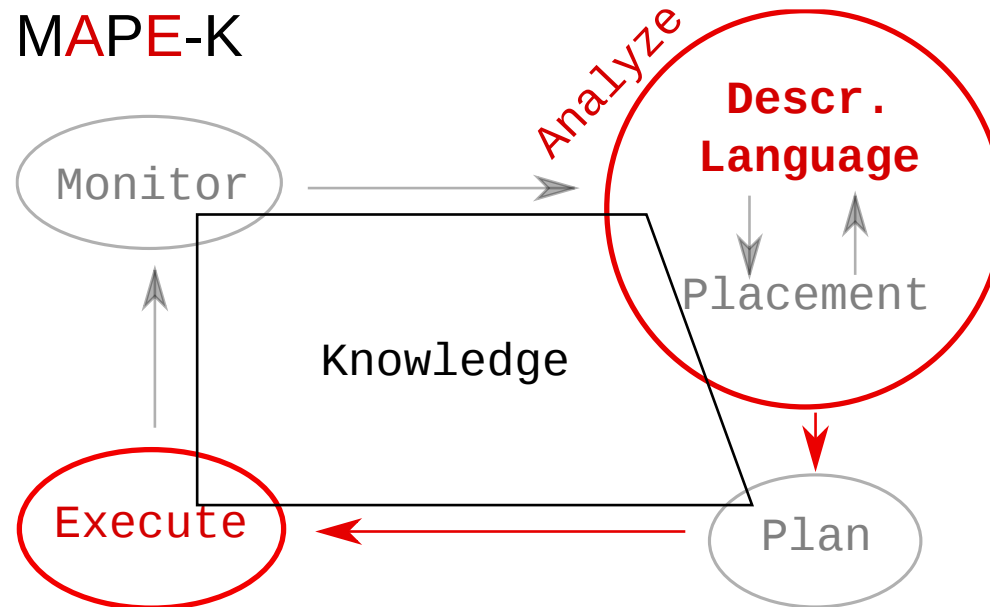
# INITIAL DEPLOYMENT

MAPE-K



# INITIAL DEPLOYMENT

MAPE-K



# INITIAL DEPLOYMENT

## GENERIC

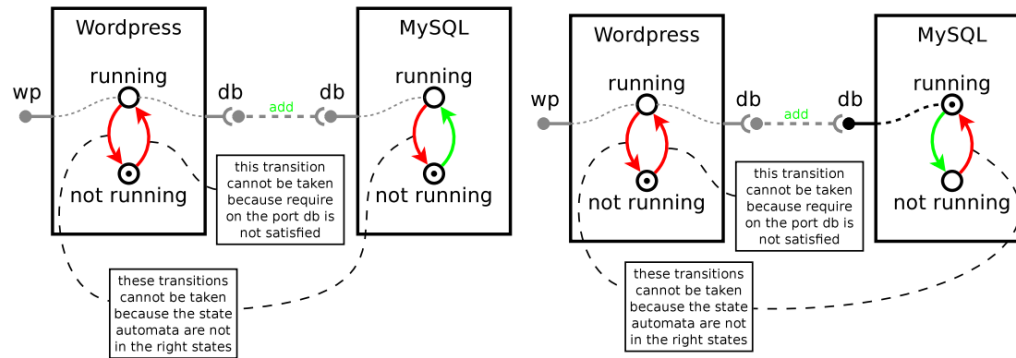
- Production tools: Juju, K8s, ansible
- Academic:
  - *Architecture Description Languages*: Tosca,
  - *Component models*: CCM, SCA, Fractal/Pro-active, L2C/HLCM etc.
  - *Deployment*: Deployware, PaaSage, MuScADeL, **Aeolus** etc.

# INITIAL DEPLOYMENT

## AEOLUS

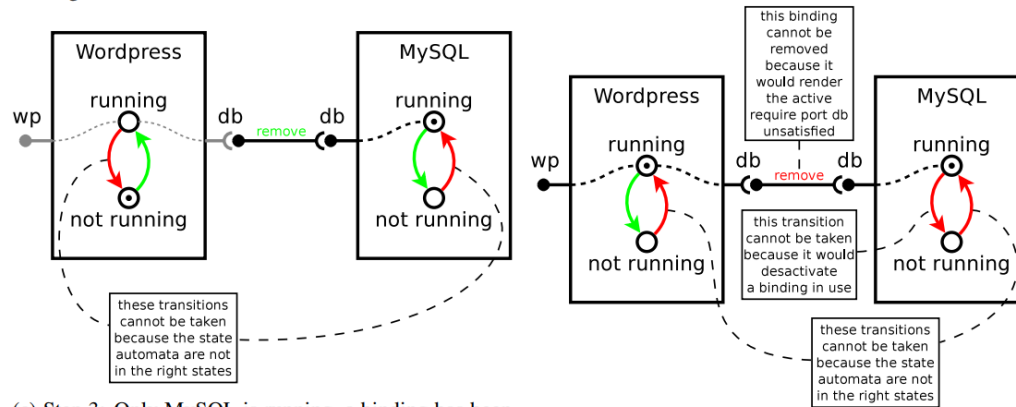
Di Cosmo et al. 2014, *Information and Computation*

- components = services/modules
- functional dependencies = use-provide
- temporal constraints = states and transitions



(a) Step 1: both components are not running, there is no binding between them.

(b) Step 2: Only MySQL is running, there is still no binding.



(c) Step 3: Only MySQL is running, a binding has been established.

(d) Step 4: Both components are running, the binding is present.

	provide port	require port	activation relation
inactive			-----
active			-----




# CONTRIBUTIONS

## MULTI-TOKEN PETRI NETS, DATAFLOW


- Multi-token
  - correct internal dependencies
  - expose parallelism
- Petri-net
  - clear semantic between transitions and states
- Dataflow

# CONTRIBUTIONS

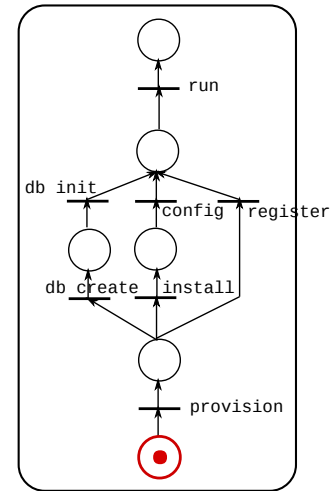
## EXAMPLE



MariaDB



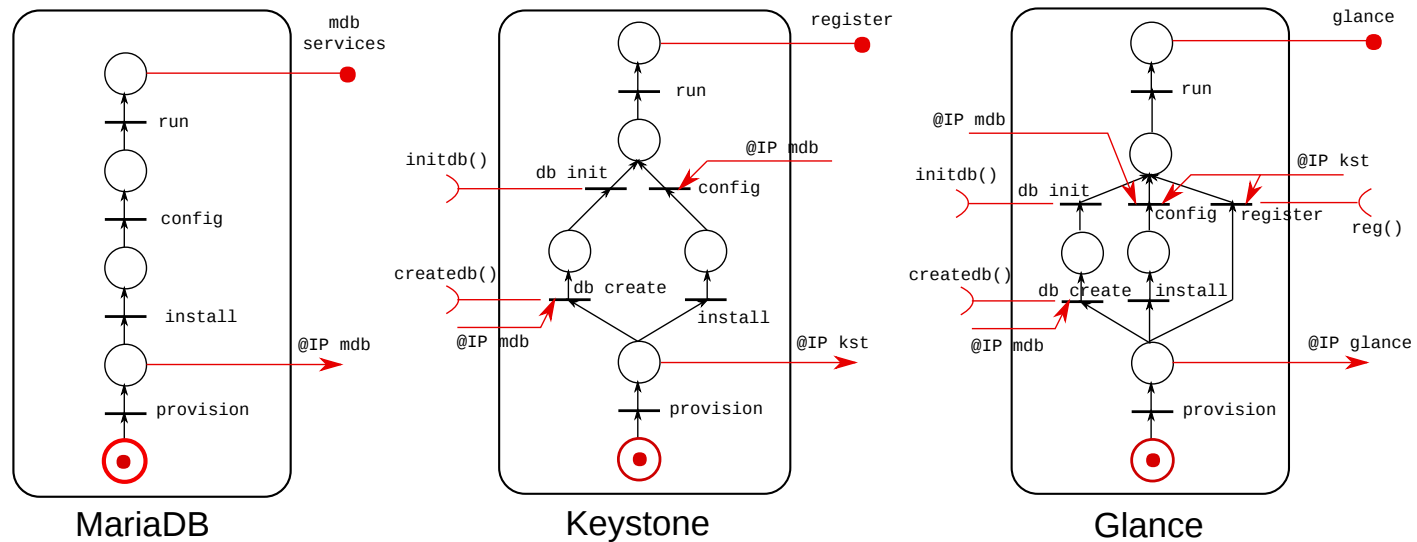
Keystone

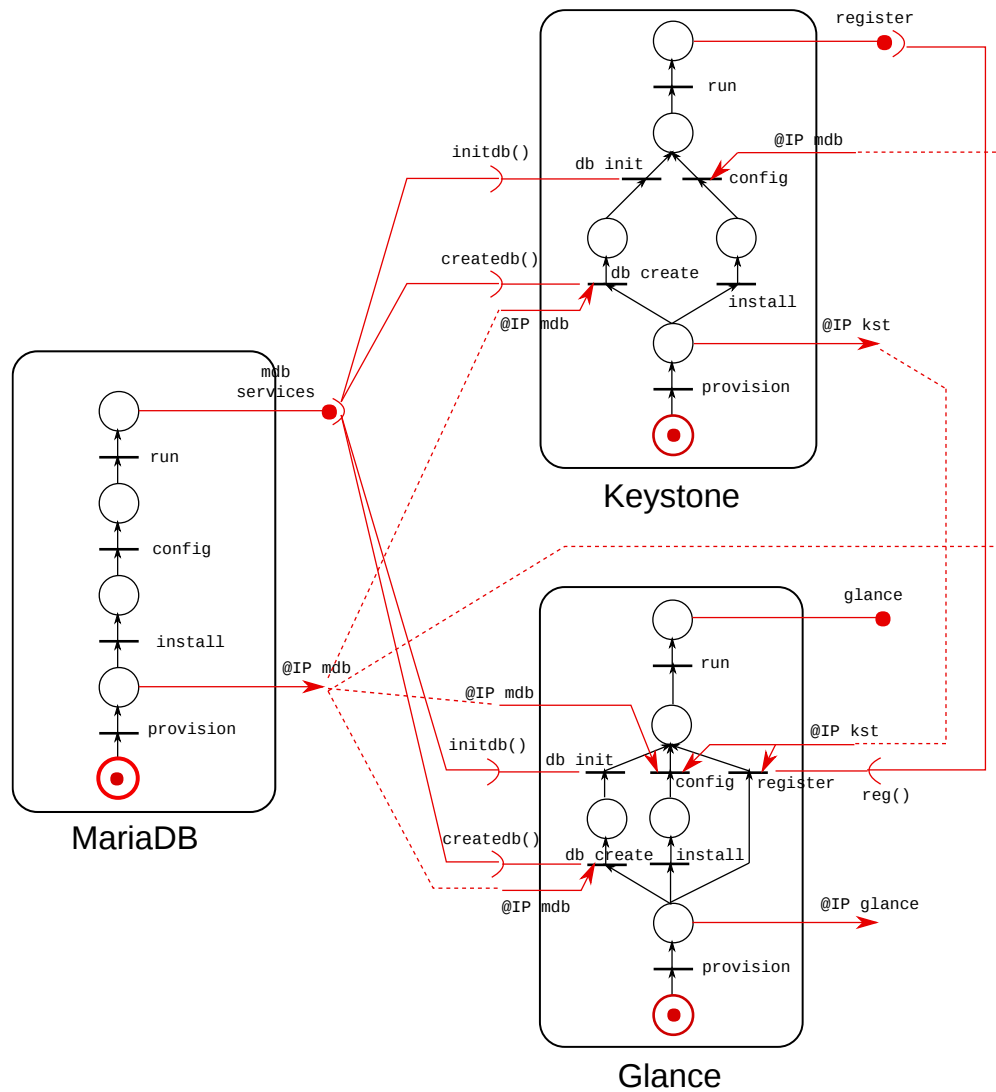


Glance

# CONTRIBUTIONS

## EXAMPLE






# OTHER CONTRIBUTIONS


- Conditions (Workflow?)
- Colored petri-nets
- Hierarchy and cardinality

# CONTRIBUTIONS

## COLORS



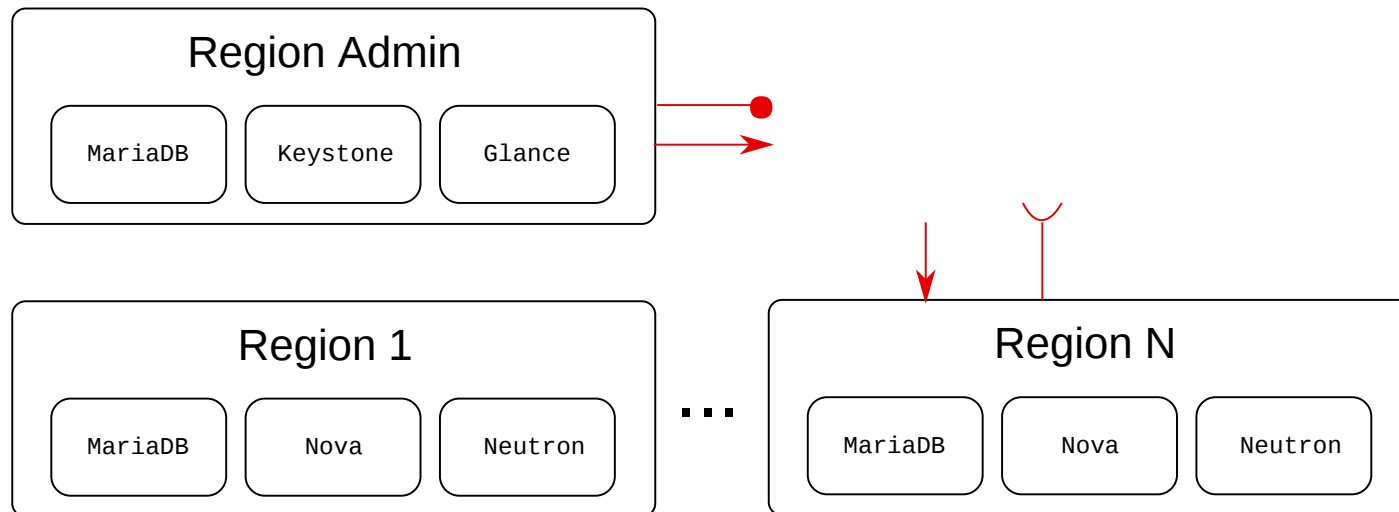
Keystone



Keystone

# CONTRIBUTIONS

## HIERARCHY

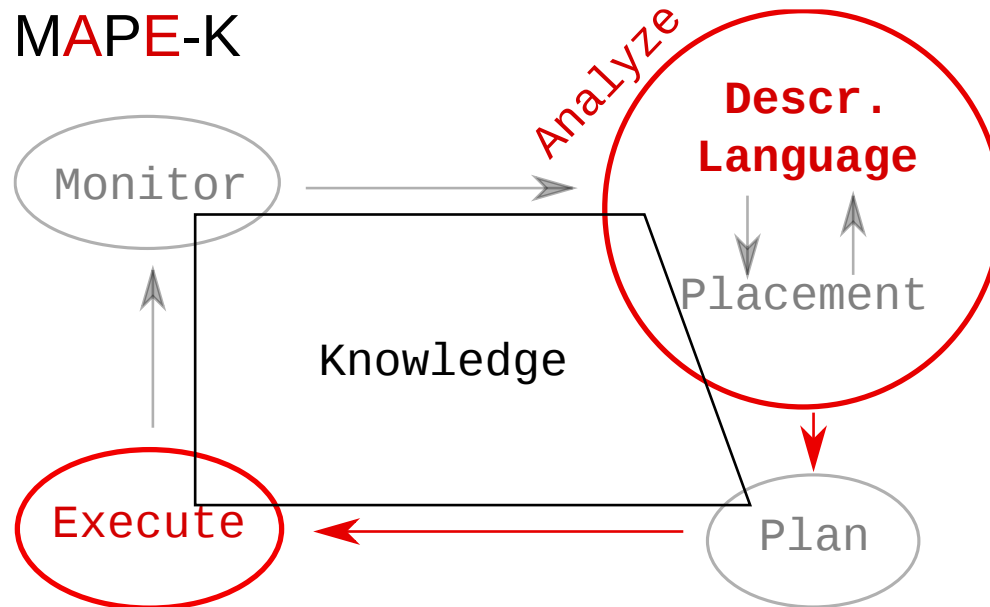


# RECONFIGURATION



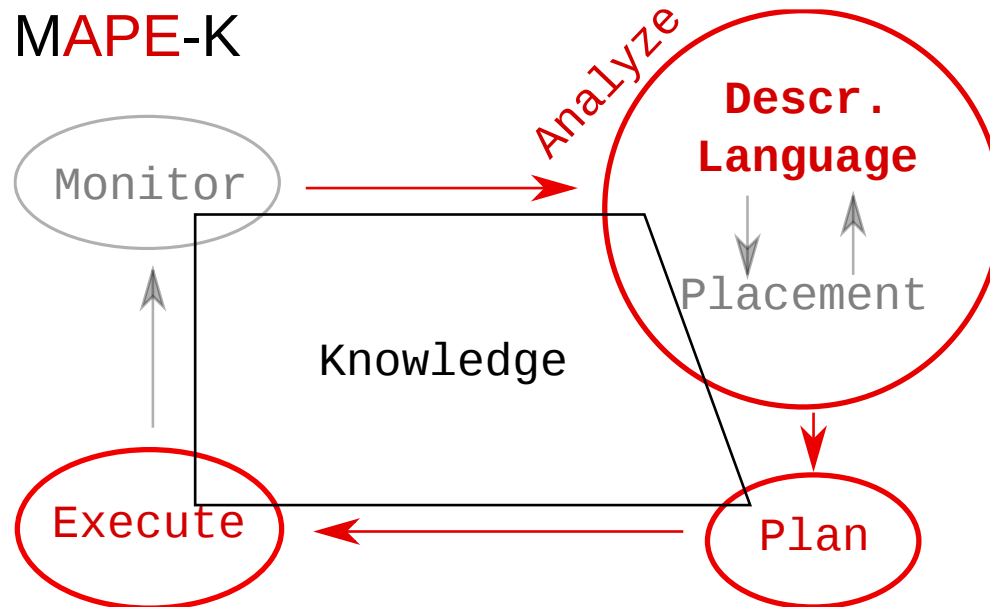
# RECONFIGURATION

MAPE-K



# RECONFIGURATION

MAPE-K



# RECONFIGURATION

## APPLICATION SIDE

- Fault Tolerance
- Automatic scaling (scale out, scale up)
- Updates (maintainability)
- Automatic software changes for various external reasons

# RECONFIGURATION

## INFRASTRUCTURE SIDE

- Massively distributed
- Failures
- Enter/leave
- Heterogeneity

# RECONFIGURATION

## CHALLENGES

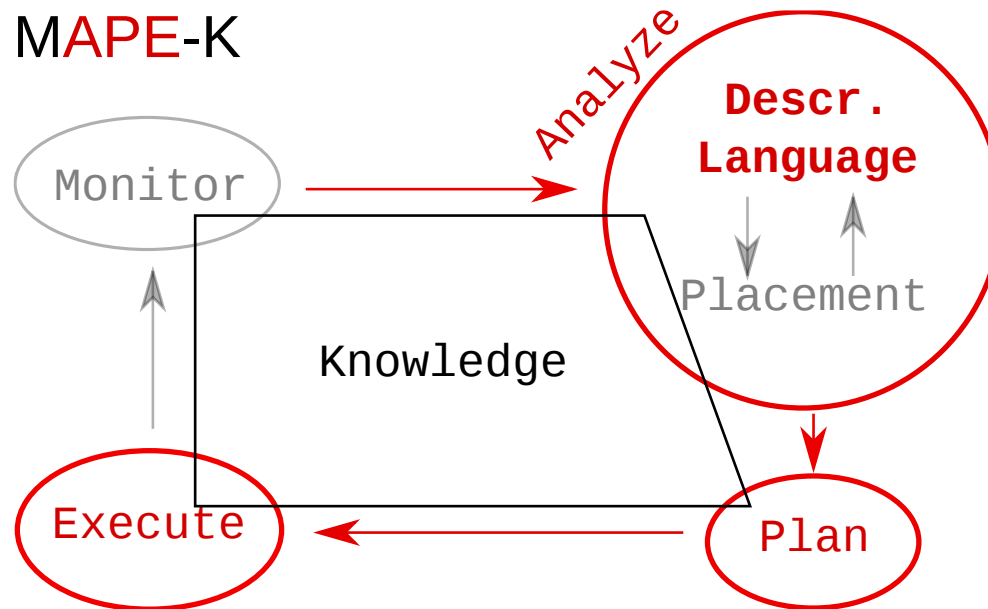
- Reconfiguration expressivity
- How to perform the reconfiguration? (Plan)
- Performances of the reconfiguration

# PERSPECTIVES

# PERSPECTIVES

## RESEARCH

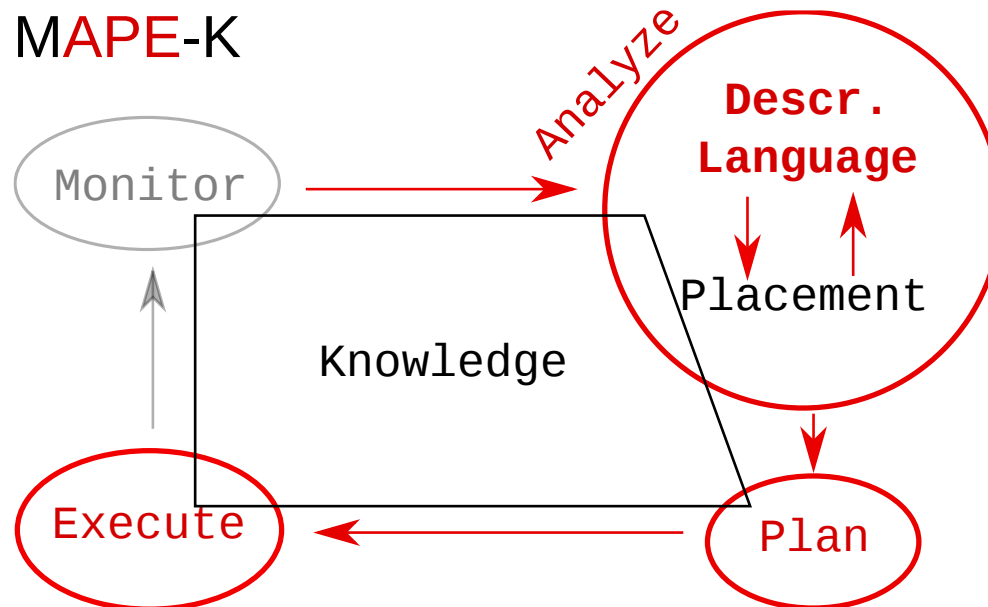
MAPE-K



# PERSPECTIVES

## RESEARCH

MAPE-K





# PERSPECTIVES

## RESEARCH

- Formalism and proofs
- Higher abstraction level models:
  - Generic model as a back-end
  - Being more specific to applications and systems
  - DSLs
  - Compilation

# PERSPECTIVES

## COLLABORATIONS

- Roberto Di Cosmo: Formalism and proofs
- Marcos Dias: User-centric use-cases (Avalon)
- Mario Studholt: Security policies (Ascola)
- Erik Elmroth and Elastisys

# **IPL ROLE AND IMPACT**

# IPL ROLE

- EPC [Avalon](#), Lyon: Component models
- EPC [Ascola](#), Nantes: Distributed Systems, OpenStack
- Use-case: [Discovery](#) ! (decentralized OpenStack)

# IMPACT

- To deploy OpenStack and automatically manage its behavior
- Within OpenStack to improve application deployment on VMs

# THANK YOU !

