

Koala protocol

@ Discovery Midterm Review

Genc Tato

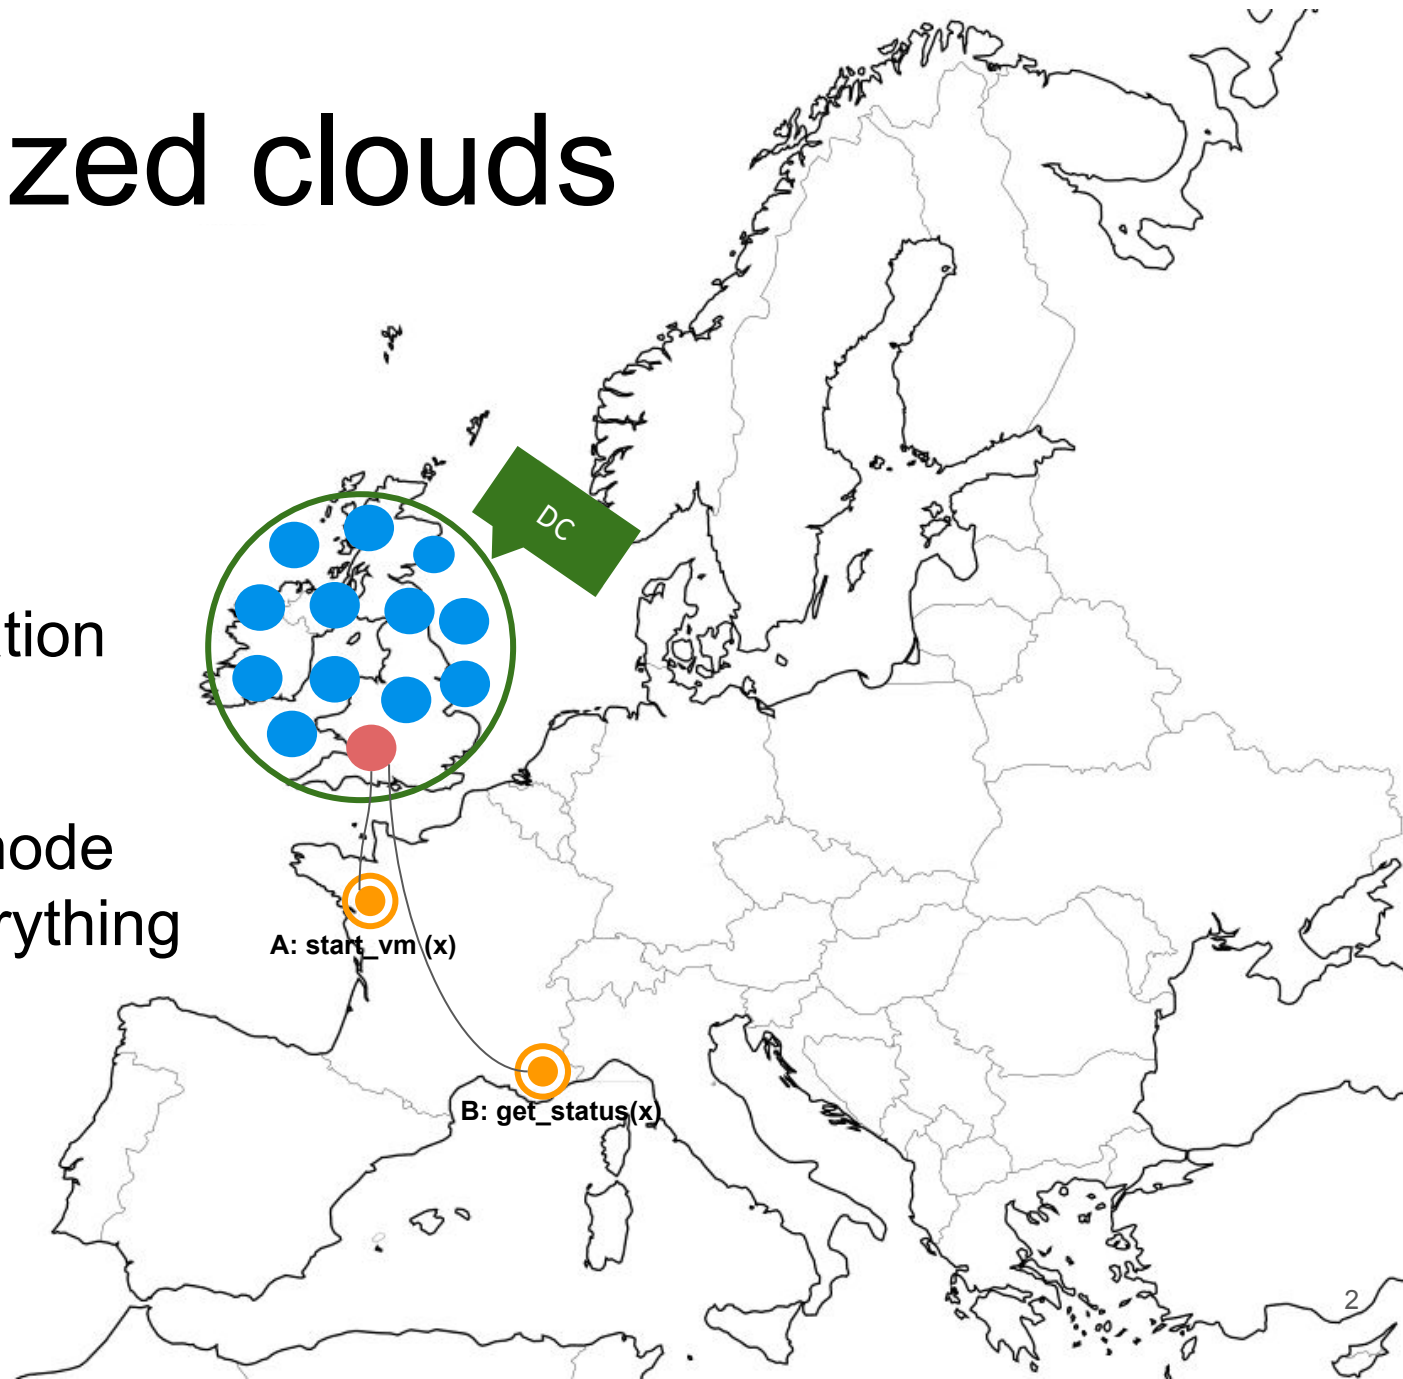
Supervised by:

- Cédric Tedeschi
- Marin Bertier

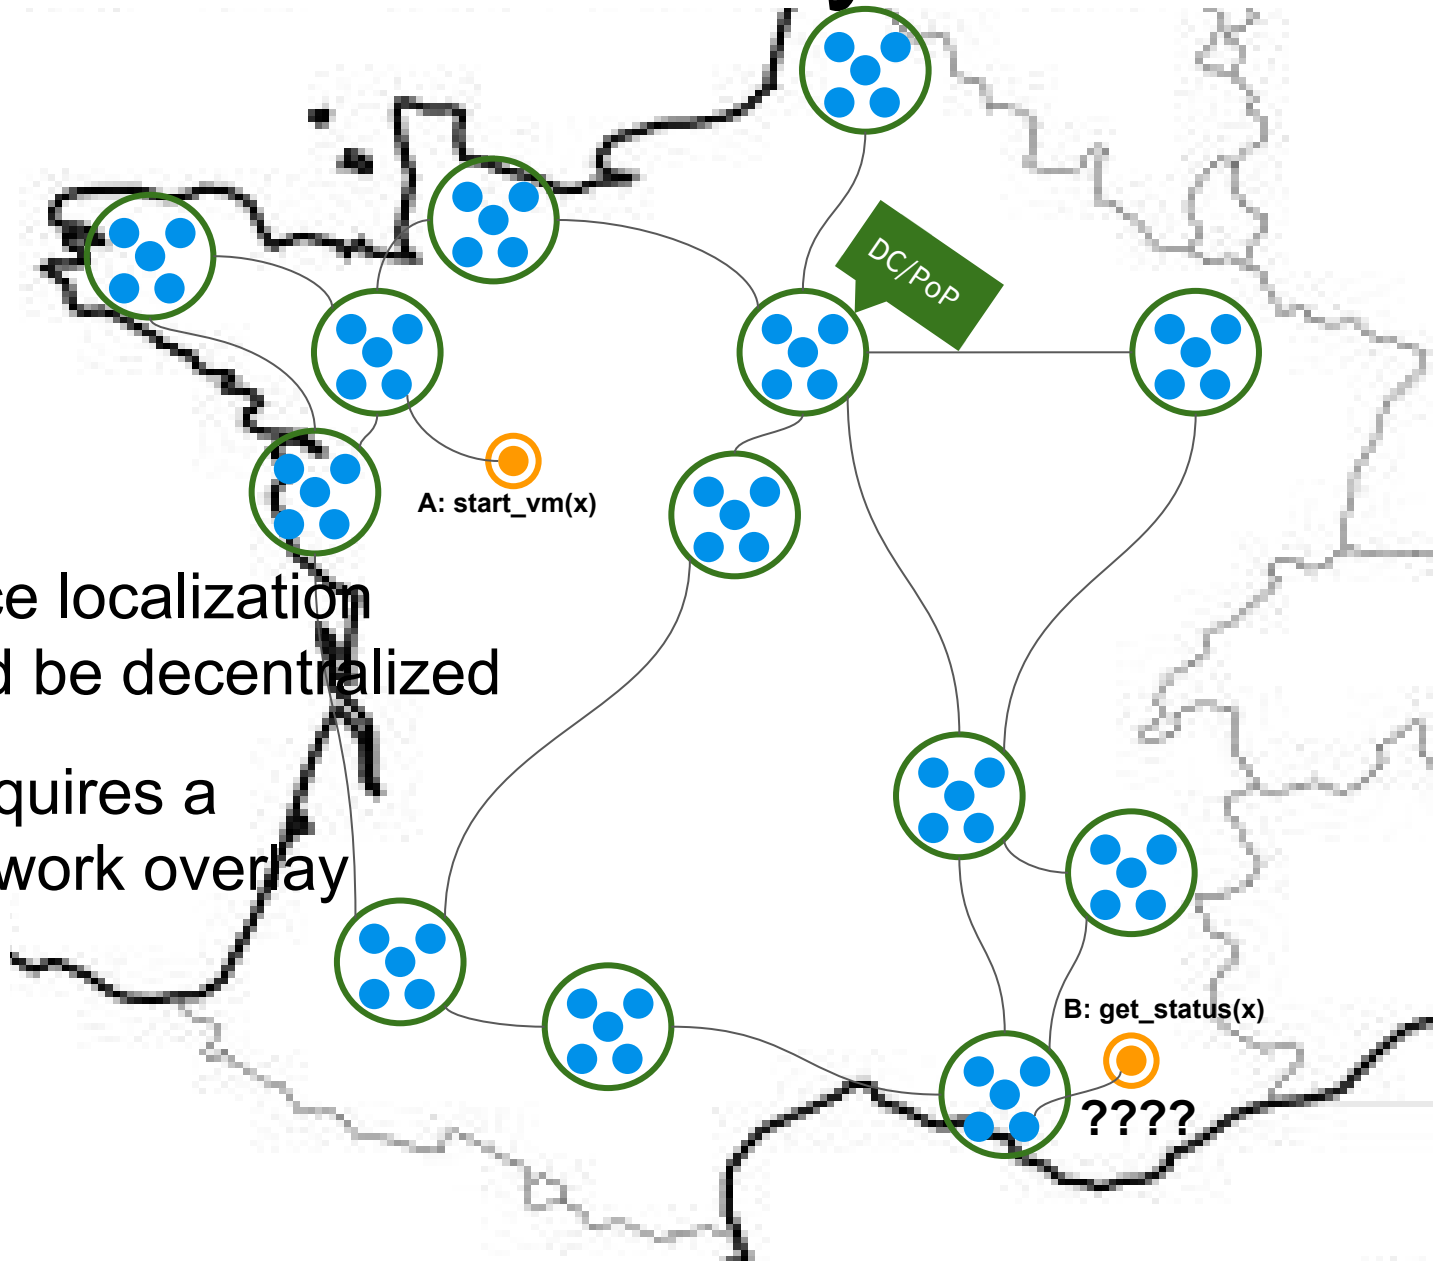
Centralized clouds

Service localization
is centralized

- Controller node
knows everything



Discovery Cloud



Service localization
should be decentralized

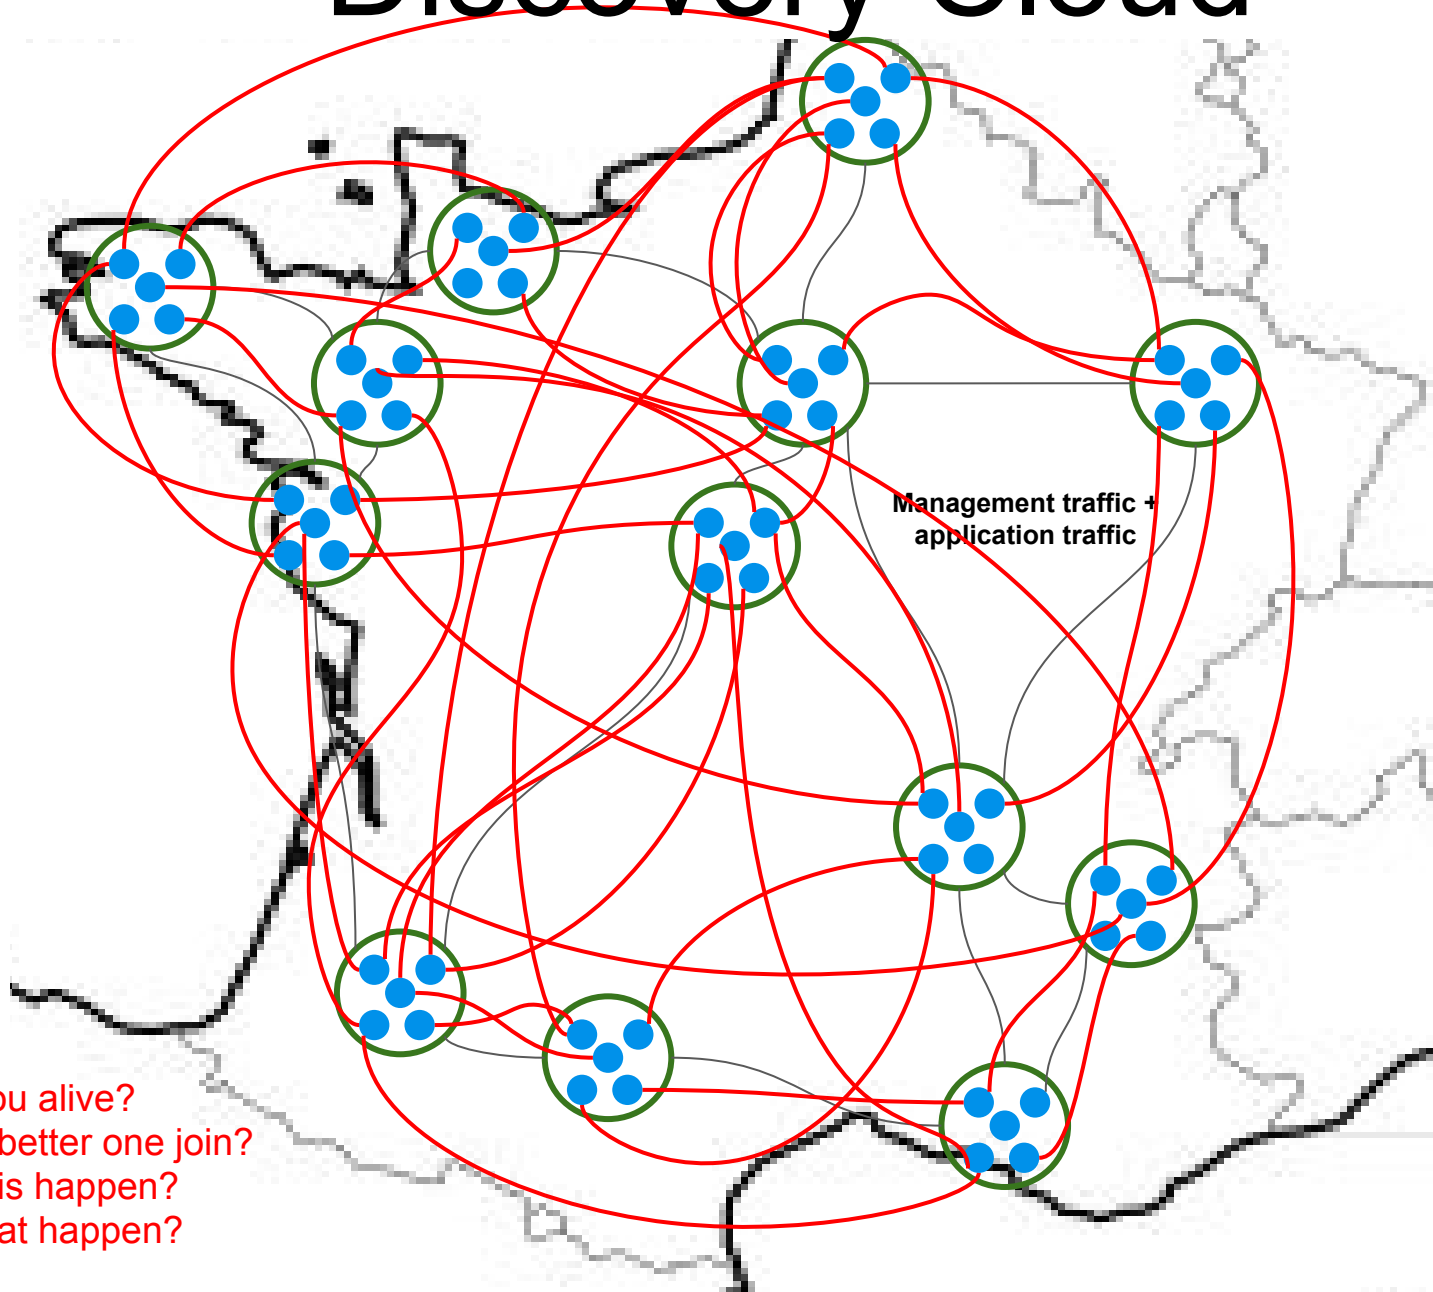
- Requires a
network overlay

Koala: An overlay for decentralized clouds

...but how should this overlay be?



Discovery Cloud



Are you alive?
Did a better one join?
Did this happen?
Did that happen?

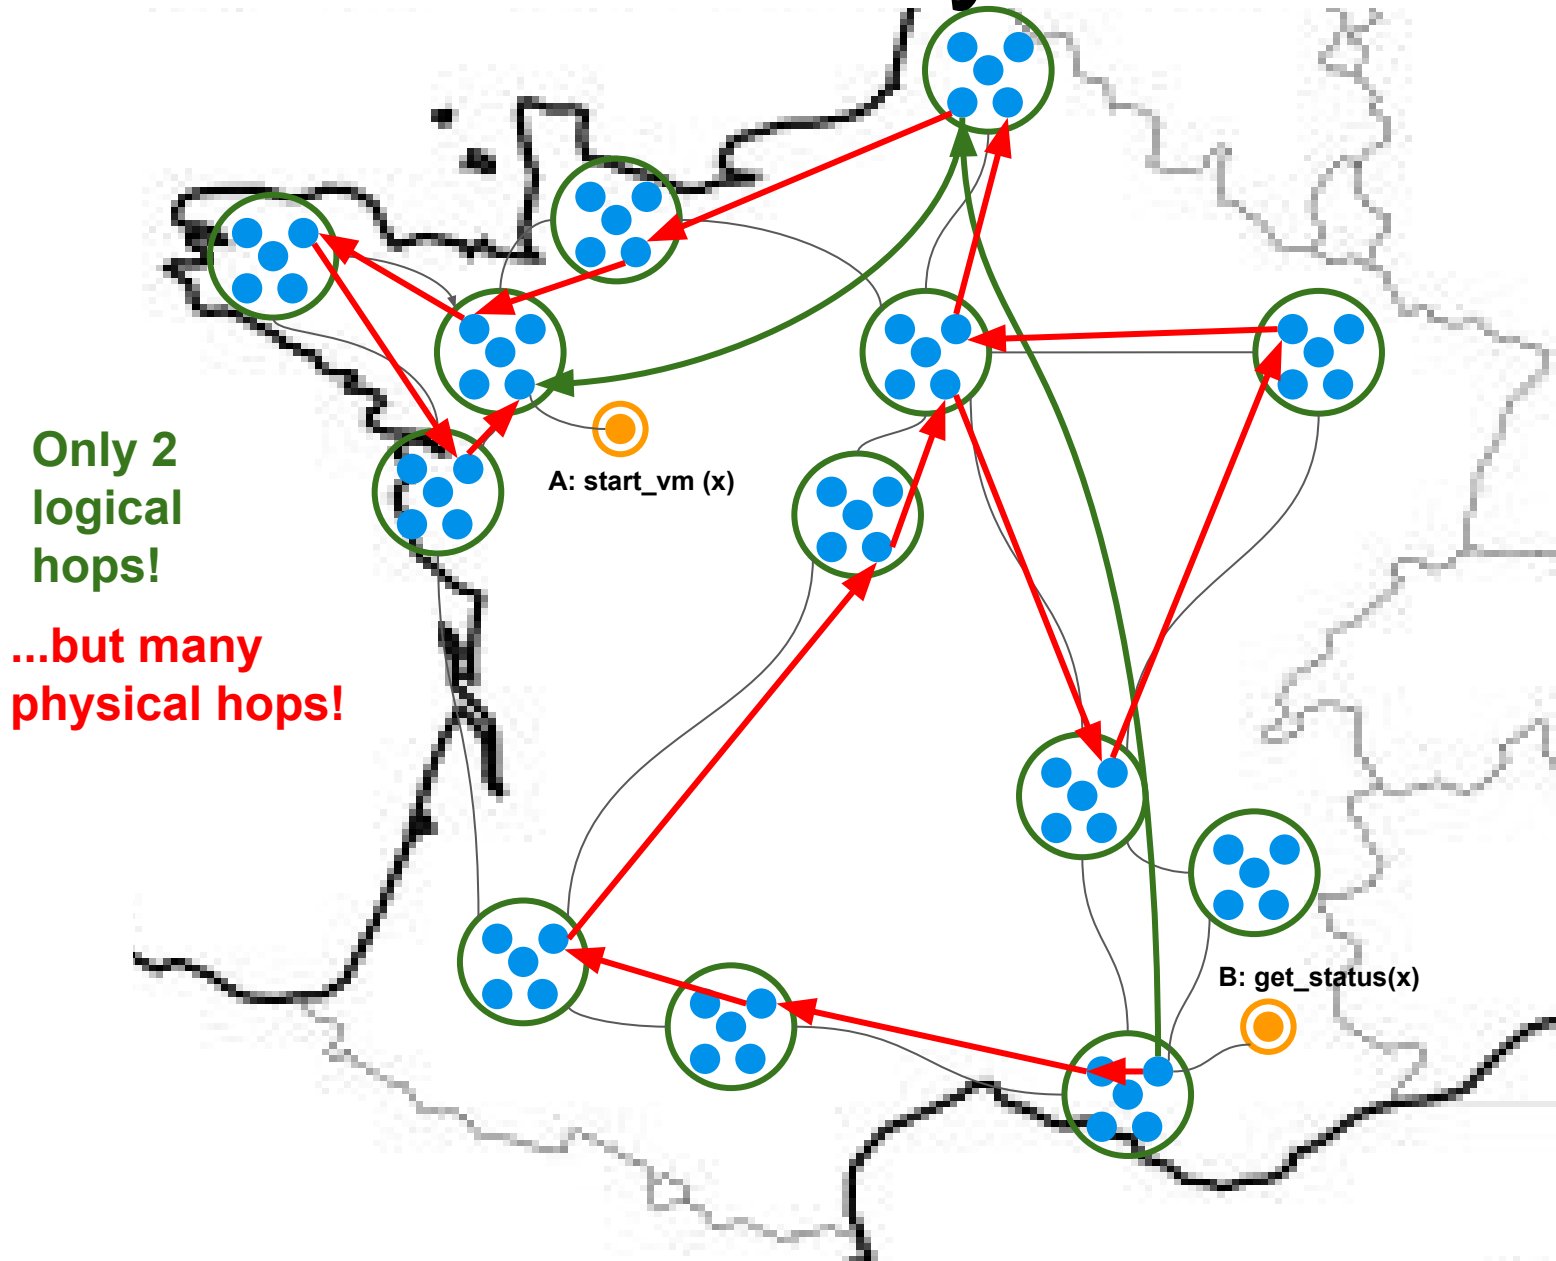
Koala: An overlay for decentralized clouds

...but how should this overlay be?

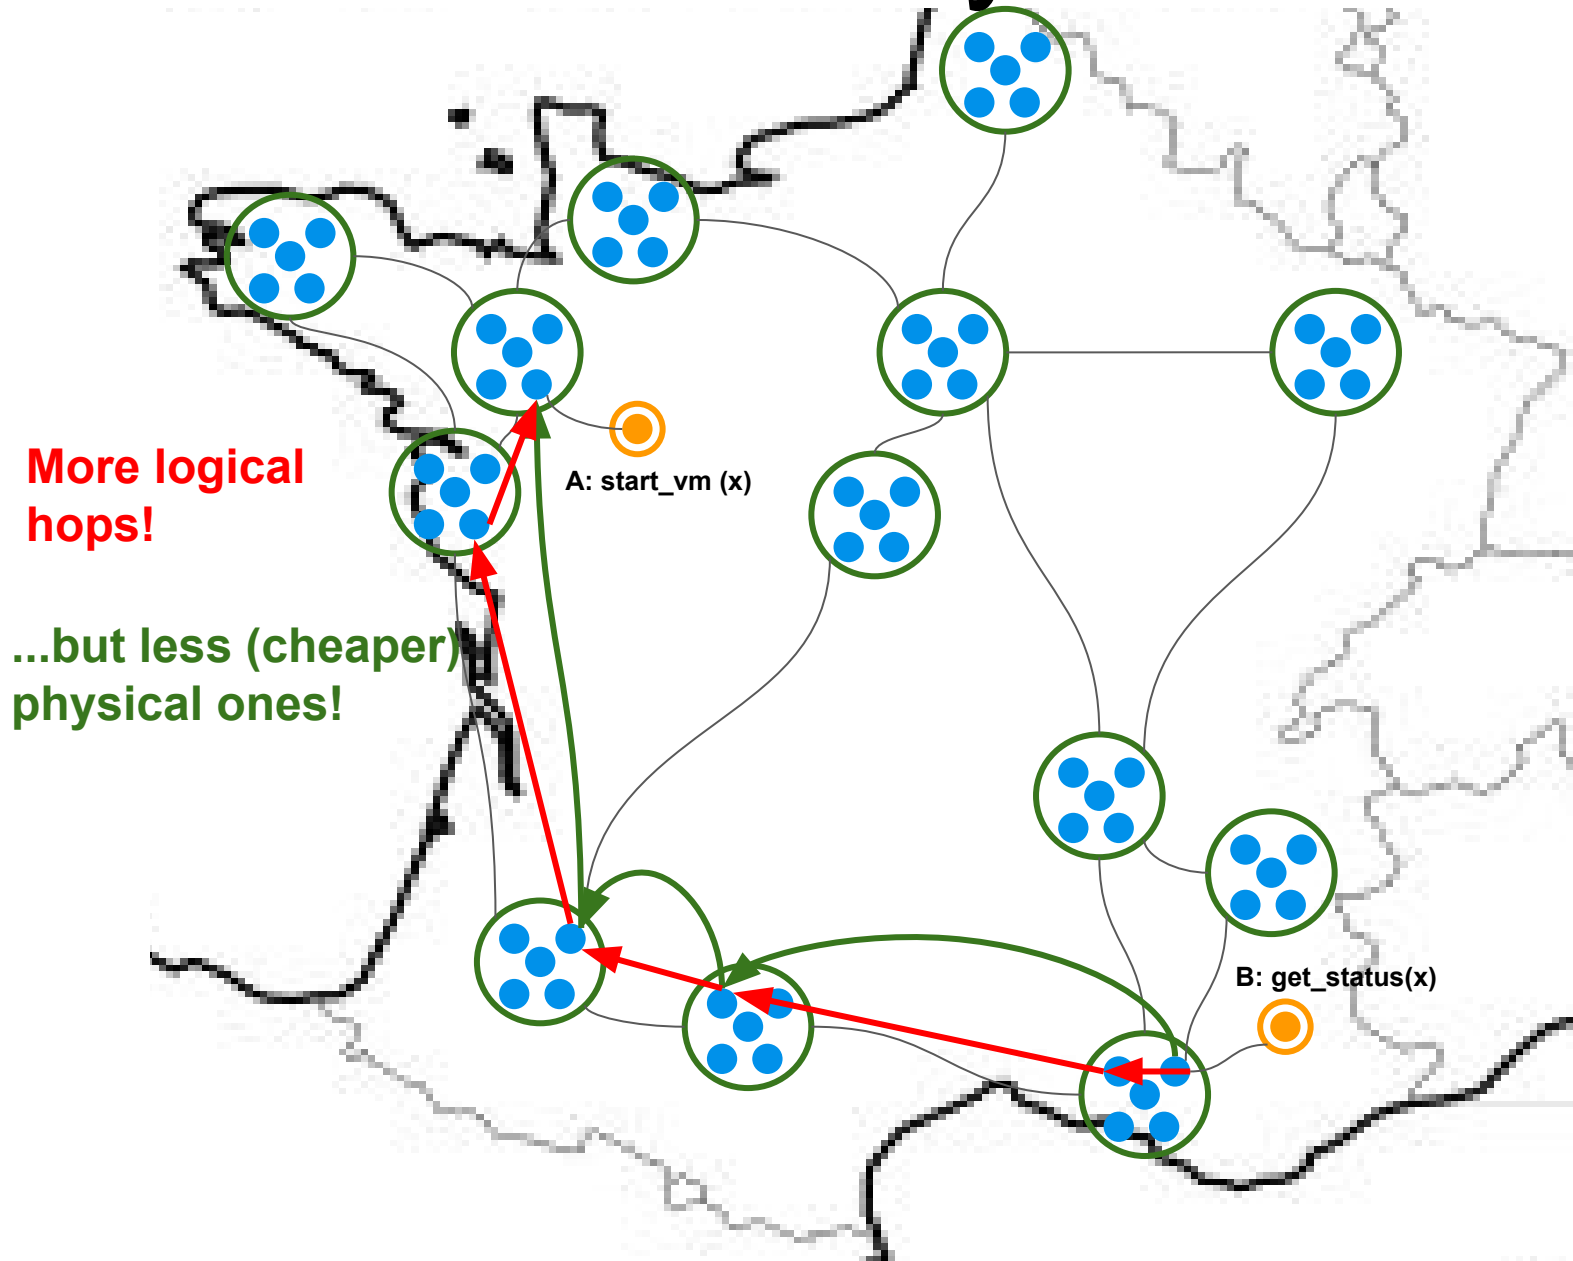


1. Lazy (update overlay only when needed)

Discovery Cloud



Discovery Cloud



Koala: An overlay for decentralized clouds

...but how should this overlay be?

1. Lazy (update overlay only when needed)
2. **Locality-aware**



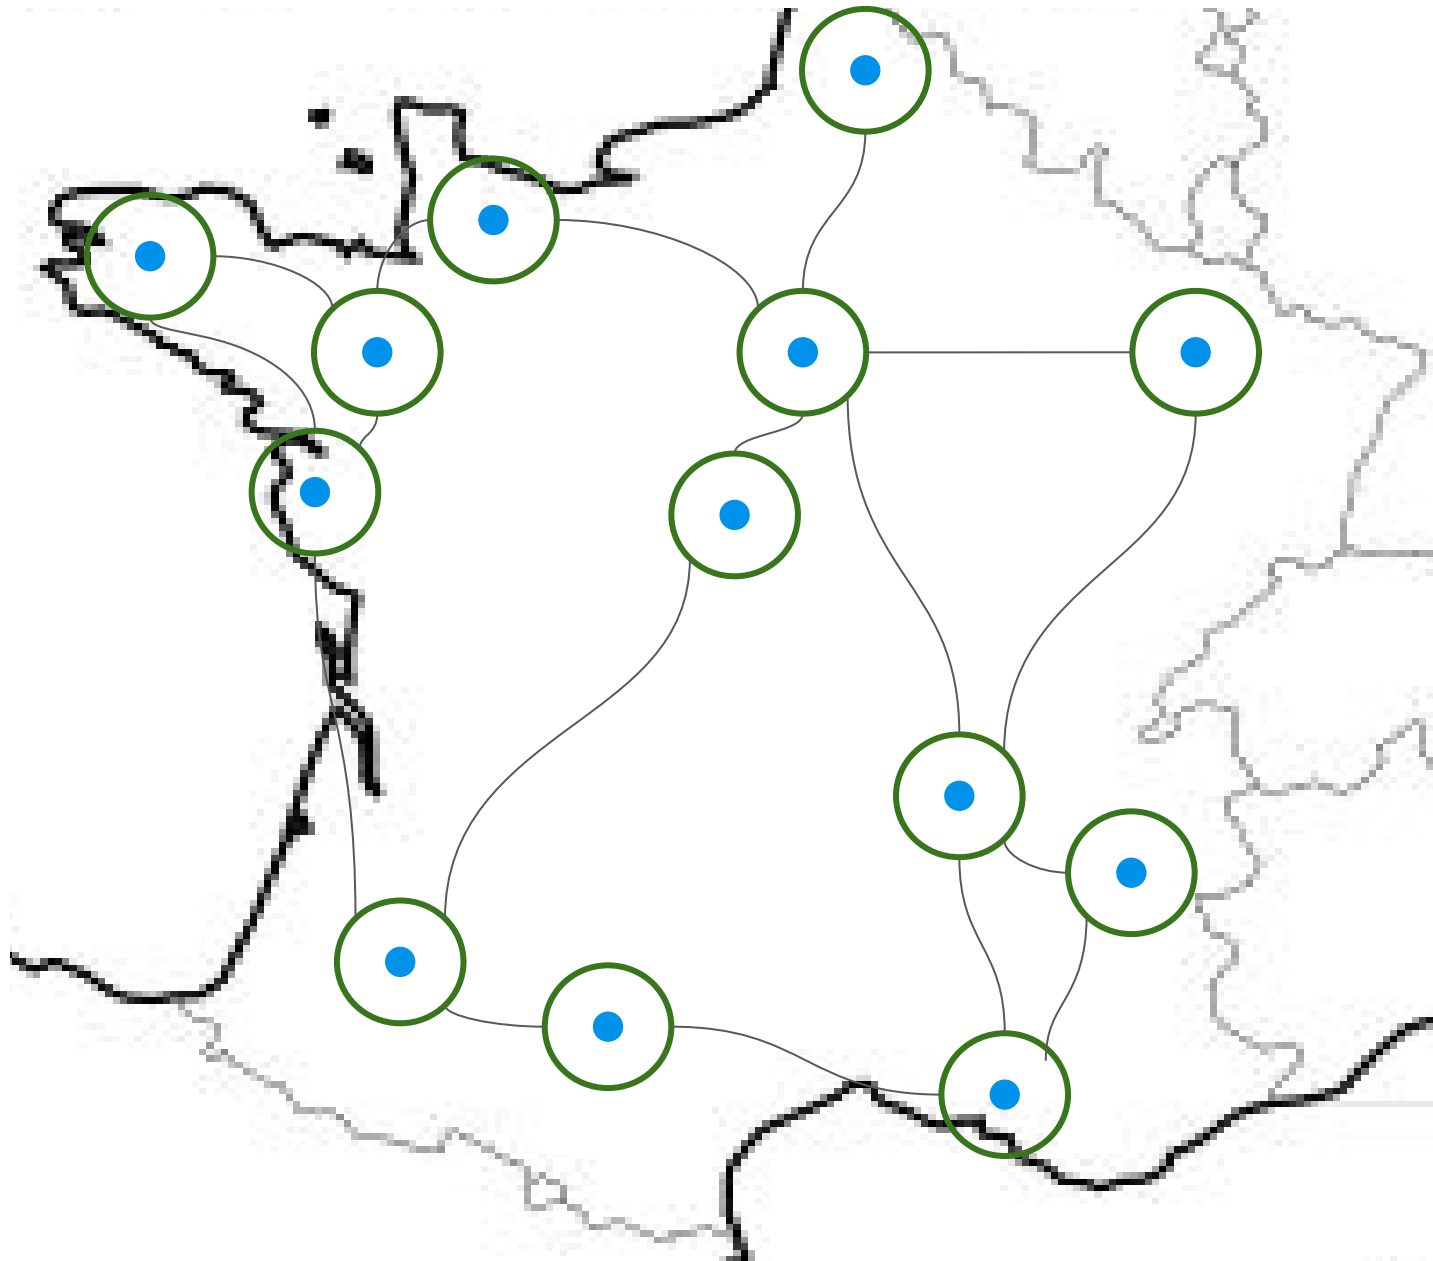
Laziness and locality-awareness

1. How do we implement laziness?
2. How do we integrate locality-awareness?

First, focus on communication between different PoPs.


(nodes in same PoPs is discussed later)

Flat structure



Koala's basics

Nodes are organized in a ring and are identified by a circular id.



Routing table of node 0

ID	IP	RTT	IID
9	a.a.a.a	50	-
1	b.b.b.b	150	-
3	x.x.x.x	125	4
6	y.y.y.y	250	7


Neighbors

Long links

Ideal long links

Still want to be an **$O(\log N)$** hops protocol.

Tell me who your long links are, I will tell you how efficient your routing is!



Continuous Kleinberg distribution:

$$p(d) = 1/(d * \ln N)$$

Where:


d = logical distance,
 N = total nr. of nodes

Generate IDs using $p(d)$ -> **Ideal IDs**

Ideal long links

Generating IDs does not mean contacting them, as these nodes may not exist.

Koala is **lazy**: it does not search, it waits to learn.



Continuous Kleinberg distribution:

$$p(d) = 1/(d * \ln N)$$

Where:

d = logical distance,
N = total nr. of nodes

Close to ideal is still ideal.

Laziness: Piggybacking

Embed information about nodes within the message

2 sources:

1. Nodes in the path
2. Nodes in the routing table of the nodes in the path

Locality-awareness

Find the cheapest logical path (latency-wise)

We could choose systematically the cheapest hop.

...but the logical distance needs to be reduced as well

A tradeoff between logical distance and latency


For each entry we calculate:

$$Q(re) = 1/(\alpha \times re.distance + (1-\alpha) \times norm(re.RTT))$$

Locality-awareness: Routing

Route from **100** to **500**

N=1000

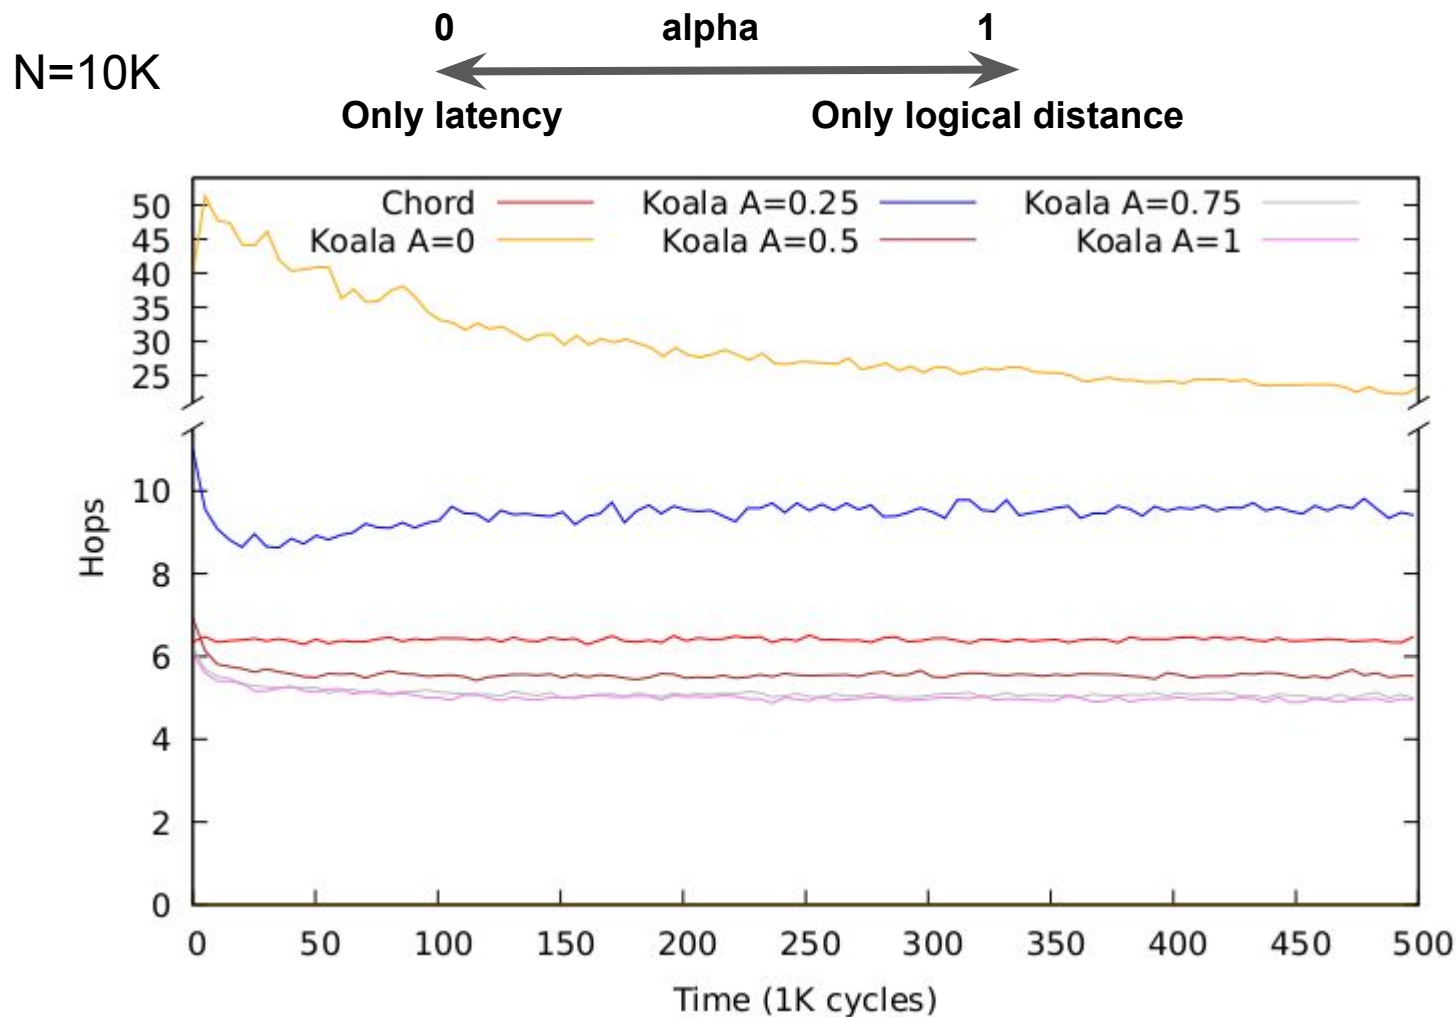


Routing table of node 100

ID	RTT	DistToDest
15	340	485
257	105	243
349	194	151
598	1230	98

$$Q(re) = 1/(\alpha \times re.distance + (1-\alpha) \times norm(re.RTT))$$

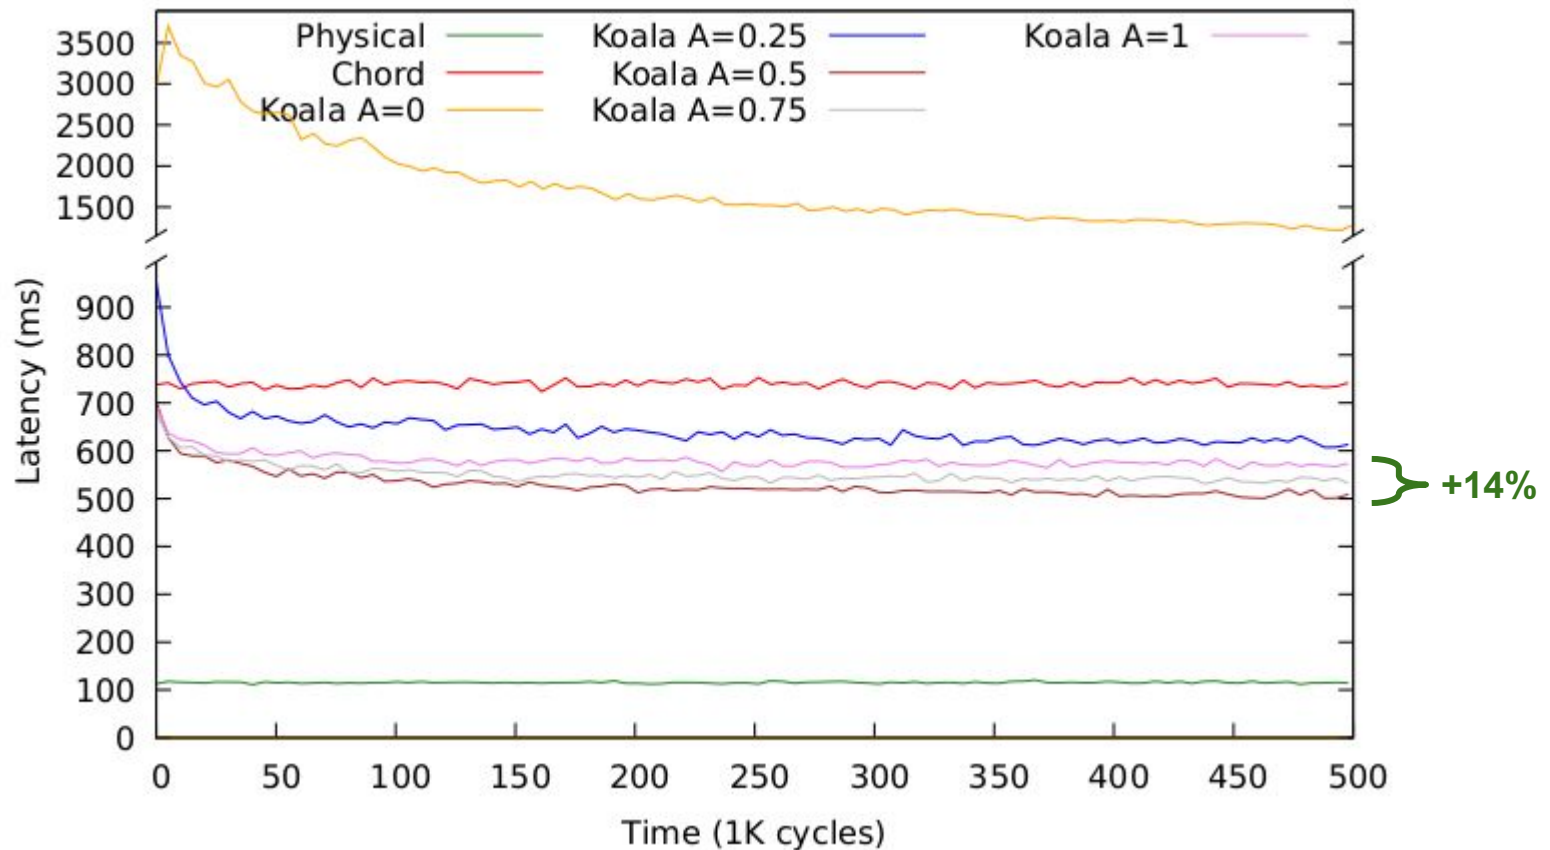
Locality-aware vs Greedy (hops)



Locality-aware vs Greedy (latency)

N=10K

0 ← alpha → 1
Only latency Only logical distance




Conclusion, current and future work

Further evaluation concerning scalability, locality-awareness and resilience have reported in a paper submitted to CloudCom2017.

Currently we are working on the adaptation of our protocol to address the two-layers nature of our physical topology (intra-PoP and inter-PoP) while still remaining flat (no hierarchies)

In the context of the Discovery, we aim at using Koala as a communication bus for OpenStack (a decentralized broker)



Thank you!

Questions?