Efficient Heuristics for the Placement of Large Scale Distributed Applications on Multiple Clouds

Pedro Silva
Frédéric Desprez, Christian Perez

Application Placement on (Multi) Cloud

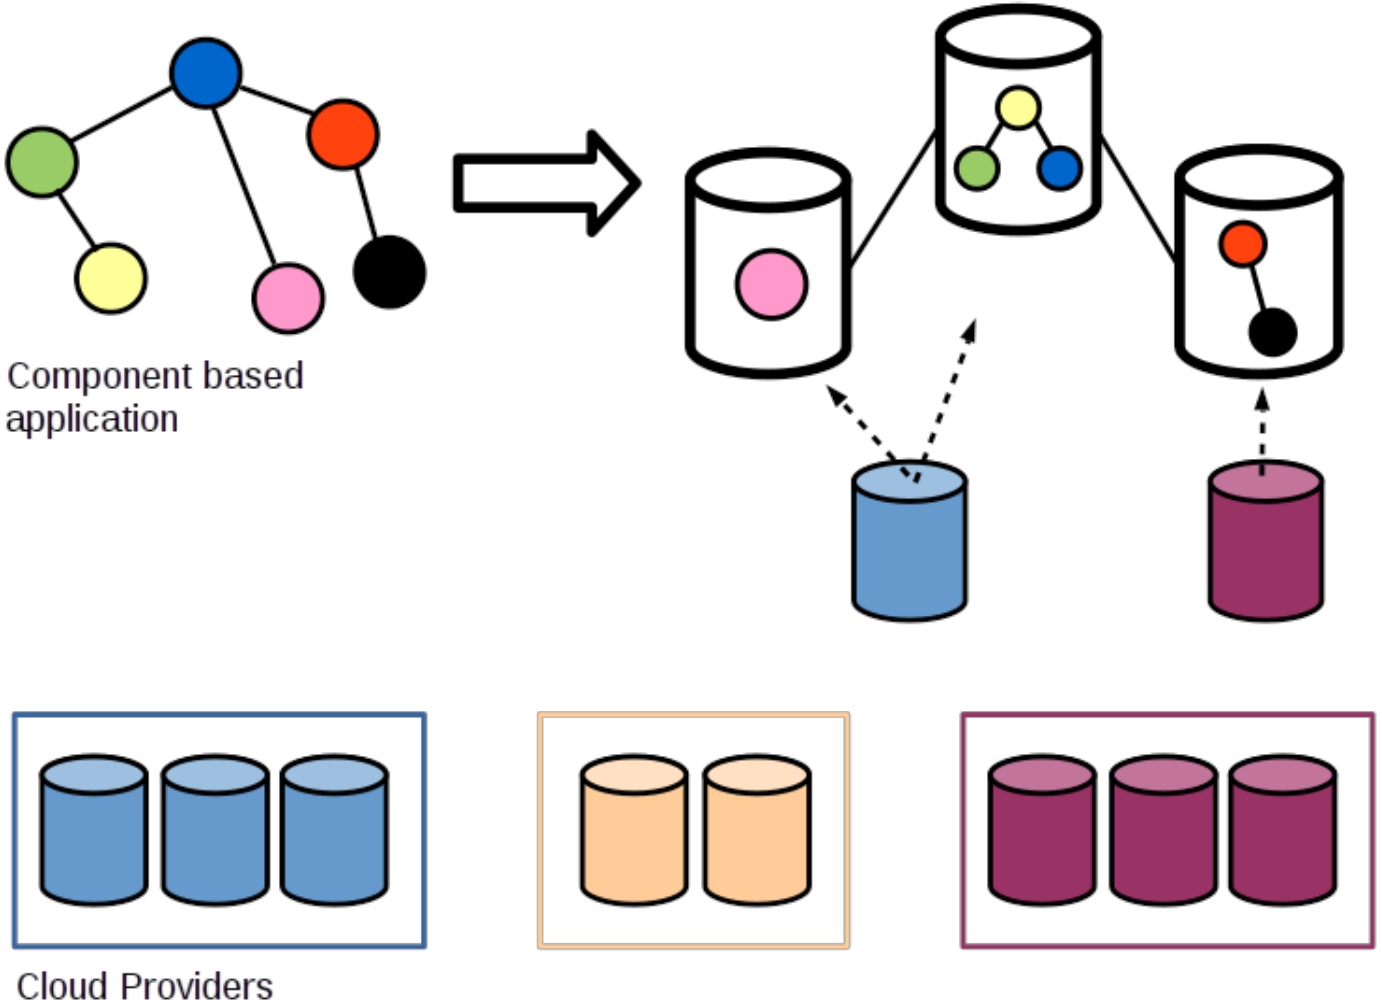

Resource Constraints

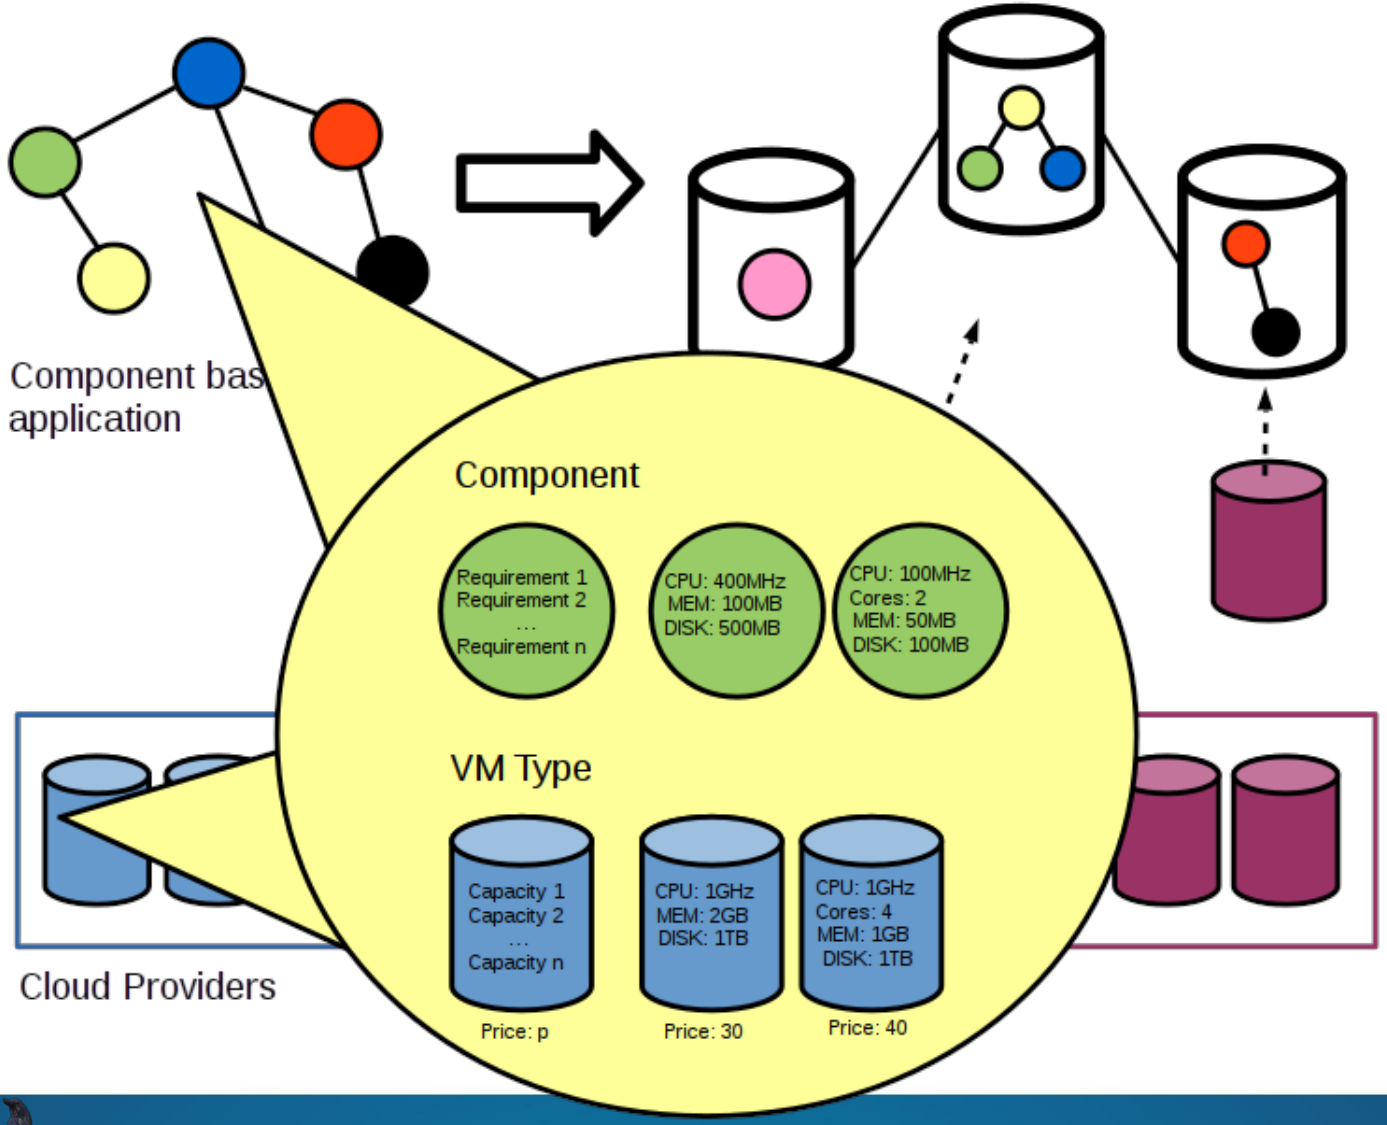

Communication Constraints

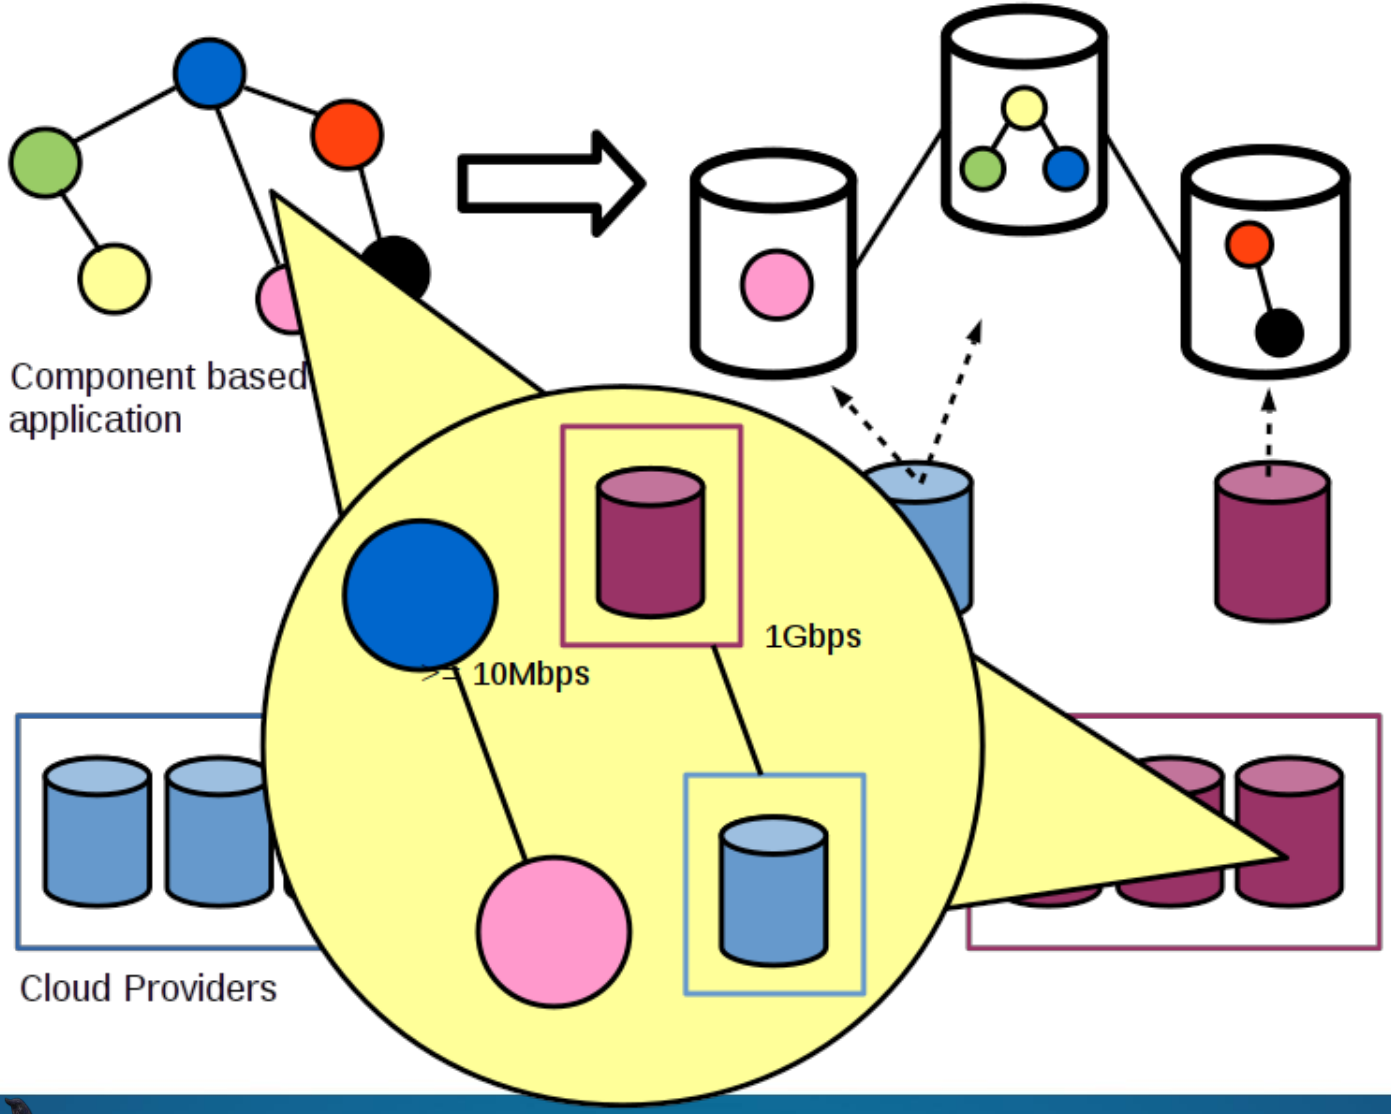

Problem Statement / Methodology

Minimize cost and satisfy resource and communication constraints

Problem divided into three sub problems

- Initial placement, communication oblivious, concrete application
- Initial placement, communication aware, concrete application
- Reconfiguration, communication aware, concrete application
- Toward parametric/abstract application

Evaluation through comparison to the state of the art

Approaches to solve the placement problem

Exact Algorithms

- Solvers
- Model details, optimal
- Scalability

Meta-heuristics

- Genetic algorithm, Simulated annealing, Particle Swarm optimization, etc.
- Model details, generic
- Execution time

Heuristics

- Greedy heuristics, graph based (clustering, partitioning, etc)
- Scalability
- Model details, solution quality

Initial Placement, Communication Oblivious

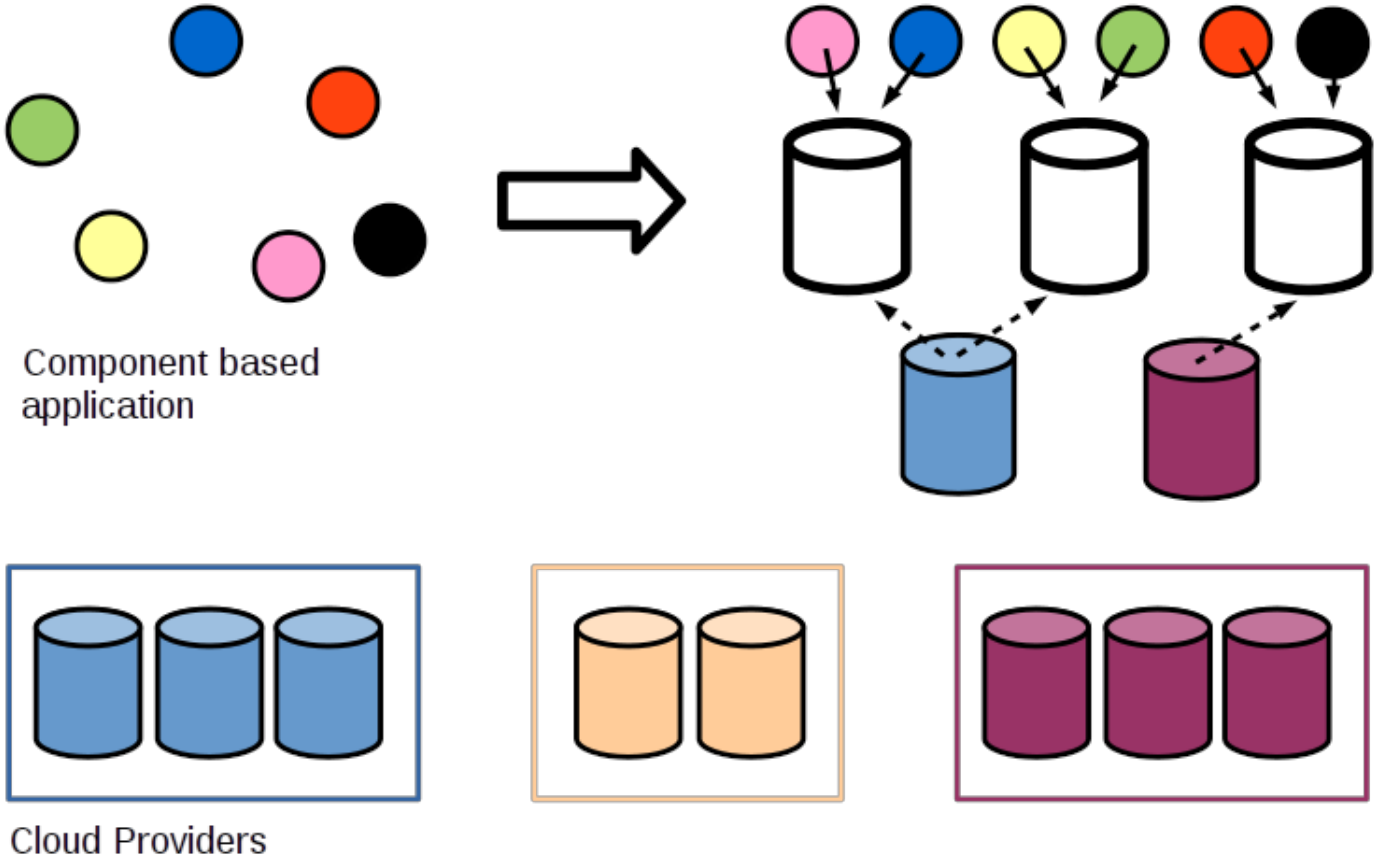

P. Silva, C. Pérez, F. Desprez. Efficient Heuristics for Placing Large-Scale Distributed Applications on Multiple Clouds. In 16th IEEE/ACM International Symposium on Cluster, Cloud and Grid Computing (CCGrid'16), May 2016, Cartagena, Colombia. 2016, <10.1109/CCGrid.2016.77>. <hal-01301382>

Contribution: Greedy Heuristics

Adaptation of three heuristics

- First Fit Decreasing Priority
  - The size is the maximal normalized value of dimensions.
- First Fit Windowed Multi-Capacity
  - It tries to balance the usage of resources from bins.
- Best Fit Dot Product:
  - Most-adapted bin: Bin that maximizes the dot product between item and bins:

Adaptations

- Cost-awareness
  - Ordered bin types by size/price ratio
- Heterogeneity of bin types
  - Created an ideal bin type as reference

Problem Classes

Compared approaches

- **Exact algorithm**: IPL Solver (SCIP solver)
  - 30h per problem
- **Meta heuristics**: Simulated annealing (SA)
  - 10 min per problem
- **Block of greedy heuristics**
  - Only the lowest cost is kept
  - The execution times are summed up

Generated problems (Uniform random)

	Resource/Dimensions	Component/Items	VMs/bin types	Total
A	1, 2, 4, 8	1, 2, …, 19	100, 200, …., 1000	1520
B	1, 2, 4, 8	10, 20, ..., 100	1k, 2k, …., 10 k	800

Evaluation

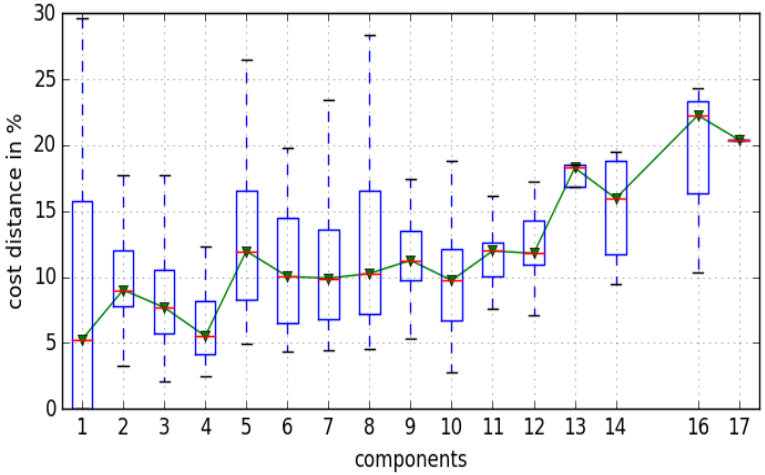

Solver vs. Heuristics on Class A

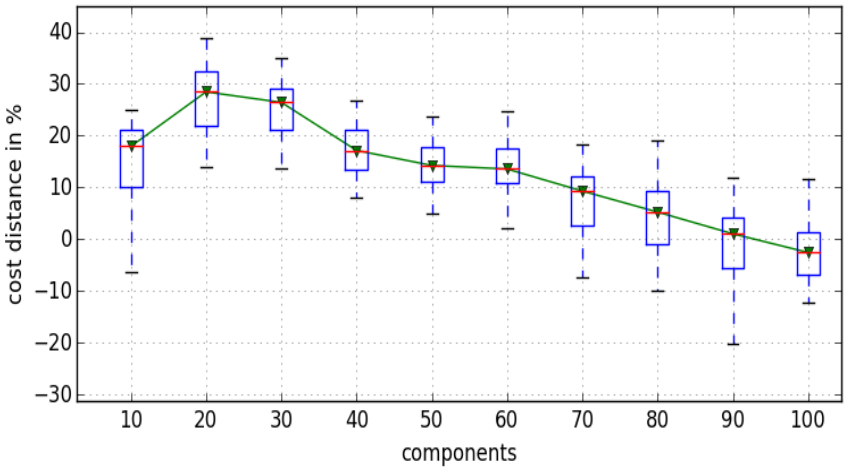

S.A. vs Heuristics on Class B

Solver

- Solve around 34% of Class A

	Average (s)	Maximum (s)
BF-DP	1.40	5.70
FFD-P	0.03	0.23
FFD-WMC	0.10	0.35

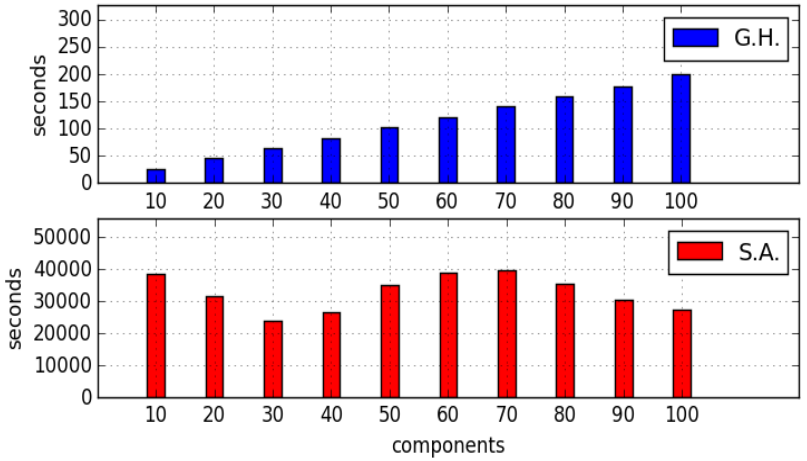

Initial Placement, Communication Aware

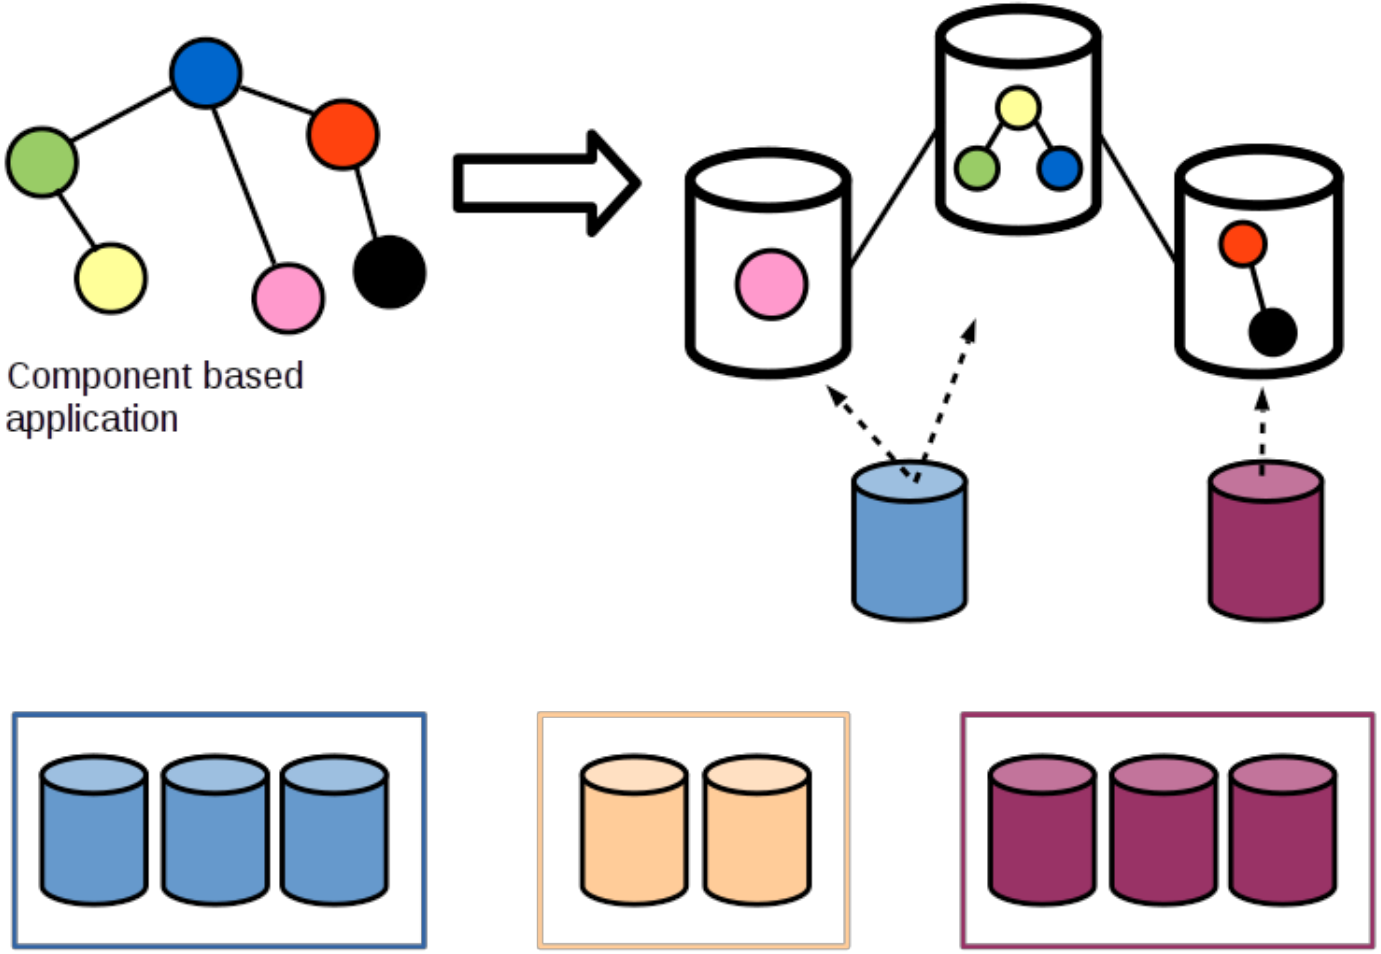

P. Silva, C. Perez, F. Desprez, An Efficient Communication Aware Heuristic for Multiple Cloud Application Placement Submitted to Euro-Par 2017

A Hierarchical Resource Model

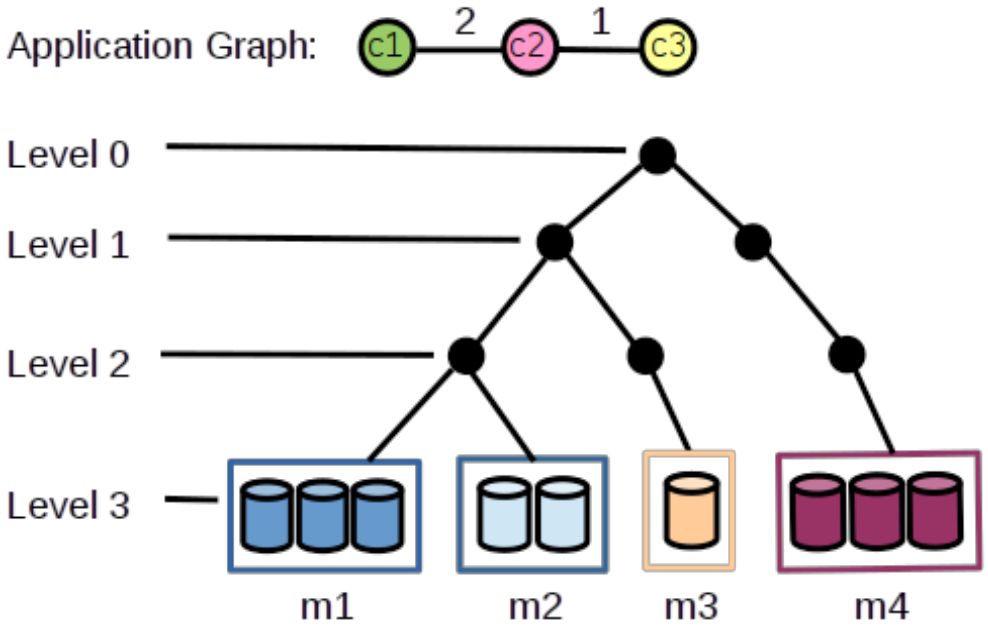

A Greedy Heuristic: 2PCAP

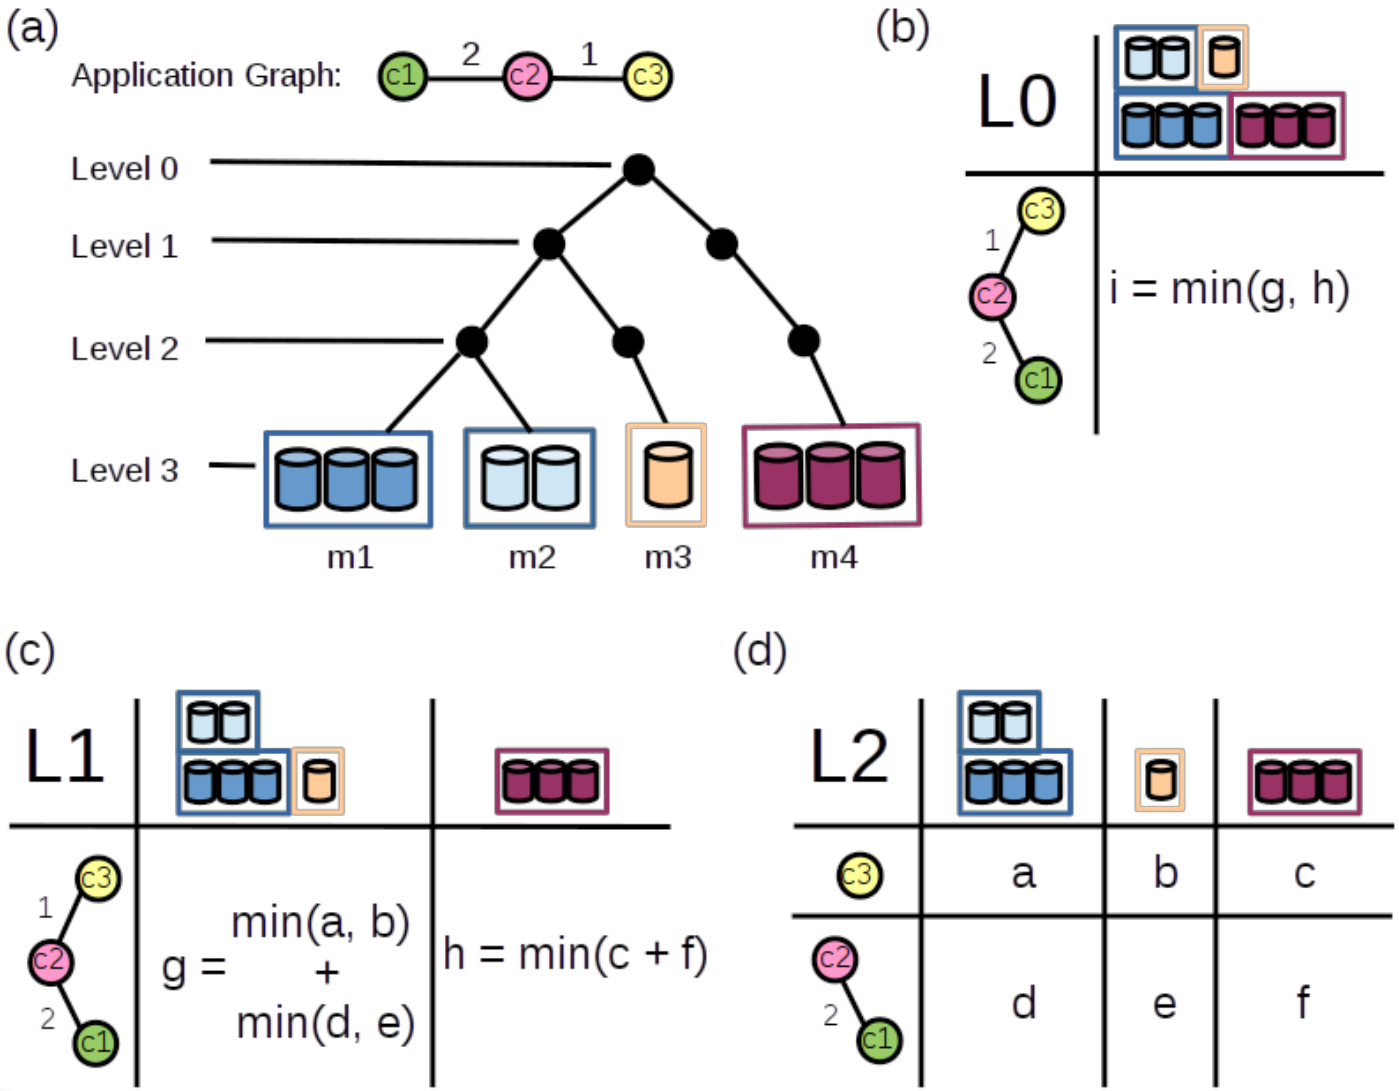

Evaluation: Exact Solver vs SA vs 2PCAP

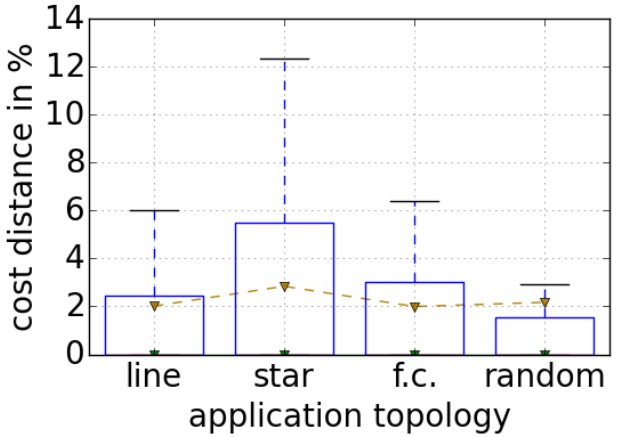

Cost: 2PCAP vs MIP - Small

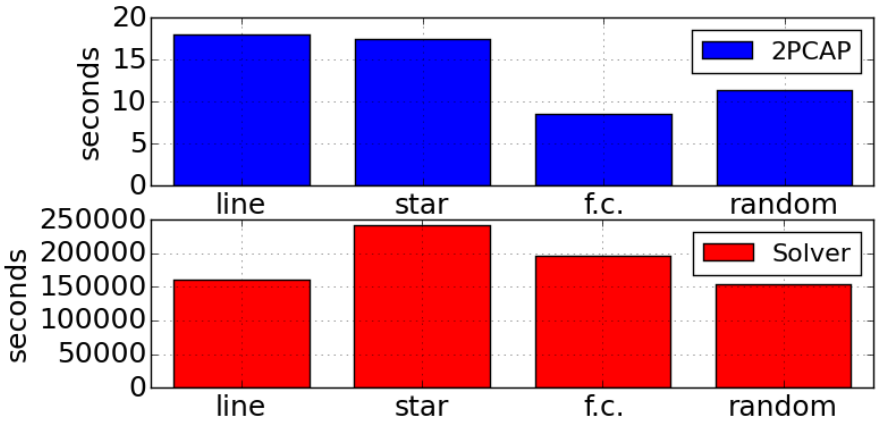

Exec. Time: 2PCAP vs MIP - Small

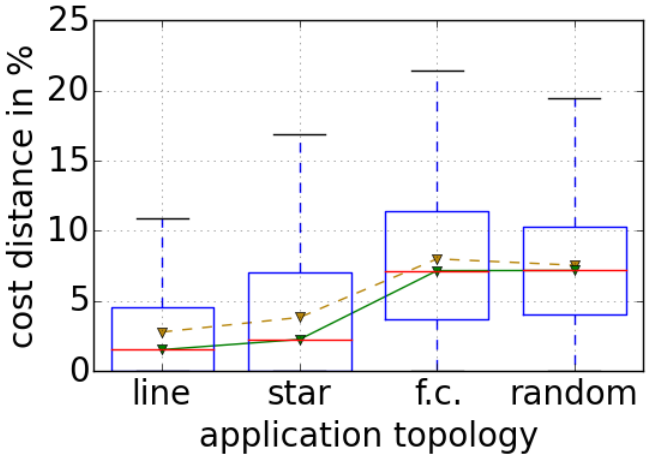

2PCAP vs Baseline Large (No communication)

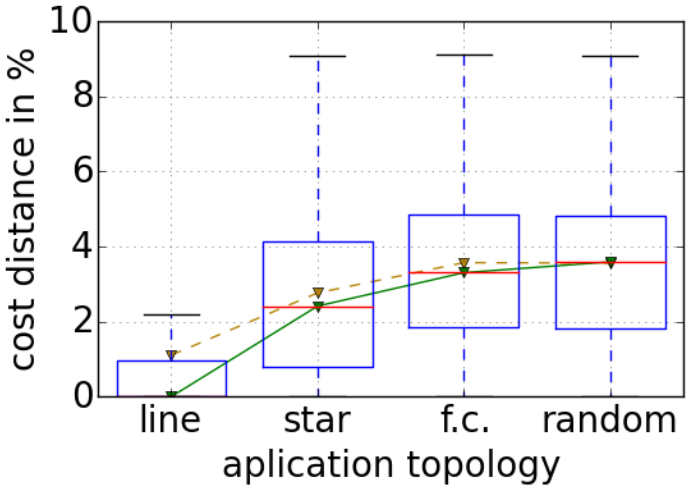

2PCAP + S.A. - Large

Reconfiguration (WIP)

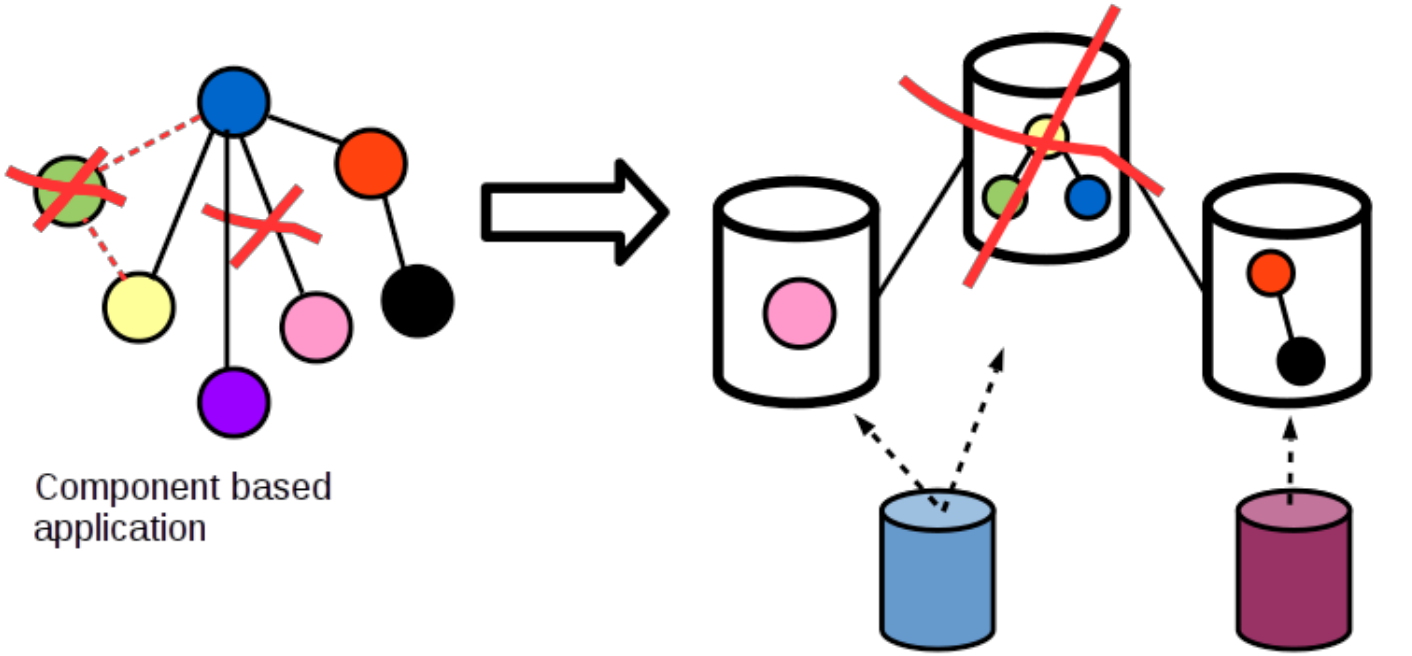

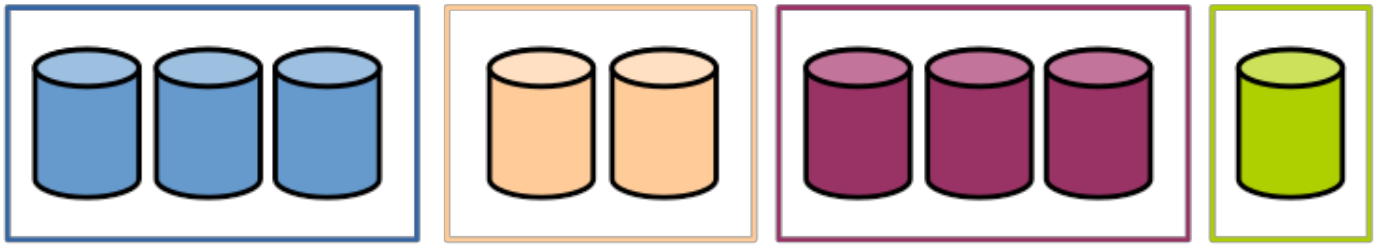

Current Status

Adaptation of 2PCAP

- Adding a migrating cost
- Not the smartest one but based on the cheap computation of a solution

First version of the algorithm implemented

Preparing experiments for evaluation

Parametrizing Application Model

Component Cardinality

- How many instances of each component class needs to satisfy a given performance constraint?
- Need a performance model
- Need a connector model

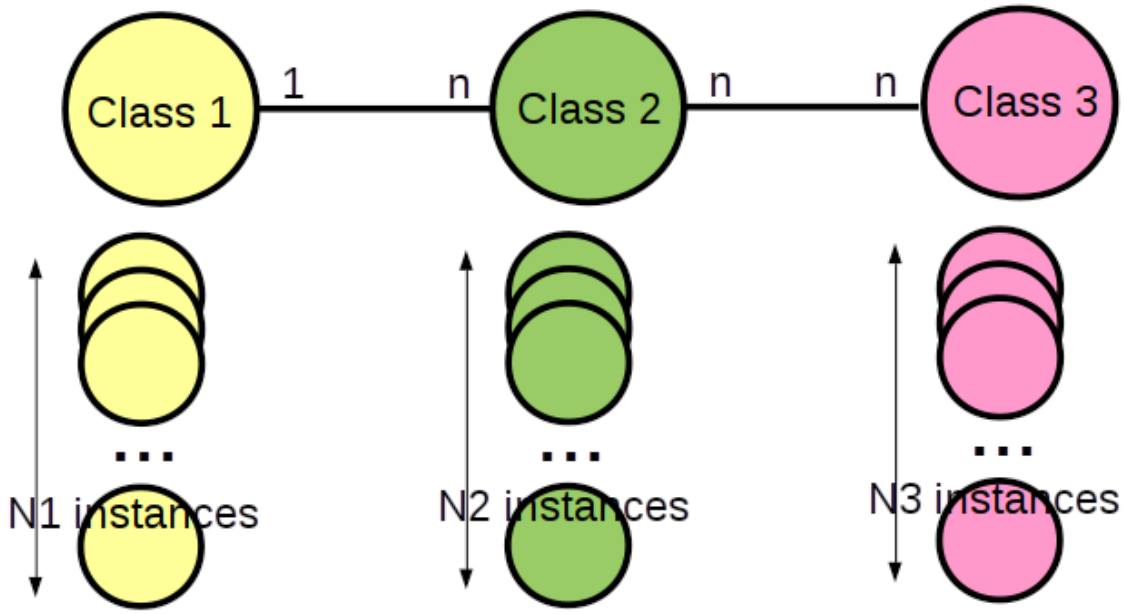

Status

- 1st modeling done
- 2PCAP modified
- Implementation in progress

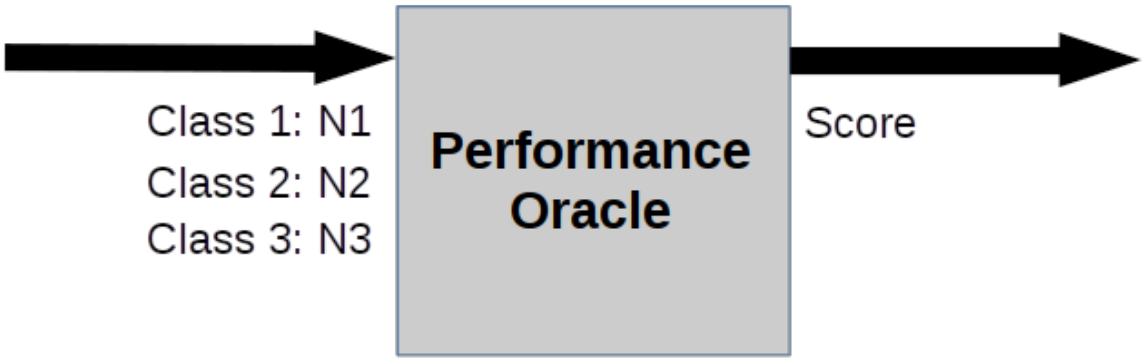

Conclusion and Relationship with Discovery

Heuristics to quickly compute a solution

- Large number of VM types (10,000')
- Moderate application size (100')
- Parametrized number of dimension (~8)
- Hierarchical communication constraint model
- Any concrete application graph
  - Extension toward component cardinality in progress

Relationship with Discovery

- Set of efficient algorithms to deal with initial placement
- Direction of future work
  - Application model and Discovery model
  - Hierarchal modeling and localization constraints
  - Distributed reconfiguration vs "centralized" heuristics
- Determine use cases to understand objective functions