

P2P OPENSTACK

SCALABILITY AND THE EDGE USE-CASE



João Monteiro Soares
Ericsson Research

DISTRIBUTED CLOUD AND THE EDGE USE-CASE



Hierarchical vs P2P

› Hierarchical

- + Scalable
- + Simpler logic
- + Effective/deterministic
- + Efficient
- Single entry point (root)
- Failure handling difficult
- Can't handle churn



› (Unstructured) peer-to-peer


- + Extremely scalable
- + No single point (of failure or bottleneck)
- + Can handle churn well
- Probabilistic/approximate
- More complex logic
- More overhead



DISTRIBUTED CLOUD AND THE EDGE USE-CASE



- › How to manage a cloud infrastructure with such scale?



DISTRIBUTED CLOUD AND THE EDGE USE-CASE




> How to manage a cloud infrastructure with such scale?

– Seeking key design principles:


- > Avoid centralized components (as much possible)
- > Avoid (Minimize) any changes to the existing software

– Proposal:


- > Use P2P techniques to horizontally scale OpenStack
- > Provide a single-cloud abstraction



SINGLE CLOUD ABSTRACTION




SINGLE CLOUD ABSTRACTION



A P2P APPROACH




- › Divide an OpenStack deployment into a (possibly large) number of smaller OpenStack ‘cloudlets’, each with its own controller and compute nodes
- › Associate an agent with each cloudlet
(Note: agent functionalities could be directly embedded in OpenStack)
- › Agents forward requests to other cloudlets (reverse proxy)
- › Implement the single-cloud abstraction
- › OpenStack APIs ‘untouched’ (remain the same)



A P2P APPROACH



- › Projects/tenants are mapped to agents
- › When an agent receives a request it
 - first chooses the cloudlet responsible for that request
 - forwards the request to the controller of the chosen cloudlet
 - updates a local mapping table upon receipt of the response
 - returns the reply to the user




A P2P APPROACH



- › An agent implements three key functions
 - Service Proxy (Keystone, Nova, Glance, Neutron, ...)
 - Overlay Management
 - State Management

› Agent-to-Agent Interface



A P2P APPROACH



- › AA(A) - Keystone Federation
 - An Identity Provider (IDP) outside of the cloud
 - Keystone in each cloudlet configured as a Service Provider (SP)
- › Images:
 - external storage, only add location to Glance
- › Block storage (tbd)
- › Networks (tbd):
 - e.g. OpenStack Tricircle

A P2P APPROACH: SCHEDULING



- › Randomized algorithms for:
 - Peer sampling: choose random peers
 - › Peer Sampling Protocol, **Cyclon**, Newscast, SCAMP
 - Distributed search: determine which node takes the workload
 - › **“The Power of Two Choices”**

CYCLON



- › Each node maintains a random ‘cache’ of nodes
 - › Periodically this cache is exchanged with the oldest node in the cache
 - › Protocol attempts to keep the size of the cache ‘constant’
- After the protocol stabilizes (in $O(\log n)$ steps), each node’s cache holds a random subset of the other nodes in the system



Voulgaris, S., Gavidia, D. and van Steen, M. (2005) ‘CYCLON: Inexpensive membership management for unstructured P2P overlays’, *Journal of Network and Systems Management*, 13(2), pp. 197–217. doi: 10.1007/s10922-005-4441-x.

THE POWER OF TWO CHOICES



- › Objective is to balance load
- › Very simple algorithm:
 - Select two nodes at random
 - Send the workload to the node with smallest load
- › Compared to randomly choosing a node, the maximum load decreases from $O(\log n)$ to $O(\log \log n)$

Mitzenmacher, M. (2001) 'The power of two choices in randomized load balancing',
IEEE Transactions on Parallel and Distributed Systems, 12(10), pp. 1094–1104. doi: 10.1109/71.963420.

P2P VS. STANDARD

SYSTEM PERFORMANCE





› Standard:

- 1 Controller (8 cores/16GB RAM)
- 64 Compute nodes (1 core/ 2GB RAM)

› P2P :

- 64 Cloudlets
- 1 Controller (2 Cores/ 4GB RAM+1GB SWAP)
- 1 Compute (1 core/ 2GB RAM)
- 1 Services (2 cores/4GB RAM)
 - › IDP, file server, Cyclon Introducer

P2P VS. STANDARD SYSTEM PERFORMANCE



MOVING FORWARD



- › Current approaches use hierarchical techniques to achieve OpenStack scalability
 - Distributed cloud and the edge case bring another dimension to the problem

- › What about P2P techniques to horizontally scale OpenStack ?
 - A novel approach to address cloud scalability
 - Reference implementation proves feasibility of the approach
 - Several challenges still remain

- › Next steps:
 - Support location requirements
 - Open Source reference implementation



ERICSSON