



Plenary Discovery 6. Nov 2017

Study of OpenStack Internal bus in a Fog/Edge Context

Orange/Inria - started 1. Oct 17

Abdelhadi Chari (Orange)

Adrien Lebre (IMT Nantes)

Ali Sanhaji (Orange)

Matthieu Simonin (Inria)

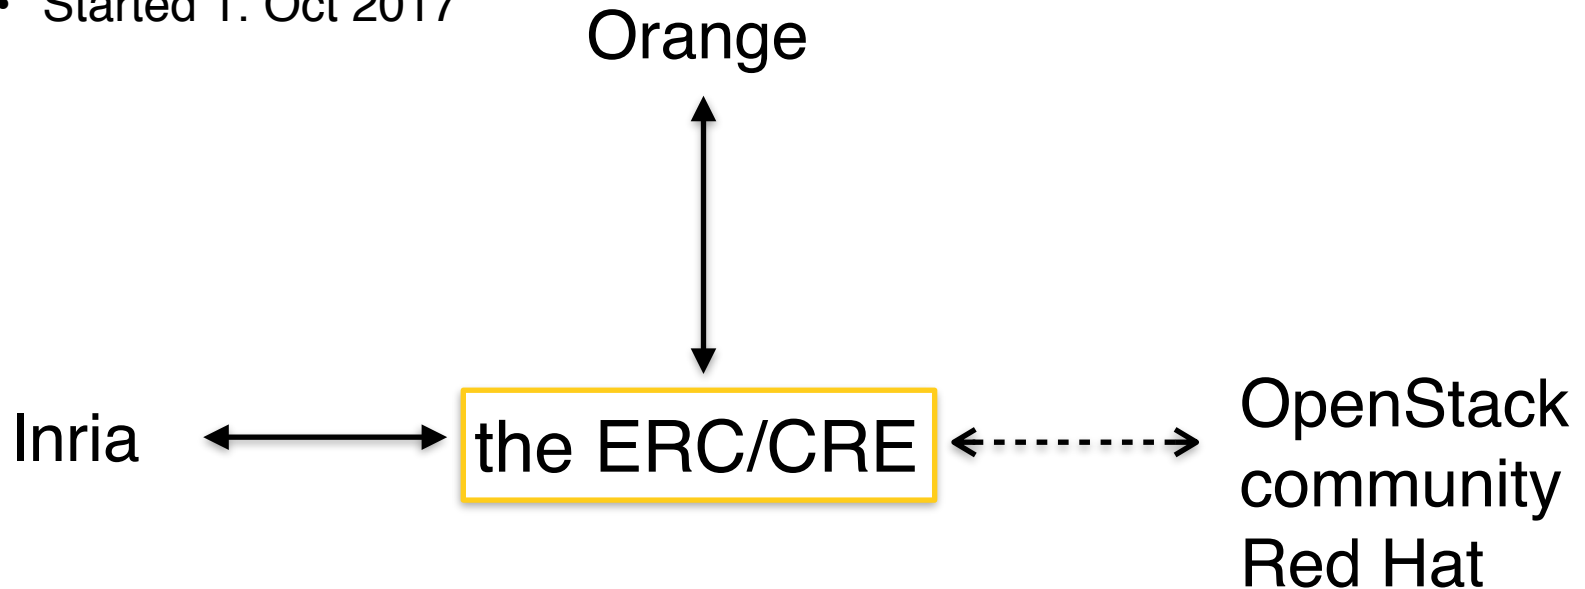
Alexandre Van Kempen (Inria)

01

Context

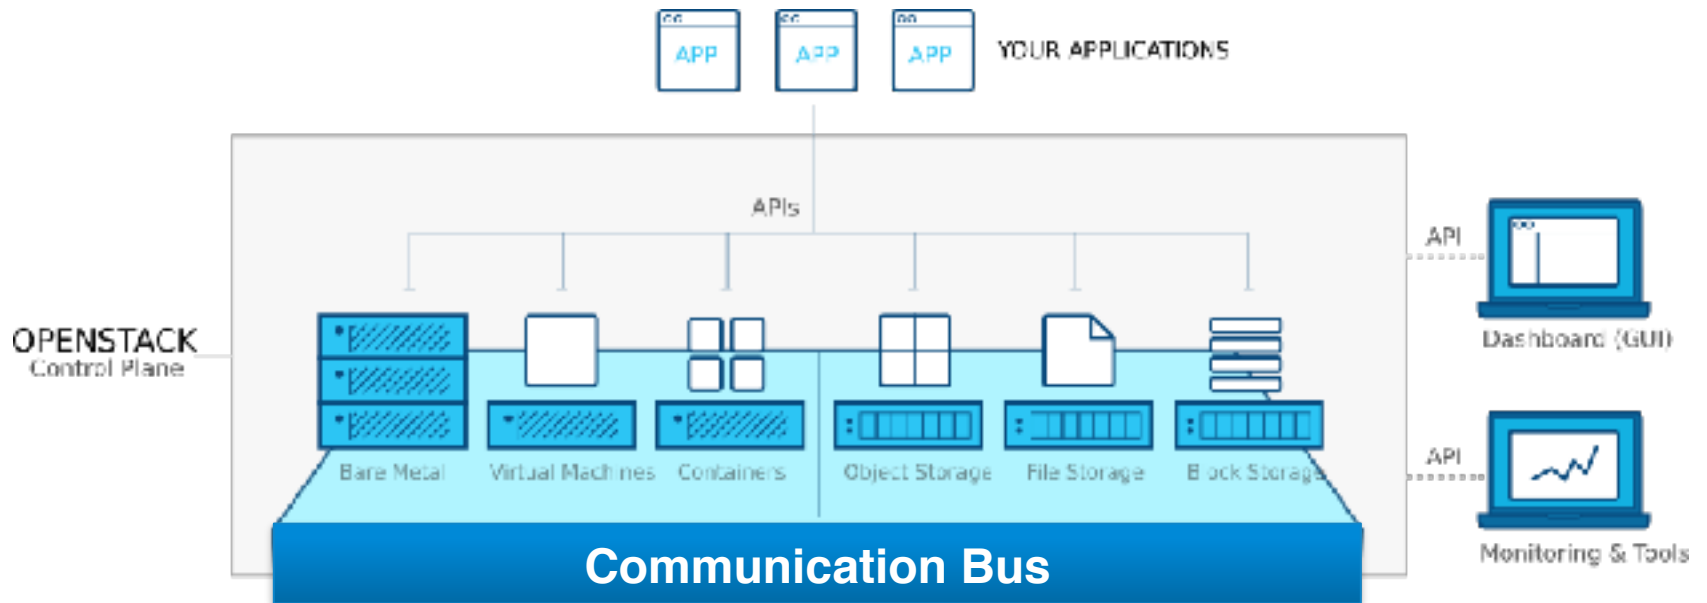
We got an ERC¹ !

- (en)ERC: Externalised Research Contract
- (fr)CRE: Contrat de Recherche Externalisée
- Signed between Orange and Inria
- Started 1. Oct 2017




1. this is not an ERC grant from the European Research Council

Technical environment : Openstack

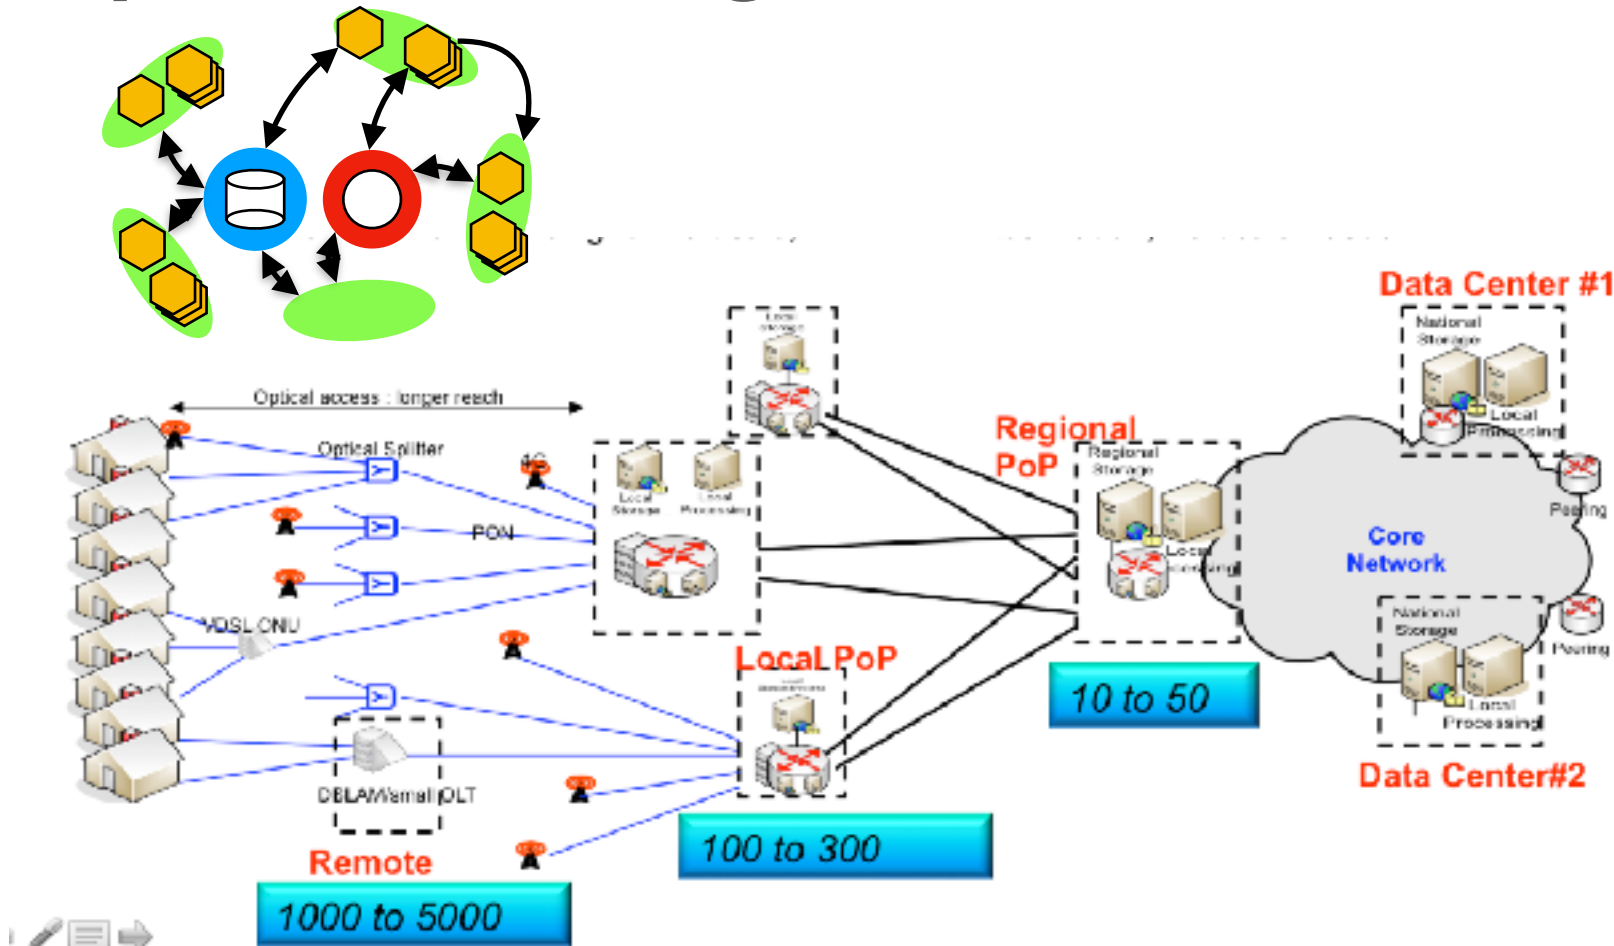


Is the communication bus of OpenStack Fog/Edge ready ?

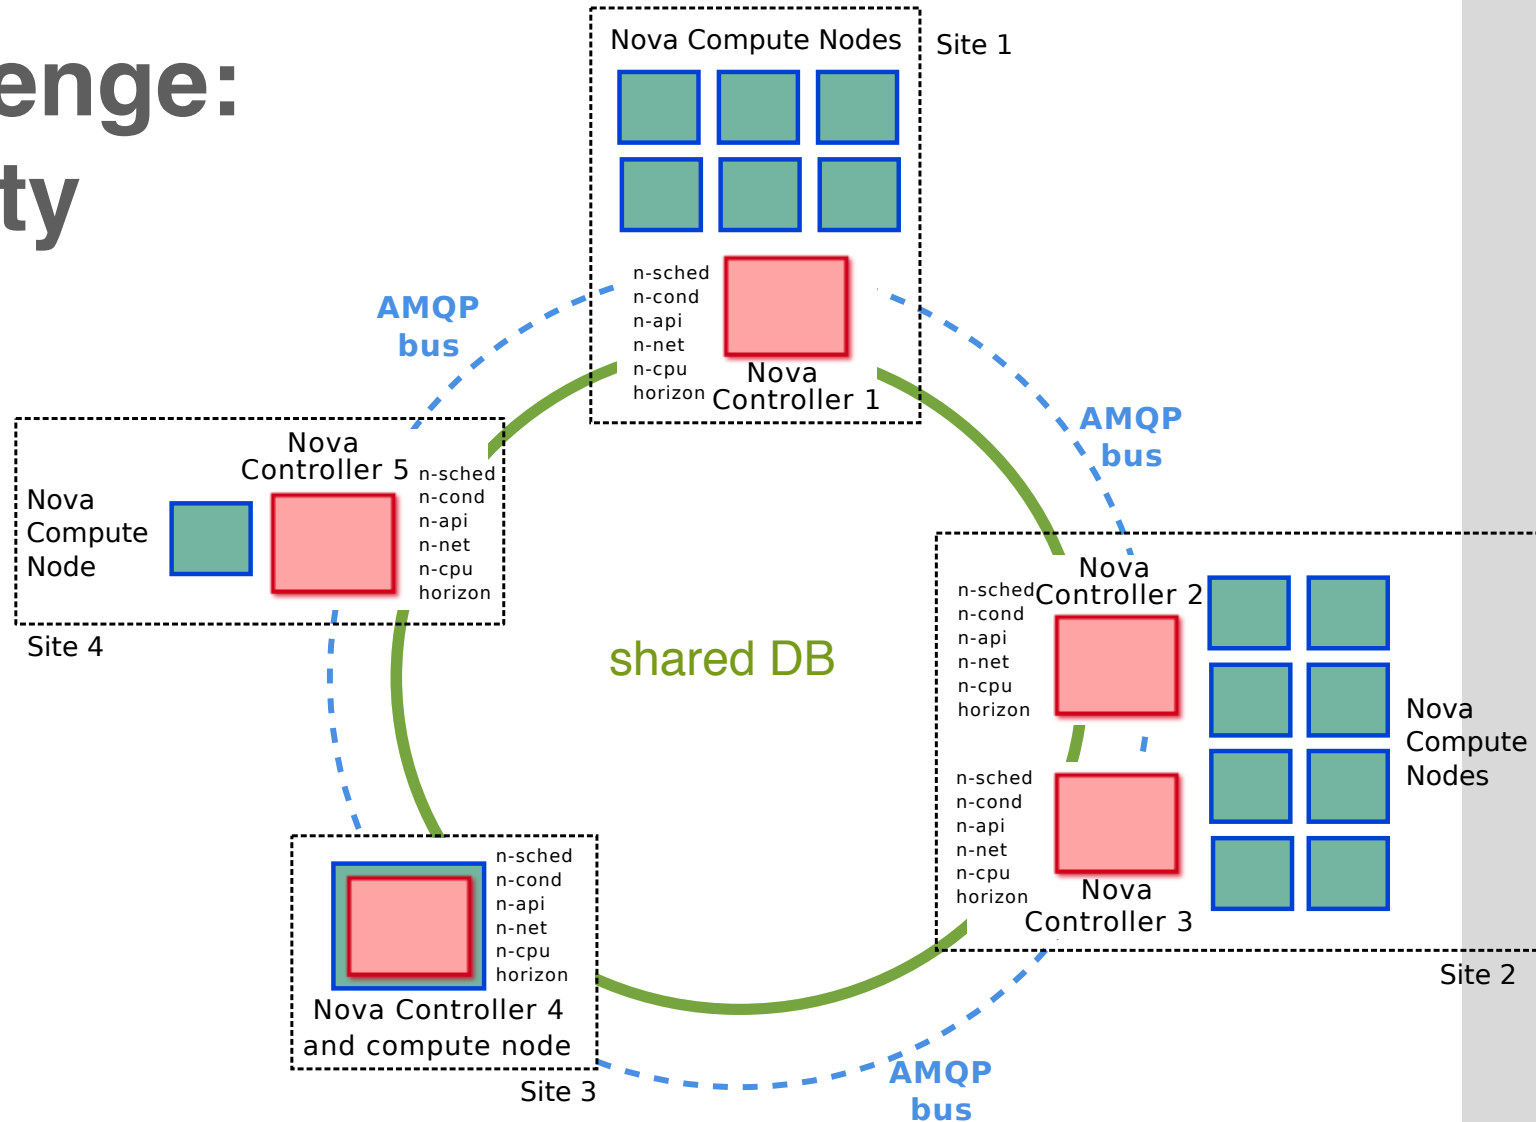
OpenStack



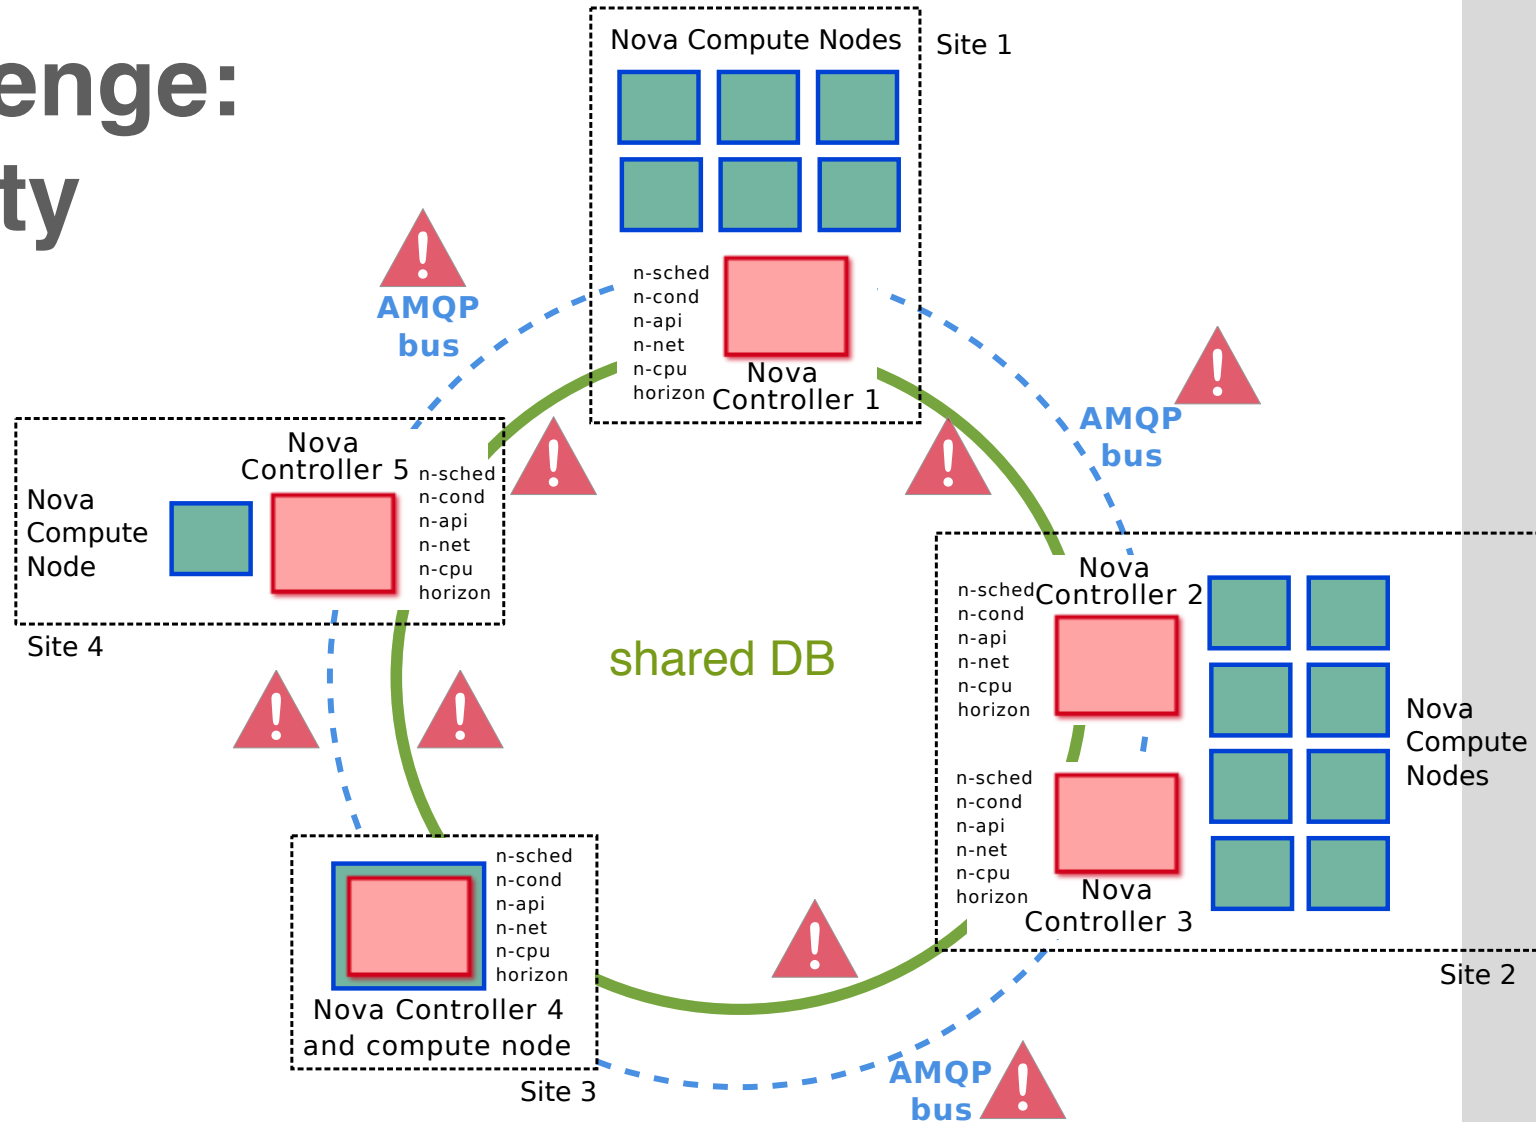
OpenStack: target infrastructure



Challenge: locality



Challenge: locality




Challenge: locality

- Distributing the database :
 - ~~Ronan's talk tomorrow~~ Next Talk from Ronan
- Distributing the message bus
 - Current talk

02

Communication bus of OpenStack

OpenStack : bus implementation



common library

also.messaging

communication
pattern

RPC

Notification

transport

RabbitMQ

ZeroMQ


AMQP1.0

RabbitMQ


AMQP1.0

Kafka


OpenStack : RPC patterns



unicast RPC
blocking & non blocking



anycast RPC
blocking & non blocking



multicast RPC
non blocking only

OpenStack : RPC transport

- RabbitMQ driver (main stream)
 - Centralized broker solution
 - First bottleneck when scaling the deployment
- ZeroMQ driver
 - « embarrassingly » distributed - brokerless
 - only suitable for small deployments
 - less supported
- AMQP1.0
 - new driver (beta version)
 - can use a mesh of routers to route messages

03

Planned Activity

CRE: planned activities

RPC Transport :

- Which RPC transport protocol in the Fog/Edge context ?

Communication patterns and service agents deployment:

- Anti-patterns in the Fog/Edge clouds context?
- Locality requirements

CRE: planned activities

Experimental protocol and objectives are being approved

<https://review.openstack.org/#/c/491818>

- In brief:
 - Emulate a target fog/edge infrastructure on Grid'5000
 - Choose a message bus and deploy it
 - Evaluate all the messaging patterns
 - Evaluate the resilience of the message bus
 - Evaluate the overall performance of (a distributed) OpenStack


CRE: tools

- EnOSlib (Inria - Discovery)
 - Experimental workflow description/deployment/execution
- Ombt (kgiusti - oslo.messaging core dev)
 - Benchmarking & orchestration tool for oslo.messaging
- Os-fault (OpenStack - Performance team)
 - Failure injection framework (agnostic)
- Osprofiler (OpenStack - Performance team)
 - Distributed tracing system
 - -> Identification of the communication patterns

CRE: short term

Short Term: Evaluation of RabbitMQ


- Central RabbitMQ and many edge servers
- Distributed RabbitMQ through federations



CRE: short term

Mid Term: Evaluation of AMQP1.0 (qpid-dispatch-router)

- Collaboration with Red Hat



Vancouver (05/18): Distributing OpenStack

Distribution of the message broker

- Joint work with
 - Red Hat (qpid dispatch router)
 - [oslo.messaging](#)

=> CRE Orange



Distribution of the database

- Joint work with
 - CockroachDB (newSQL implementation)
 - [oslo.db](#)

=> next talk tomorrow

Study of OpenStack Internal bus in a Fog/Edge Context

Orange/Inria - started 1. Oct 17

Abdelhadi Chari (Orange)


Adrien Lebre (IMT Nantes)

Ali Sanhaji (Orange)

Matthieu Simonin (Inria)

Alexandre Van Kempen (Inria)

CRE: using EnOS(lib)



The Discovery Inria Project Lab

Cloud computing

- centralized
- small number of large datacenter
- user locality unaware

Edge computing

- distributed
- large number of small datacenter
- user locality aware



Fog computing

<https://beyondtheclouds.github.io/>

OpenStack Internal bus in a Fog/Edge Context

CRE-Orange/Inria - started 1. Oct 17

Abdelhadi Chari (Orange)
Adrien Lebre (IMT Nantes)
Matthieu Simonin (Inria)

Thank you!

05

Conclusion

Foundation

Board of
directors

Technical
committee

User
committee

Project Teams

Neutron

Glance

Swift

Nova

Keystone

+50 more

Working groups

API*

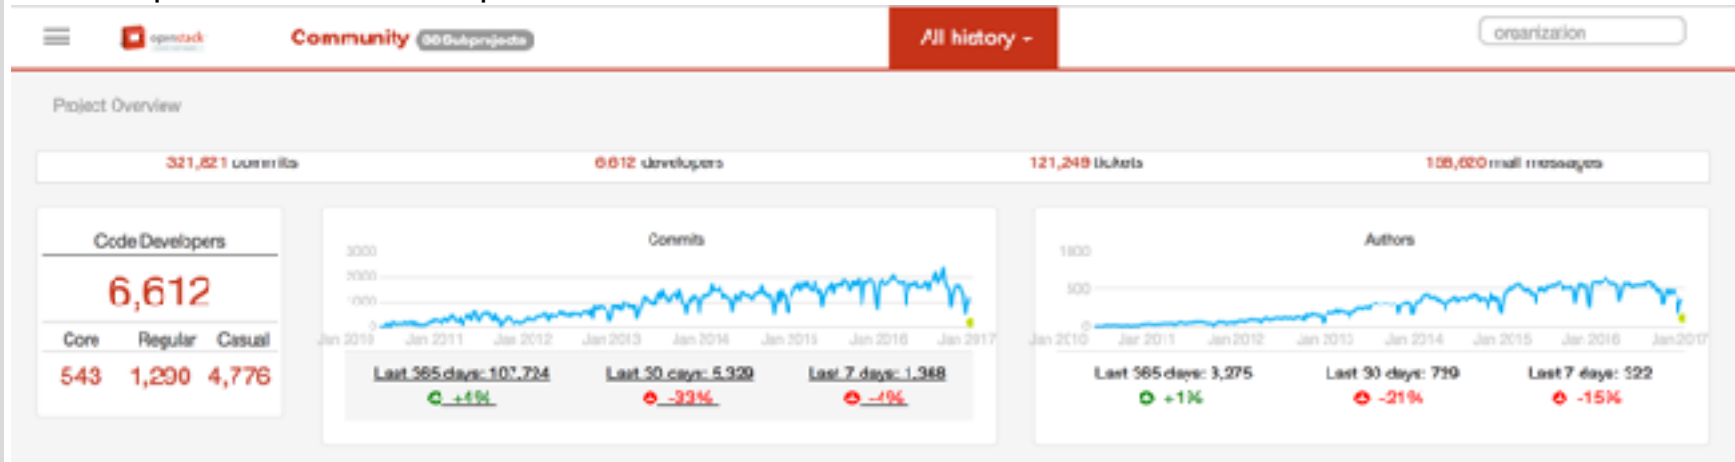
Large Deployment
Team

Fog/Edge Massively
distributed clouds

Scientific*

+20 more

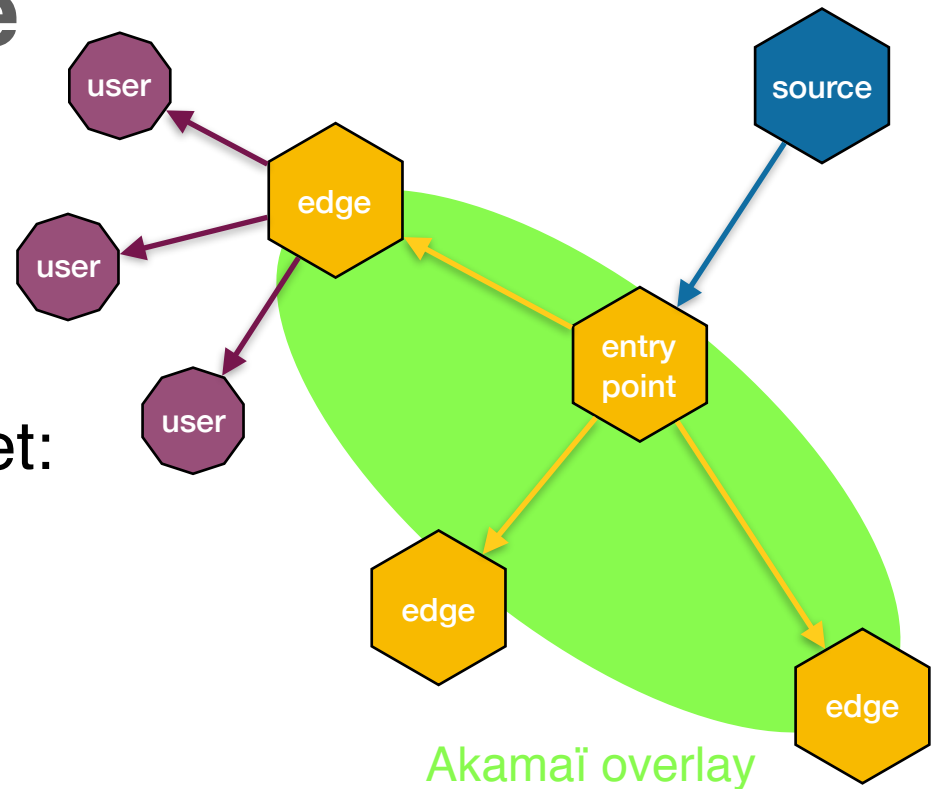
* SIG :Special Interest Group



The Akamai case

Akamai Internet in Internet:

- more distributed
- faster
- more reliable
- but only for content delivery



<https://www.akamai.com/us/en/multimedia/documents/technical-publication/the-akamai-network-a-platform-for-high-performance-internet-applications-technical-publication.pdf>

The Akamai case

Peak Usage*:

- 10+ millions simultaneous video streams
- traffic :50+ Tbps

* estimation made in 2009 for the next 5 years

<https://www.akamai.com/us/en/multimedia/documents/technical-publication/the-akamai-network-a-platform-for-high-performance-internet-applications-technical-publication.pdf>

02

EnOS: performance toolkit for OpenStack

From the performance angle

Goal:

- Discovery/Inria to be visible in the OpenStack community
- Fog/Edge use case to be visible

Mean:

- Working groups (Performance + FEMDC)
- Summit presentations
- **A toolkit for the performance study of OpenStack**

Challenge:

- Small task force

« The PhD student problem »

This includes :

- Dealing with OpenStack deployment
- Maintaining code + deployment
- Ensuring the reproducibility (at least reusability) of the experiments
- Scaling the experiment

In summary :

- **What makes an experimental validation technically trustworthy**

EnOS

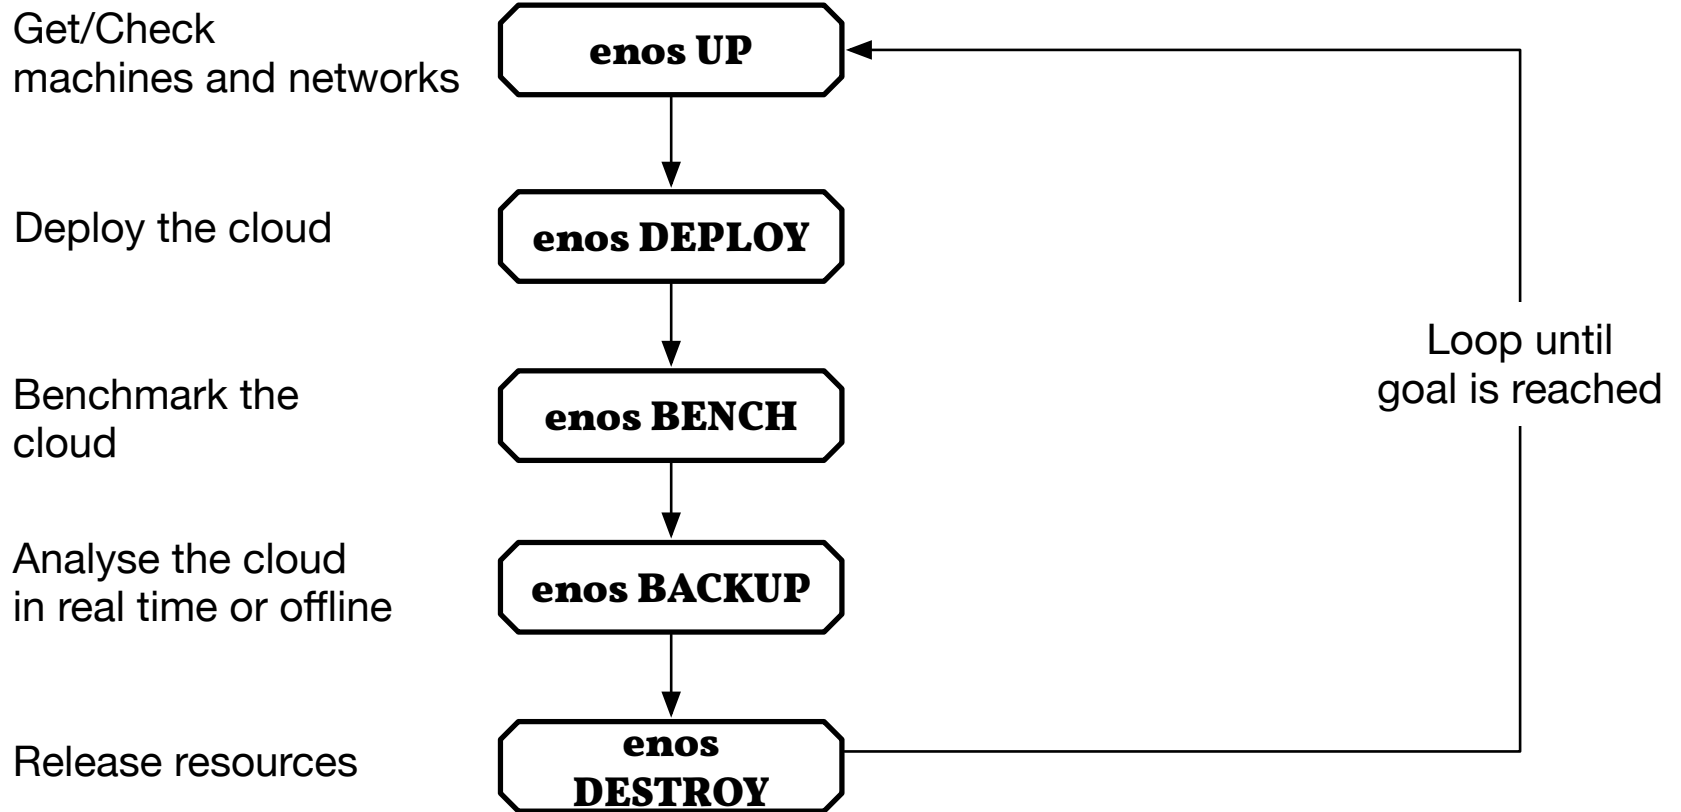
EnOS : Experimental ENvironment for OpenStack

A common good for the IPL :


- Take care of the deployment
- Can be customized easily
- Allow the emulation of geo-distributed ressources

Ronan-Alexandre Cherrueau, Dimitri Pertin, Anthony Simonet, Adrien Lèbre, Matthieu Simonin:
Toward a Holistic Framework for Conducting Scientific Evaluations of OpenStack. CCGrid2017

EnOS workflow



EnOS workflow : up



Different providers :

- Local machines
 - virtual box
 - libvirt
- Testbeds
 - Grid'5000
 - OpenStack
 - Chameleon Cloud

EnOS workflow : deploy

Configuration, Environment




A cloud deployed

small-sized deployments
(100 machines)

- approx 500 agents to deploy

EnOS workflow : deploy

Configuration, Environment




A cloud deployed

EnOS is flexible:

- Custom topology
- Different scales

EnOS workflow : bench

Configuration, Environment



Benchmarks report

Integrated evaluation tools:

- Rally (control plane)
- Shaker (data plane)
- OSProfiler (tracing)

EnOS workflow : backup

Configuration, Environment



enos BACKUP




A tarball with settings/results

Backups include


- Logs/Configurations
- Benchmark reports
- Metrics gathered

EnOS workflow : netem

Configuration, Environment




Network emulated



03

EnOS: Case studies

Monitoring functions placement



Mohamed Abderrahim, Meryem Ouzzif, Karine Guillouard, Jerome Francois, Adrien Lèbre. A Holistic Monitoring Service for Fog/Edge Infrastructures: a Foresight Study. *The IEEE 5th International Conference on Future Internet of Things and Cloud (FiCloud 2017)*, Aug 2017, Prague, Czech Republic.



Large scale deployment

Achievements:

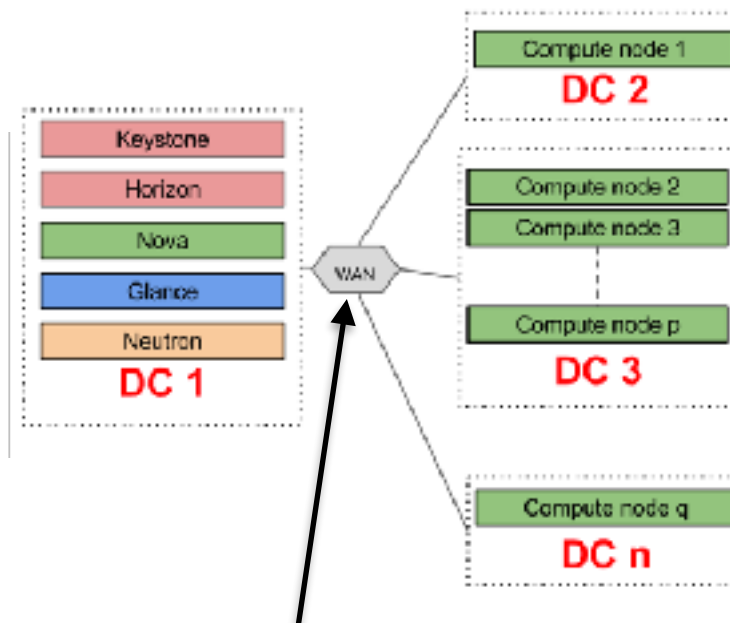

- « Chasing 1000 nodes scalability »
- Joint Work with Mirantis
- G5K official listed as official testbed for OpenStack performance evaluation



OpenStack WANWide



Collaboration




Network emulation :
latency/bandwidth/loss

OpenStack WANWide



Collaboration



Achievements:

- Experiments run on **Grid'5000** and **Chameleon** independently
- Fully automatized
- 250 benchmarks (approx. 100 running hours) on each testbeds
- Results followed the same trends
- experimental setup : <https://github.com/beyondtheclouds/enos-scenarios>
- results : <https://enos.irisa.fr/html>

OpenStack IOT

FBK (Italy) - FEMDC active members



- OpenStack with EnOS on an IOT use case
- Results will be presented @Openstack day Italy (Milan 28 Sept. 17)
- EnOS contributions

Fed4fire+



- Benchmarks comparisons between Open Nebula and OpenStack
- Results will be presented in the next engineering conference (Volos 4-6 Oct. 17)
- EnOS contributions