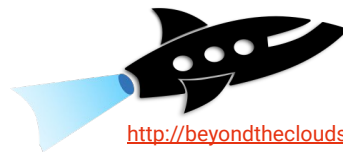


Collaborative On-Demand OpenStack Clouds and Beyond

— Collaboration by composition with the scope-lang, IPL Discovery



<http://beyondthecLOUDS.github.io>

Managing Resources of an Edge Infrastructure?

— What does it mean?

Edge Infra?

A kind of Distributed Cloud Infrastructure

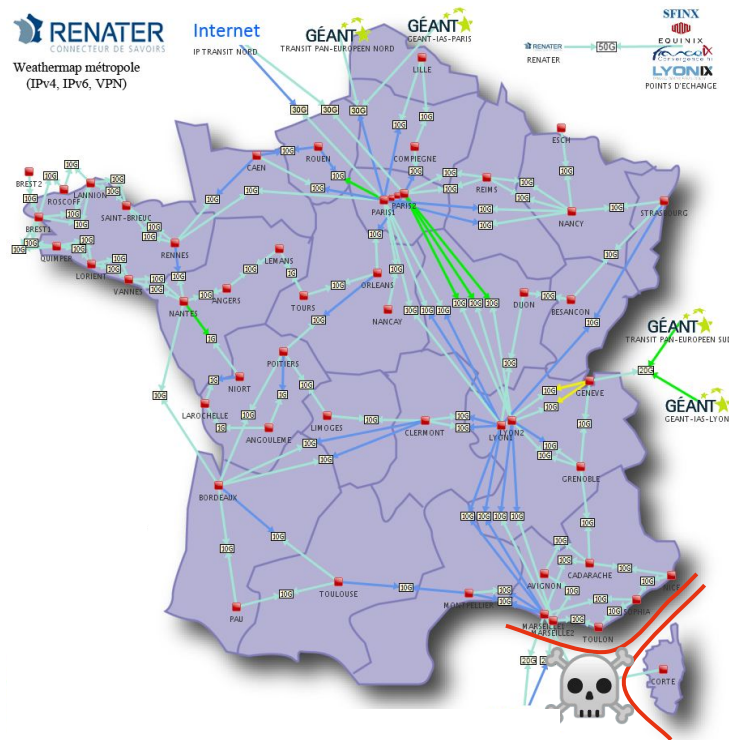
Particularities

- 100s/1000s of locations (*i.e.*, Data Centers)
- Dozen of servers per Data Center
- WAN links (10 to 300 ms RTT)
- Intermittent connectivity
- Network partitioning issues

Example of an Edge Infrastructure

Renater backbone

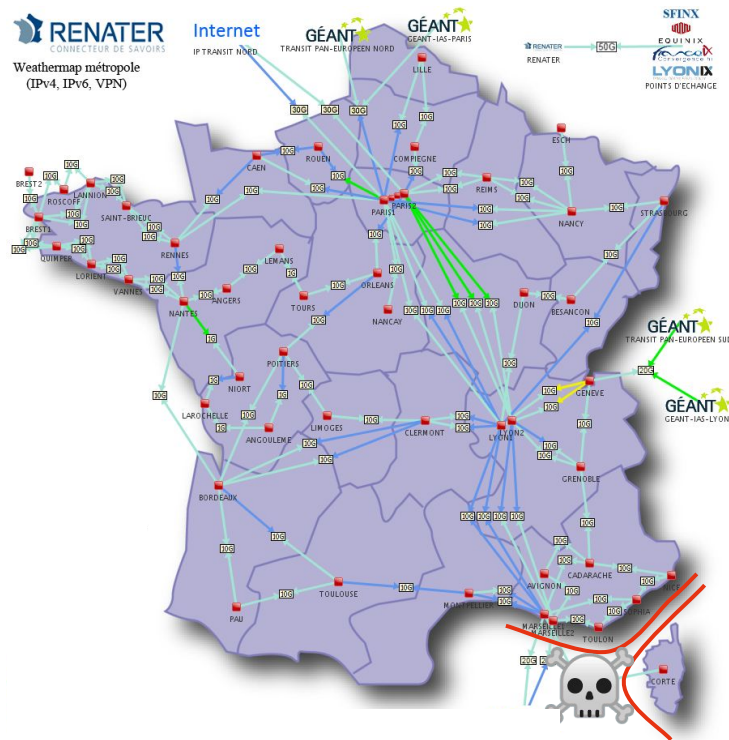
- Point of Presence (PoP) in red
- Micro DC in each PoP
 - Dozen of servers
- WAN links interconnect PoPs
 - 10ms, Paris ↔ Marseille
 - 300ms, Paris ↔ Vancouver
- Net. partitions risks (💀) between PoPs
 - Marseille/Corte



Managing Resources of an Edge Infra?

Same as in Cloud Computing. **Tuned** for the Edge.[‡]

1. Operate/use a **single DC**
 - Manage users, flavors, quotas
 - Provision computes, storages, nets
2. Operate/use **several DCs**
 - **Cross-DC** collaborative provisioning (intra/inter services)
 - Manage **multiple DC simultaneously**
3. Robustness w.r.t. network delay & **disconnections**
 - Access/Manage reachable resources (full isolation)



Managing Resources of an Edge Infra with OpenStack

— Review past and ongoing actions.

Red Thread: Boot of a VM

```
openstack server create my-vm --image Debian
```

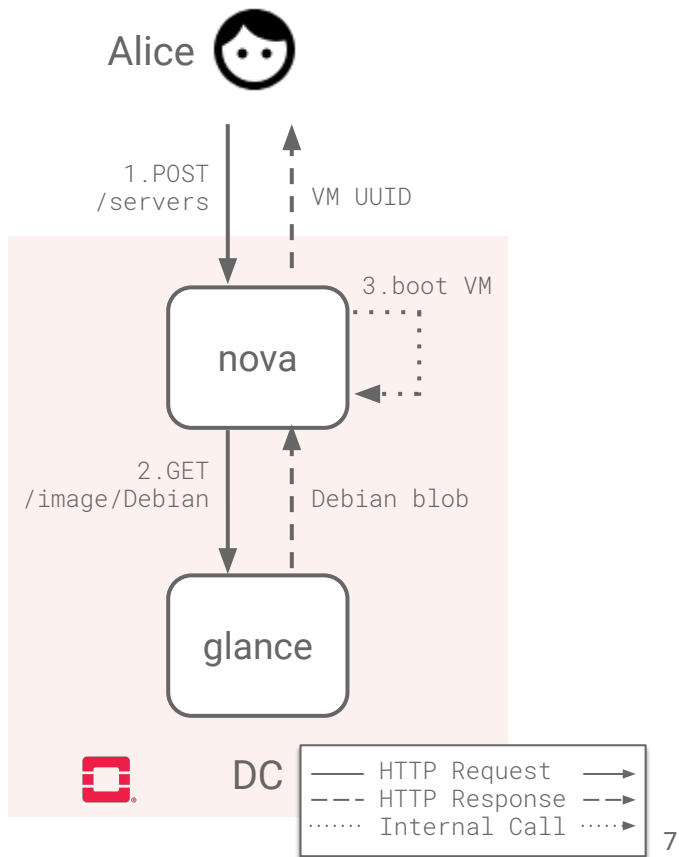
Boot of a Debian VM (simplified)

1. Operator requests a boot to nova
2. Nova contacts glance to get Debian
3. Nova boots VM internally

Boot VM scenarios

- in a single DC (**1-DC**)
- in one DC with an image from another DC (**x-DC**)
- in multiple DC (***-DC**)
- **Globally** vs. **partially** connected infra.

Boot VM	1-DC	x-DC	*-DC
global	??	??	??
partial	??	??	??



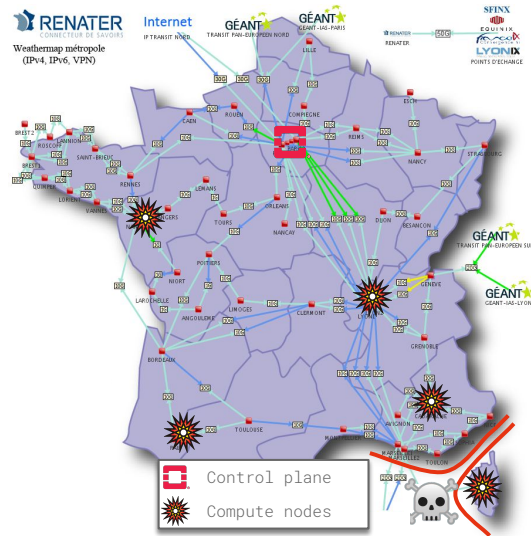
Approach 1: Centralized Management

*One DC hosts the **control plane**; Other DCs host **compute nodes***

Theoretically, a normal OpenStack deployment

Practically, a lot of issues/challenges^{‡ †}

- Impact of latency, throughput, intermittent connectivity
- What are the deployment rules for each service?
- Deployment/Upgrade of the system



Boot VM	1-DC	x-DC	*-DC
global	✓	-	-
partial	x	x	x

→ **Operational, but focuses on specific use-cases**

[‡]. <https://www.openstack.org/videos/summits/boston-2017/toward-fog-edge-and-nfv-deployments-evaluating-openstack-wanwide>

[†]. <https://www.openstack.org/videos/summits/Paris-2018/rabbitmq-or-qpid-dispatch-router-pushing-openstack-to-the-edge>

Approach 2: Distributed Management

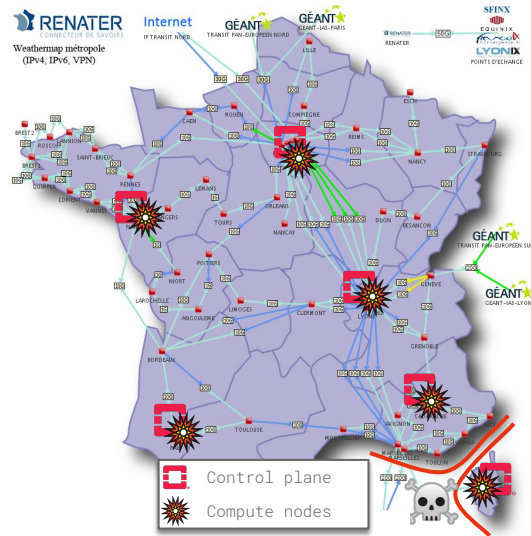
Every DC is one OpenStack that **collaborates** with others (à la p2p)

Theoretically, should fulfill our needs

- One **autonomous** control plane per DC (*partial* ✓)
- DCs are **collaborative** with each other
 - Share resources with others (*benefits from natural sharding*)
 - Replicate resources at some locations (*preserve from delay/partial*)

Practically, a sophisticated solution

- Implementing collaboration is a **conundrum**
- *OpenStack doesn't provide a general solution*



Boot VM	1-DC	x-DC	*-DC
global	✓	✓	✓
partial	✓	✓	✓

Theoretical

Service-to-Service Collaboration

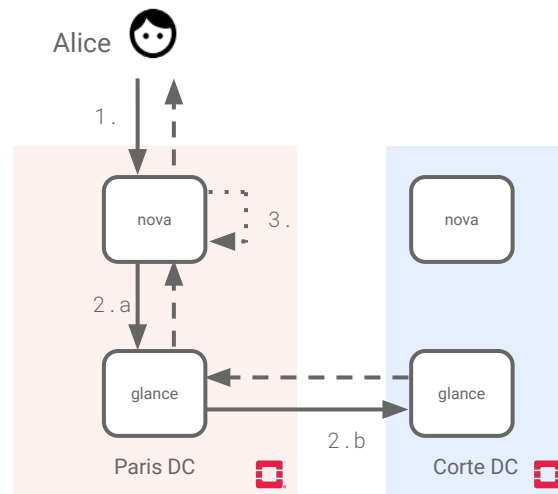
Make the *service natively collaborative* ($K2K^\ddagger$, Glance to Glance[†])

Pro

- Efficient/Optimal implementation (optimistically)

Issues

- **Tangle sophisticated collaboration code** with **vanilla code**
 - Force core developers to maintain collaboration code, make new features collaborative
 - Intrusive collaboration is not an option for some services (not everyone wants to do edge or need collaboration)



→ **Collaboration code should be decoupled from vanilla code**

Boot VM	1-DC	x-DC	*-DC
global	✓	✓?	✓?
partial	✓?	✓?	✓?

[‡]. <https://docs.openstack.org/security-guide/identity/federated-keystone.html>

[†]. https://wiki.openstack.org/wiki/Image_handling_in_edge_environment

Broker Collaboration

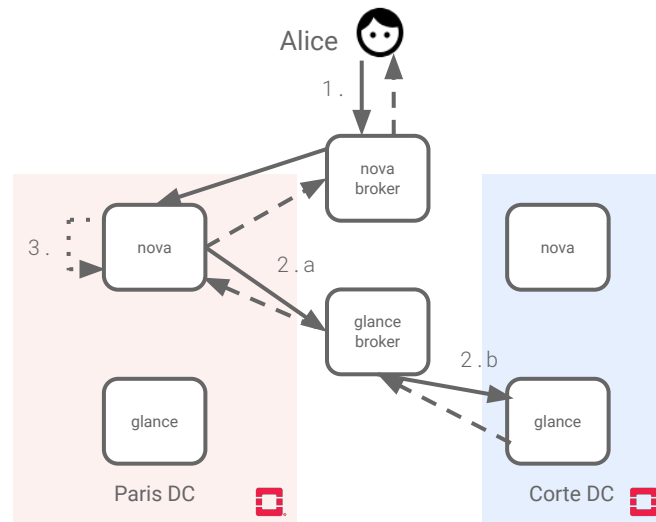
Broker on top orchestrates the collaboration (Tricircle[‡], Mixmatch[†], ...)

Pro

- Put **collaboration code outside** of vanilla code (in the broker)
- Enable **enhancement of APIs** for sharing/replication

Issues

- Current implementations
 - Rely on a central broken (partial: **X**)
 - Miss mechanism for replication (*-DC: **X**)
- **Broker** has to be **exhaustive** with the underlying APIs
 - **Lot of code** to simply **expose APIs** at **broker level**



→ **Broker should not reimplement API to the risk of developing a new OpenStack on top of OpenStack**

Boot VM	1-DC	x-DC	*-DC
global	✓	✓	X
partial	X	X	X

[‡]. <https://wiki.openstack.org/wiki/Tricircle>

[†]. <https://mixmatch.readthedocs.io/en/latest/>

DataBase Collaboration

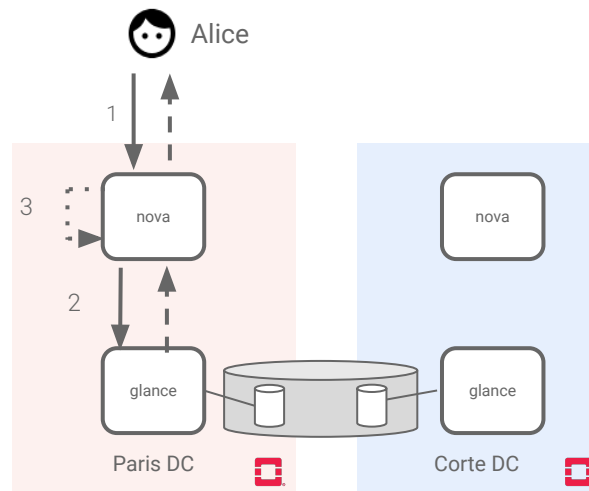
Implement **collaboration** by making **resources reachable** via the **DB** (Rome, active-active Galera, CockroachDB, ...)†‡

Pro

- Do not need to modify OpenStack code
- Allow data replication (mitigate delay/partial ✓)

Issues

- Galera: replicate data at **every location** does not scale
- Maintain **consistency** of **all data across all DCs** forbids writes in network partition (partial ✗)
- DataBase **only considers data**
 - A **resource** is made of **data and effects**
 - Collaboration via DB **misses effects** (x-DC/*-DC: Keystone ✓, Neutron ✗, ...)



Boot VM	1-DC	x-DC	*-DC
global	✓	✓, ✗	✓, ✗
partial	✗	✗	✗

→ **Resources could not be global**
→ **Resources have to come with there side effects**

†. <https://www.openstack.org/videos/summits/vancouver-2018/keystone-in-the-context-of-fogedge-massively-distributed-clouds>

‡. <https://hal.archives-ouvertes.fr/tel-01416099v1>

Collab. techno.	Issue	What we want
Service-to-service	Tangled invasive collab. code → Not always an option	Ad-hoc collaboration → Modular confined collab. code → Plugged if need be
Broker	Re-implement API and logic at Broker → Do the job twice → Specific code for one app.	Generic collab. for any app. → Don't do the job!
DataBase	Replicate data everywhere (Galera) → Do not scale at Edge Maintain consistency everywhere → Forbid partial writes Omit effects → Work for data only app.	On-demand replication → Keep clusters small for scaling On-demand collaboration → Keep clusters small for resiliency Include effects → Work for any app.

Draw Conclusions

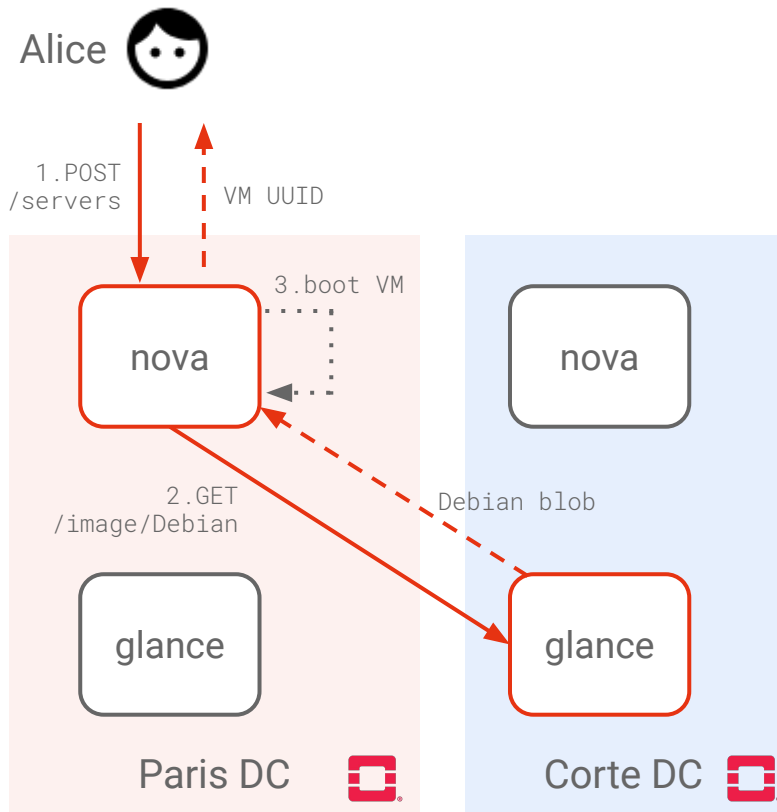
Distributed Management with the scope-lang

— Ad-hoc, generic, on-demand and effectful collaboration.

scope-lang

Alice defines the **scope** of the request into the CLI. The **scope** specifies **where the request applies**

```
openstack server create my-vm
--image Debian
--scope { nova:    Paris
        , glance: Corte }
```



**Boot VM in Paris with Debian
from Corte**

scope-lang in Actions

Scope for **1-DC** operations

```
OS@Paris$ openstack server create my-vm --image Debian  
scope {nova: Paris, glance: Paris}
```

Scope for **x-DC** operations

```
OS@Paris$ openstack server create my-vm --image Debian  
--scope {nova: Paris, glance: Corte}
```

Scope for ***-DC** operations

```
OS@Paris$ openstack server create my-vm --image Debian  
--scope {nova: Paris&Corte, glance: Corte}
```


*-DC: and '&'

Do the operation **here** and **there**

- Create a user in Paris and Corte

```
openstack user create Alice  
  --password-prompt  
  --scope {keystone: Paris&Corte}
```

- List VMs in Paris and Corte

```
openstack server list  
  --scope {nova: Paris&Corte}
```

Properties

- **On-demand** partial replication
 - Replication at scope locations
 - Keep cluster small for scaling and resiliency
- Query **multiple** DCs at once

*-DC: or '|'

Do the operation **here** or **there**

- Boot a VM in Paris with image from Corte or Marseille

```
openstack server create my-vm
  --image Debian
  --scope { nova: Paris
           , glance: Corte|Marseille }
```

Properties

- Let the operator **implements retries workflow**

No matter if one is down or don't have the image, till the other is up and has the image.

Manage	Needs	scope
Single DC	Manage resources locally <ul style="list-style-type: none"> boot VM in Paris List VMs in Corte 	1-DC^{†‡} <ul style="list-style-type: none"> {nova:Paris, glance:Paris} {nova:Corte}
Multiple DCs	Cross-DC collaboration <ul style="list-style-type: none"> boot VM in Paris with Debian from Corte Manage resources simultaneously <ul style="list-style-type: none"> create image in Paris and Corte boot VM in Paris with Debian from Corte or Marseille 	x-DC[†] <ul style="list-style-type: none"> {nova:Paris, glance:Corte} *-DC <ul style="list-style-type: none"> {glance:Paris&Corte} {nova:Paris, glance:Corte Marseille}

†. **PoC**: <http://github.com/BeyondTheClouds/openstackoid>

‡. Unnecessary; scope is implicitly local

Managing Resources of an Edge Infrastructure with the scope-lang

How to Implement the scope-lang?

— Collaboration by composition

Collaboration by Composition Principle

Service is an **abstraction** that

- **Encapsulates** the management of one or various **resources**
- Provides a **transparent access** to these resources
- Relies on **composition** to build an **app**.

**Modularity says we can
*interchange one service by
another one with the same
abstraction***

E.g., the OpenStack app.

- Nova/Glance web services **encapsulate** the management of **computes/images**
- Web services **REST API provide access** to these resources
- Web services **composition** builds the **manager**

Collaboration by Composition + scope-lang

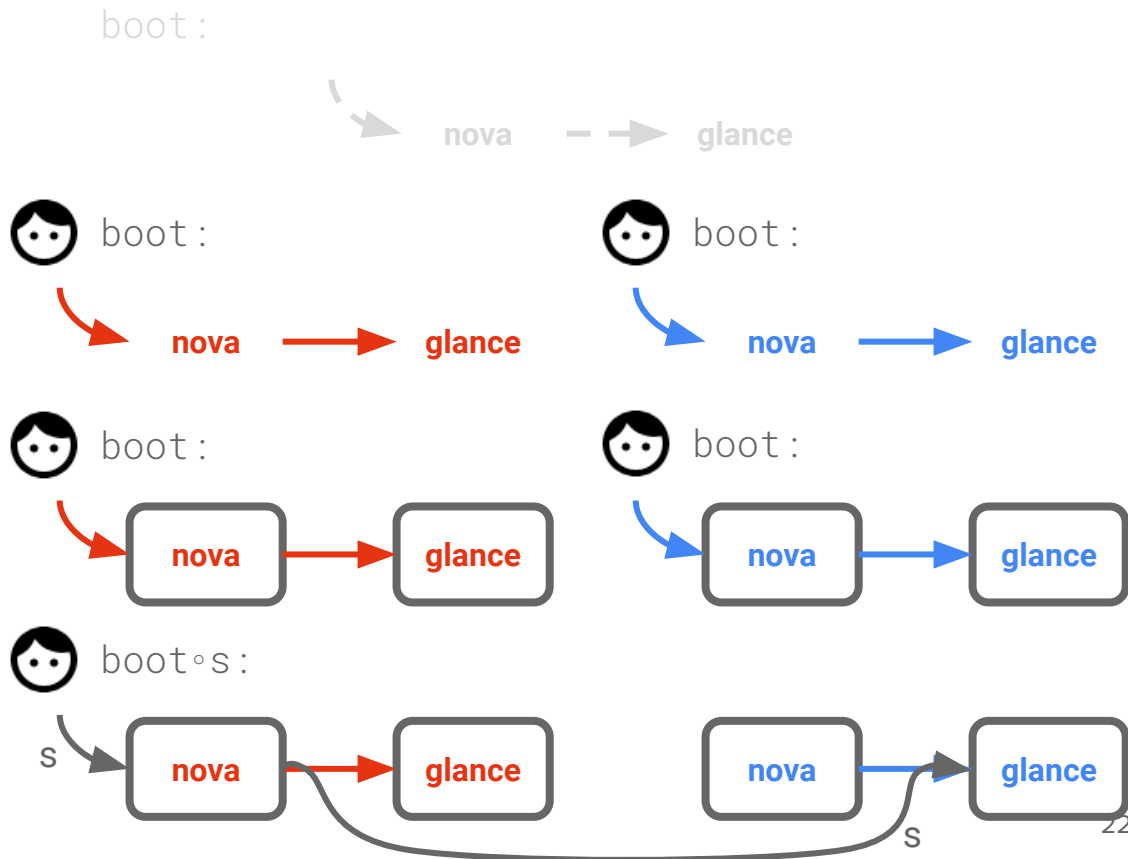
CLI operation: boot

Two DCs:

- Paris
- Corte

Wrappers intercept service calls. But, proceed by default.

A **scope** specifies the OS composition:
 $s = \{\text{nova: Paris, glance: Corte}\}$. In presence of a scope, wrappers forward service calls to the correct OS.



The Next 700 Brokers

The scope-lang wrapper/interpreter is

- **Generic**
 - Changing the composition does not require a code specific to an app.
- **Ad-hoc**
 - Changing the composition does not presume any intent of collaboration at start
- **Effectful**
 - Encapsulation property of services ensures to tracks effects
- **On-demand**
 - Define the scope of a collaboration

→ **Works for any app. built by service composition and under well modularization**

Limitations & Problems

Limited to **inter-service** collaborations

- Boot VM in Paris with image in Corte (Nova to Glance service compo. ✓)
- Live-migrate a VM from Paris to Corte (Nova to Nova intra-service operation ✗)

Risk of **bad collaborations** (x-DC)

- **Resources unreachability:** boot a VM in Paris with local network in Corte; it is not yet possible to extend network resources across DCs (API limitations, technical issues)
- **Local state:** verify in Keystone of Paris the Glance service token from Corte

Ph. D. thesis starts on Oct. 2019

- Formally define the semantic of the scope-lang (&, |, ...)
- Statically reject bad collaborations

Wrap Up

Takeaway

- **Collaboration** between Edge should be done **on-demand** (and only if needed)
 - Cope with thousands of independent sites
 - Deal with Intermittent connections
- **scope-lang** for **on-demand** collaboration
 - 1-DC: {nova:Paris, glance:Paris}
 - x-DC: {nova:Paris, glance:Corte}
 - *-DC: {nova:Paris&Corte}
 - x/*-DC: {nova:Paris&Corte, glance:Corte|Marseille}
- Harness **service abstraction** for a **collaboration by composition**
 - Generic, Ad-hoc and effectfull approach
 - **Tested** on OpenStack[‡]. **Should work** for K8S, ShareLaTeX, ...